

From Zero to Senior Data Science

A Guide into
Advanced Data Science Standards
and Big Data Approaches



By Andreas Traut

From Zero to Senior Data Science

A Guide into Advanced Data Science Standards and Big Data
Approaches

Andreas Traut

CC BY-NC-SA 4.0 Licence

This work is licensed under the Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License.

To view a copy of this license, visit

<http://creativecommons.org/licenses/by-nc-sa/4.0/>

or send a letter to

Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.

1. Edition 2021

Author: Andreas Traut

E-Mail: andreas.traut@freedommail.ch

GitHub: <https://github.com/AndreasTraut>

Homepage: <https://andreas-traut.gitbook.io/>

Contents

1	Introduction	1
1.1	Aim	1
1.2	Who is a “Data Scientist”?	1
1.2.1	Term “Data Scientist” is not Protected	1
1.2.2	Generally Accepted Definition of “Data Science”	2
1.2.3	Minimum Required Years of Working Experience for a “(Senior) Data Scientist”	3
1.3	Structure of this Book	4
1.3.1	Installation of Data Science Tools	4
1.3.2	Visualization of Different Datasets	4
1.3.3	Machine Learning with Python: Comparison Small Data vs Big Data	5
1.3.4	Big Data: Map-Reduce and K-Means Clustering	6
1.3.5	Use Cases of Artificial Intelligence in the Industry	6
1.4	My Qualification	6
1.5	Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License .	7
2	Installation of Data Science Tools	9
2.1	Jupyter-Notebooks and Spyder-IDE	9
2.2	Motivation for IDEs	11
2.3	Docker	11
3	Visualization of Different Datasets	13
3.1	Examples	14
3.1.1	Consumer Price Index Example	15
3.1.2	Last-FM Statistics of my Songs	24
3.1.3	Marathon Runtimes: Finding Systematics	32
3.1.4	Pedestrians during the first Corona-Lockdown	43
3.1.5	Station Elevators of Deutsche Bahn: API Example	47
3.1.6	Interactive Pivots	52
3.2	Introduction to Big Data Visualization Techniques	54
3.2.1	Basic Problems in Big Data Visualizations	54
3.2.2	Vermont Payments Example	55
3.3	Visualize Data with Data Apps	62

3.4	Professional Tools	65
3.4.1	Power BI	66
3.4.2	Tableau	67
3.4.3	QLink	70
4	Machine Learning with Python: Comparison Small Data vs Big Data	73
4.1	Movies Database Example	73
4.1.1	Import the Data	73
4.1.2	Separate “NaN”-Values	75
4.1.3	Visualization of the Data	76
4.1.4	Draw a Stratified Sample	79
4.1.5	Split of Dataset into Training-Data and Test-Data	81
4.1.6	Create a Pipeline	83
4.1.7	Fit the Model with “DecisionTreeRegressor”	84
4.1.8	Cross-Validation	85
4.1.9	Test the model	85
4.1.10	Conclusion	89
4.2	“Small Data”: Machine Learning using Scikit-Learn	90
4.2.1	Create Index (1)	93
4.2.2	Discover and Visualize the Data to Gain Insights (2)	94
4.2.3	Prepare for Machine Learning (3)	95
4.2.4	Use “Imputer” to Clean NaNs (4)	96
4.2.5	Treat “Categorial” Inputs (5)	96
4.2.6	Custom Transformer and Pipelines (6)	96
4.2.7	Select and Train Model (7)	98
4.2.8	Cross-Validation (8)	99
4.2.9	Save Model (9)	100
4.2.10	Optimize Model (10)	100
4.2.11	Evaluate final model on test dataset (11)	104
4.3	“Big Data”: Machine Learning using Spark ML Library	105
4.3.1	Get Ready	106
4.3.2	Initialize Spark (0)	110
4.3.3	Cleaning the data (1)	112
4.3.4	Model-specific preprocessing (2)	115
4.3.5	Aligning and numerating Features and Labels (3)	118
4.3.6	Pipelines (4)	118
4.3.7	Training data and Testing data (5)	119
4.3.8	Apply models and evaluate (6)	119
4.3.9	Minhash und Local-Sensitive-Hashing (7)	120

4.4	Summary Mind-Map	120
4.5	Future learnings and coding & data sources	121
5	Big Data: Map-Reduce and K-Means Clustering	123
5.1	Map-Reduce Example	123
5.1.1	Word Count Example	123
5.1.2	Term Frequency–Inverse Document Frequency (TF-idf) Example	129
5.2	K-Means Clustering Algorithm	139
6	Use Cases of Artificial Intelligence in the Industry	147
6.1	AI Explained	148
6.1.1	What kind of processes are meant here?	148
6.1.2	What data does artificial intelligence have access to?	149
6.1.3	Is Artificial Intelligence Really Intelligent?	149
6.2	Example: Benefits of AI for BMW	150
6.3	Implementation of AI Techniques	152
6.3.1	What can be seen in the picture?	152
6.3.2	Which groups can be formed?	154
6.4	Recommendations For Implementing AI	157

1 Introduction

1.1 Aim

The goal of this book is to guide you to becoming a Senior Data Scientist without too much prior knowledge. My aim is to help you to get started in the Data Science topics and explain the necessary knowledge to become a Senior Data Scientist. It is important to me to give you many practical examples to download and try out yourself, as well as to help you to install some very important Data Science tools (like Docker, Jupyter-Notebooks,...) on your own computer.

I am not afraid to refer to good tutorials or trainings, which are necessary from my point of view to become a Senior Data Scientist. For me it doesn't make sense to re-document here every aspect that is already described much better somewhere else. My aim is more to connect many of these great sources here in my book. I want to you to get started and my motivation is to enable you to continue working independently yourself with these references to further documentation. As soon as you have understood the topics from my book, it will be time to deal with the subject-specific documents anyway.

I believe in the Open-Source¹ spirit of sharing knowledge to other people. Therefore this book is subject to the CC BY-NC-SA 4.0 Licence (see sec. 1.5) and is available free of charge.

I am also willing to improve this book in the future and spend as much as I can to motivate and educate many people to become a Senior Data Scientist. Don't hesitate to contact me in case of suggestions for improving this book. I hope you enjoy reading and experimenting.

1.2 Who is a “Data Scientist”?

1.2.1 Term “Data Scientist” is not Protected

Unfortunately, the term “Data Scientist” is not a protected term! Anyone can call themselves a “Data Scientist” without any training or certificate or any kind of proof! This amazes and shocks me, because Data Science is a very promising area for the future. Even HR managers struggle in correctly evaluating the meaning and requirements of “Data Science”.

¹ Open-source model, Wikipedia, https://en.wikipedia.org/wiki/Open-source_model

What each individual understands by “Data Science” is as elastic as chewing gum. Nevertheless, it is used as a job title in many job descriptions, often with completely different interpretations. For example: one person writes “Data Scientist” in a job profile, but thinks that in-depth knowledge of the company and its manufacturing processes is urgently needed, and sorts out applicants according to this criterion. This person denies that fact, that programming skills with machine learning libraries is also very essential in order to be a “Data Scientist”.

Another person writes “Data Scientist” in a job profile and thinks that you have to have many years of profound programming experience and lists programming languages and tools that you can only learn and understand in depth in two lifetimes ignoring that mathematical skills are also required to be a Data Scientist.

People who describe themselves as a “Data Scientist,” do not have to worry about any legal or employment lawsuits. If, on the other hand, someone describes himself as a “business economist” and cannot back this up with certificates or diplomas, then he risks to be threatened with lawsuits. For a profession of the future, such as “Data Scientists” are, this is a bitter realization.

1.2.2 Generally Accepted Definition of “Data Science”

Is there a generally accepted definition for “Data Science / Data Scientist”? Yes! There is a much discussed and generally quite accepted definition: the “**Data Science Venn Diagram**”². If you are not yet familiar with this “Data Science Venn Diagram”, I would strongly recommend that you look into it! The misconceptions of what a “Data Scientist” is and what they can do are unfortunately still huge.

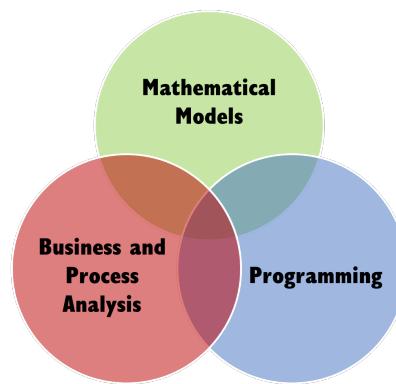


Figure 1.1: Data Science - Adapted version. [Creative Commons Legal Code](#)

For me, a Data Scientist is a person who is 30% mathematician, 30% understanding the business and processes (which requires some years of working experience, see sec. 1.2.3), 30% programming experience with machine-learning libraries and 10% passion for new exciting topics and willingness to dive into new tools.

² Data Science Venn Diagram, <http://drewconway.com/zia/2013/3/26/the-data-science-venn-diagram>

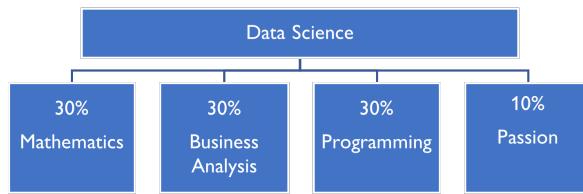


Figure 1.2: Data Science

So “Data Science” is a mix of mathematical knowledge, machine-learning programming experience and business knowledge. The balanced intersection of these areas makes a good Data Scientist.

It makes no sense to me to overvalue one of these areas and ignore one of the others! For example: someone who has profound in-depth knowledge of a company and its manufacturing processes but who lacks mathematical skills or programming skills with machine-learning libraries is in my opinion not a “Data Scientist”. You always need all parts – mathematics, programming and business knowledge – to be a Data Scientist!

1.2.3 Minimum Required Years of Working Experience for a “(Senior) Data Scientist”

How long should you have worked in order to name yourself a “Senior Data Scientist”? I think that you should have at least 1-2 years of working experience before you can name yourself a “Data Scientist” and you should have at least 2-3 years of working experience before you can name yourself a “Senior Data Scientist”. I would substantiate this assertion as follows: the **“Certification Manual for Data Science”³** from **“Fraunhofer-Institut für Intelligente Analyse- und Informationssysteme IAIS”⁴** explains their requirements for certifying someone as a “Data Scientist”. In order to become a “Data Scientist Basic Level” you need

- completed studies at a university (or equivalent)
- at least 1 year of work experience in a data science relevant field.
- the Certificate “Data Scientist Basic Level”, which you can earn after passing a 3 hours written exam at the Fraunhofer-Institute

In order to become a “Senior Data Scientist” you need

- completed studies at a university (or equivalent)
- at least 2 years of work experience with at least 1 year in a data science relevant field
- the Certificate “Data Scientist Basic Level”, which you can earn after passing a 3 hours written exam at the Fraunhofer-Institute

³ Fraunhofer Zertifizierungshandbuch “Data Science”, https://www.personenzertifizierung.fraunhofer.de/content/dam/zertifizierung/de/documents/Zertifizierungshandbuch_Data_Scientist.pdf

⁴ Fraunhofer-Institut für Intelligente Analyse- und Informationssysteme IAIS, <https://www.iais.fraunhofer.de/de/data-scientist-schulungen.html>

- a Specialized Certificate (several options like “Big Data” or “Deep Learning” and others are available), which you can earn with a 3 hours written exam at the Fraunhofer-Institute
- a 40 paged expose (similar to a bachelor thesis) approved by Fraunhofer and related to a data science topic

Besides this I found various articles among them also one on the Online-Magazine “Towards Data Science”⁵. It says, that

- after 2 years working experience have skills on Python, Jupyter and kaggle, which is required to name yourself a “Junior Data Scientist” and
- after 3-5 years you have skills on Docker, Cloud and APIs, which are required to name yourself a “Senior Data Scientist”.

This excludes any freshly graduated student from an university to name himself a “Data Scientist” or even a “Senior Data Scientist”! Underestimating work experience is one common issue, that also HR managers face, when they evaluate a “Data Scientist” job application.

I hope that one day in the future the term “Senior Data Scientist” will be protected a bit better with certificates or other proofs of qualification. At the moment, unfortunately, “Data Scientist” and “Senior Data Scientist” is a very elastic and stretchable term, which everyone likes to interpret for himself as he wants to have it.

1.3 Structure of this Book

1.3.1 Installation of Data Science Tools

In this chapter (see sec. 2) I will explain how to install the basic Data Science Tools, which you need. In my examples I use Jupyter-Notebooks⁶, which is a widespread standard today, but I also use an Integrated Development Environment⁷. In my opinion Jupyter-Notebooks are good for the first examinations of data and for documenting procedures and up to a certain degree also for sophisticated data science. But it is a good idea to learn very early how to work with an IDE. In my opinion Jupyter Notebooks are not always the best environment for learning to code! Therefore I will give you a short introduction into an IDE and highly recommend that you learn how to work with an IDE.

1.3.2 Visualization of Different Datasets

In the data scientist environment the visualization is as important as the analysis itself. I worked on different datasets with the aim to visualize the data and in the **first part** of this chapter (see sec. 3.1)

⁵ “Towards Data Science”, <https://towardsdatascience.com/becoming-a-level-3-0-data-scientist-52641ff73cb3>

⁶ Jupyter-Notebook, <https://jupyter.org/>

⁷ Integrated development environment, Wikipedia, https://en.wikipedia.org/wiki/Integrated_development_environment

I will explain these examples. I used python and libraries like e.g Matplotlib⁸ or Seaborn⁹, which are available for free. I will show you in Section sec. 2 how to install these tools on your own computer.

Each of the datasets, which I worked on, contains different topics of necessary preliminary work before I could visualize them, e.g. converting dates or numbers, adding/extracting information and so on. I will show you how this can be done.

In the **second part** (see sec. 3.2) I will explain some Big Data visualizations problems and techniques to solves these.

In the **third part** (see sec. 3.3) I will show you how to visualize and share the data with a “data app”. Data Scientist often forget, that all models, visualizations, which they have built, need to be used by someone, who is probably not as skilled in all these technical requirements. Such “data apps” are helpful to make the data accessible very quickly for everyone on all devices (also mobile phones).

In the **fourth part** sec. 3.4 I will list some common professional tools, which offer visualization functionality and more. These tools cost some money, but I recommend to have look into these: many companies use these or similar tools.

1.3.3 Machine Learning with Python: Comparison Small Data vs Big Data

After having learnt visualization techniques in Python it is time to start working on different datasets with the aim to learn and apply machine learning algorithms. In chapter sec. 4 I am particularly interested in explaining the differences and similarities of “Small Data” (Scikit-Learn) approaches versus the “Big Data” (Apache Spark) approaches. Therefore in this chapter I will to focus on this “comparison” question of “Small Data” coding vs “Big Data” coding. I haven’t seen many comparisons of “Small Data” vs “Big Data” (neither theoretically nor in coding patterns) and I think understanding this is interesting and important.

The **first example** (see sec. 4.1) is about a Movies database and the revenues, which these movies generated. My aim is to predict the revenues. I use a Jupyter-Notebook and will explain how to apply the standard procedures (e.g. data-cleaning & preparing, model-training,...).

The **second example** (see sec. 4.2) is for being used in an IDE (Integrated Developer Environment), like the Spyder-IDE¹⁰ from the Anaconda distribution¹¹ and applies the *Scikit-Learn Python Machine Learning Library*¹² (you may call this example a “Small Data” example if you want). I will show you a typical structure for a machine-learning example and put it into a mind-map. The same structure will be applied on the third example.

⁸ Matplotlib, <https://matplotlib.org/>

⁹ Seaborn, <https://seaborn.pydata.org/>

¹⁰ Spyder-IDE, <https://www.spyder-ide.org/>

¹¹ Anaconda distribution, <https://www.anaconda.com/>

¹² Scikit-Learn, <https://scikit-learn.org/>

The **third example** (see sec. 4.3) is a “Big Data” example and will use a Docker environment¹³ and apply the *Apache Machine Learning Library*¹⁴, a scalable machine learning library. The mind-map from the second part will be extended and aligned to the second example.

1.3.4 Big Data: Map-Reduce and K-Means Clustering

Big Data Approaches are very different and challenging. Someone who never heard about Big Data Approaches will be surprised about how many details and topics are very different to the standard data science approaches (like Jupyter-Notebooks on Scikit-Learn). Therefore I dedicate a separate chapter for this topic: in chapter sec. 5) I explain the Map-Reduce programming model on a Word-Count example and I show how the K-Means Clustering Algorithm in Apache Spark ML works.

1.3.5 Use Cases of Artificial Intelligence in the Industry

In the chapter sec. 6 I explain in an easily understandable way what is meant by ”artificial intelligence (AI) in the industry”, describe areas of application and illustrate with practical examples and programming applications how AI can be implemented in concrete terms.

In the **first part**, I will explain the basic concepts of AI and describe some areas where AI is already being successfully applied. I will also describe the peculiarities of big data, deep learning and process mining.

In the **second part**, I show an example of how the car manufacturer BMW has benefited from AI techniques.

In the **third part**, I show how AI techniques can be implemented in the Python programming language when dealing with the topic of image recognition and provide my programming code in the process.

In the **fourth part**, I give some recommendations on what to look for when introducing AI techniques in a company.

I think the choice of further articles to delve into the topic ”Use cases of artificial intelligence in industry” is huge and I hope this short introduction is helpful to get started.

1.4 My Qualification

I am a graduated *Diplom-Mathematician* and a *Certified Advanced Data Scientist*. I am a *Certified Data Scientist Basic Level* and a *Certified Data Scientist Specialized in Big Data Analytics*. Additionally I passed a *French bachelor degree in mathematics* during my Erasmus studies in France.

¹³ Docker environment, <https://www.docker.com/>

¹⁴ Apache Machine Learning Library, <https://spark.apache.org/mllib/>

I am also holding the *Certificate of "Data Analysis with Python: Zero to Pandas"*¹⁵ (see fig. 1.3) which covers topics like data visualization and exploratory data analysis on the basis of Python¹⁶, Numpy¹⁷, Pandas¹⁸, Matplotlib¹⁹ and Seaborn²⁰. I can recommend this course and I wish I would have found this course before I wrote this repository, because it was very helpful.



Figure 1.3: Certificate Data Analysis with Python: Zero to Pandas

1.5 Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License



Figure 1.4: Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/>

¹⁵ Certificate of "Data Analysis with Python: Zero to Pandas, <https://jovian.ai/certificate/MFQTGOBQGM>

¹⁶ Python, <https://www.python.org/>

¹⁷ Numpy, <https://numpy.org/>

¹⁸ Pandas, <https://pandas.pydata.org/>

¹⁹ Matplotlib, <https://matplotlib.org/>

²⁰ Seaborn, <https://seaborn.pydata.org/>

[sa/4.0/](https://creativecommons.org/licenses/by/4.0/) or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.

2 Installation of Data Science Tools

2.1 Jupyter-Notebooks and Spyder-IDE

I use Jupyter-Notebooks¹, which is a widespread standard today, but I also use the Spyder-IDE². The IDE stands for Integrated Development Environments³. I think you will need both of them.

The Spyder-IDE is a separate software application, where you can debug your code, which makes the development of algorithms easier - a big advantage! The disadvantage is, that you need to learn to use this software application first, which is a bit more difficult than a Jupyter-Notebook but really worth the time you spent. It looks like shown in figure fig. 2.1.

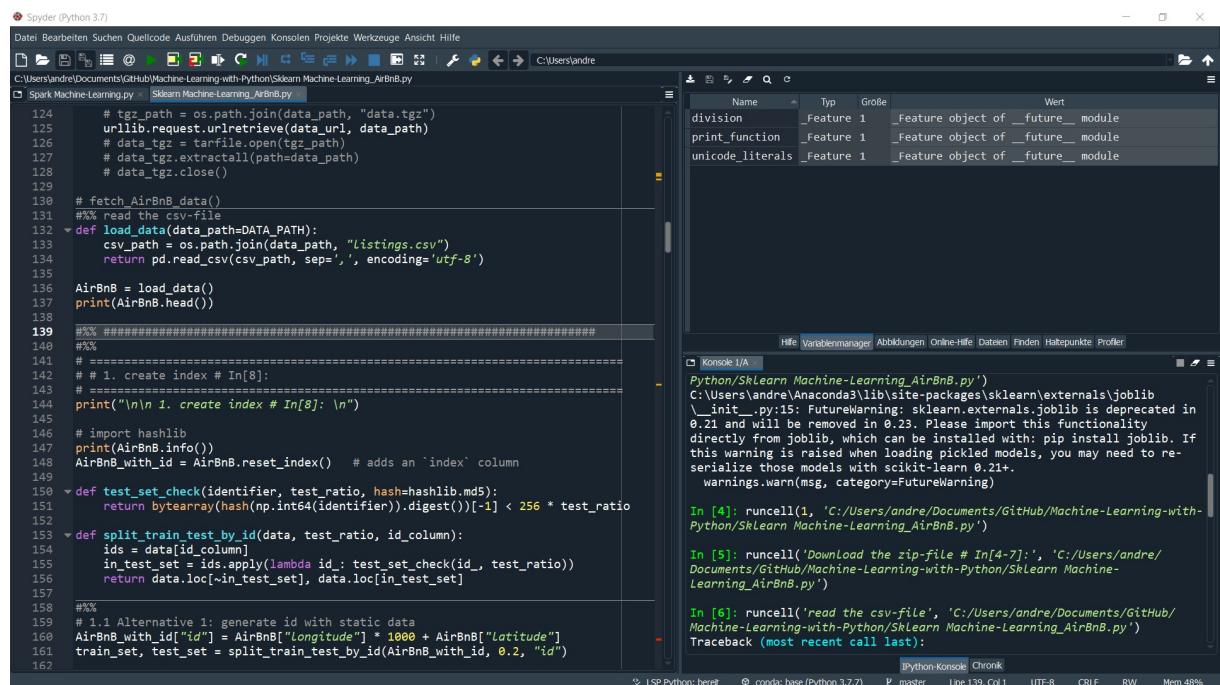


Figure 2.1: Spyder-IDE: An integrated development environment with debugger.

A Jupyter-Notebook runs in a browser (like Chrome, Edge, Firefox) and combines code and document-

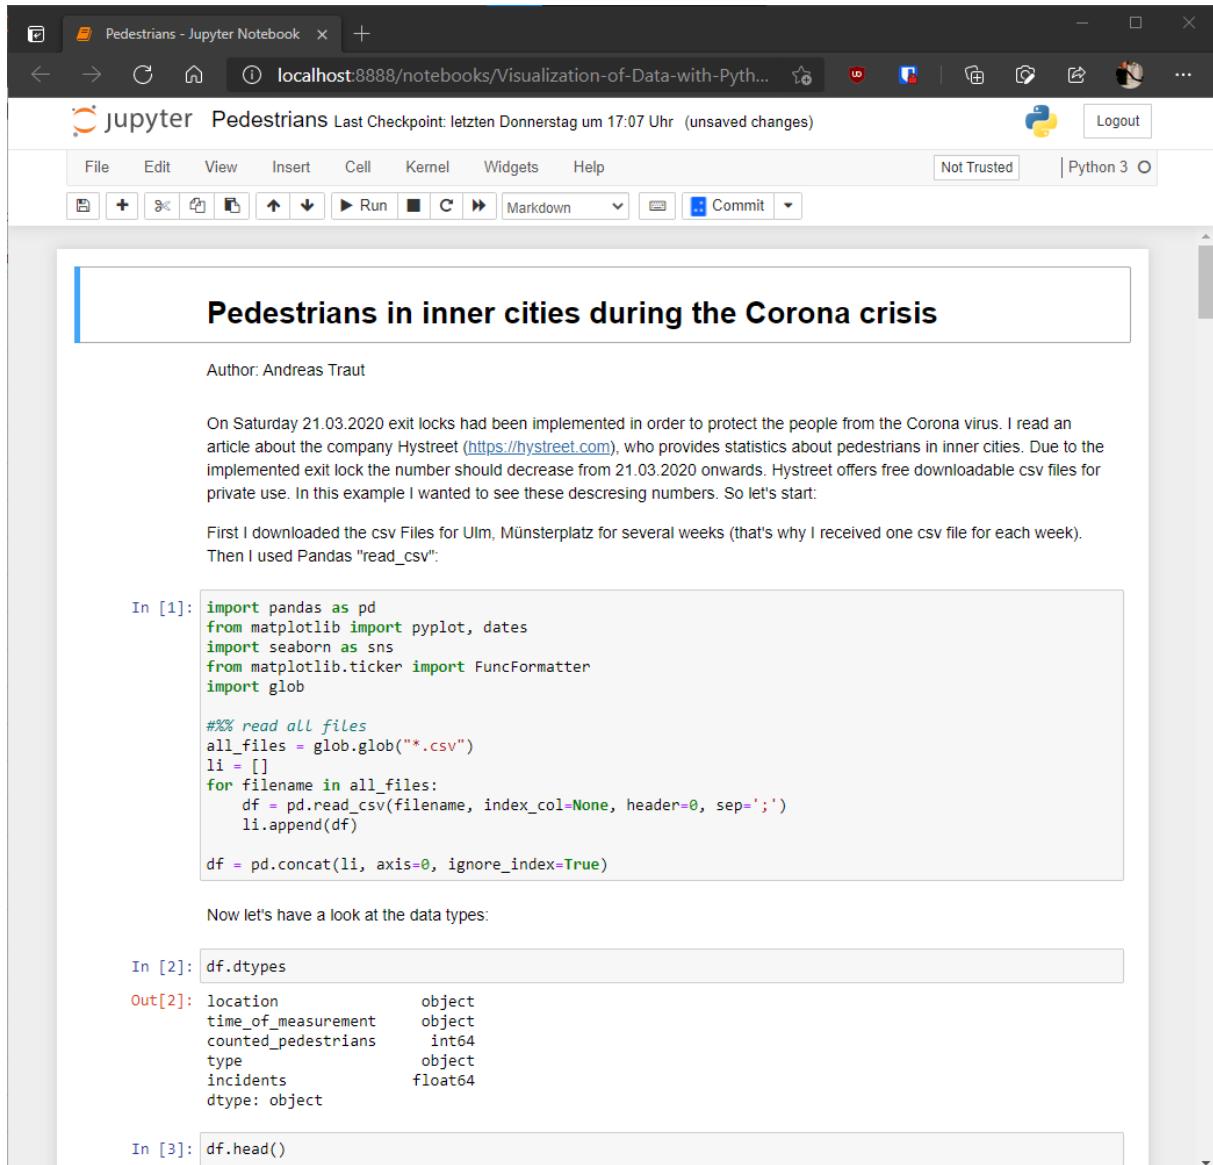
¹ Jupyter-Notebook, <https://jupyter.org/>

² Spyder-IDE, <https://www.spyder-ide.org/>

³ Integrated development environment, Wikipedia, https://en.wikipedia.org/wiki/Integrated_development_environment

tation. The advantage is, that it looks beautiful, is easy to share with other people. The disadvantage is, that there is no code-debugger. It look like shown in figure fig. 2.2.

As I think that you should get familiar with both (the Jupyter-Notebooks and the Spyder-IDE), I recommend installing the Anaconda distribution⁴. The installation is very simple and includes the Jupyter Notebooks as well as the Spyder-IDE and a lot more Data Science applications, which might be useful for you later.



The screenshot shows a Jupyter Notebook window titled "Pedestrians - Jupyter Notebook". The title bar includes the URL "localhost:8888/notebooks/Visualization-of-Data-with-Pyth..." and a "Logout" button. The menu bar has options: File, Edit, View, Insert, Cell, Kernel, Widgets, Help. The toolbar includes icons for file operations, run, and commit. A header bar shows "jupyter Pedestrians Last Checkpoint: letzten Donnerstag um 17:07 Uhr (unsaved changes)" and a Python 3 kernel indicator. The main area contains a section titled "Pedestrians in inner cities during the Corona crisis" with author "Andreas Traut". The text describes downloading CSV files from Hystreet and concatenating them. Below this, In [1] shows the Python code for reading multiple CSV files and concatenating them. In [2] shows the resulting DataFrame's dtypes, and In [3] shows the first few rows of the DataFrame.

```
In [1]: import pandas as pd
from matplotlib import pyplot, dates
import seaborn as sns
from matplotlib.ticker import FuncFormatter
import glob

# %% read all files
all_files = glob.glob("*.csv")
li = []
for filename in all_files:
    df = pd.read_csv(filename, index_col=None, header=0, sep=';')
    li.append(df)

df = pd.concat(li, axis=0, ignore_index=True)
```

Now let's have a look at the data types:

```
In [2]: df.dtypes
```

```
Out[2]: location          object
time_of_measurement   object
counted_pedestrians    int64
type                  object
incidents            float64
dtype: object
```

```
In [3]: df.head()
```

Figure 2.2: Jupyter-Notebook: Code and Documentation in one place.

⁴ Anaconda Distribution Installation, <https://www.anaconda.com/products/individual#Downloads>

2.2 Motivation for IDEs

The first Jupyter-Notebooks have been developed 5 years ago (in 2015). Since my first programming experience was more than 25 years ago (I started with GW-Basic⁵ then Turbo-Pascal⁶ and so on and I am also familiar with MS-DOS⁷). I quickly learnt the advantages of using Jupyter-Notebooks. **But** I missed the comfort of an IDE from the very first days!

Why is it important for me to mention the IDEs out so early in a learning process? In my opinion Jupyter-Notebooks are good for the first examinations of data and for documenting procedures and up to a certain degree also for sophisticated data science. But it is a good idea to learn very early how to work with an IDE. I point this out here, because after having read several e-Books and having participated in seminars I see that IDEs are not in the focus.

In my opinion Jupyter Notebooks are **not** always the best environment for learning to code! I agree, that Jupyter Notebooks are nice for doing documentation of python code. It really looks beautiful. But I prefer debugging python code in an IDE instead of a Jupyter-Notebook: having the possibility to set a breakpoint can be a pleasure for my nerves, specially if you have longer programs. Some of my longer Jupyter Notebooks feel from the hundreds line of code onwards more like pain than like anything helpful. Using an IDE makes it easier for you to split a program into several subprograms.

In an IDE I also appreciate having a “help window” or a “variable explorer”, which is smoothly integrated into the IDE user interface. And there are a lot more advantages why getting familiar with an IDE is a big advantage compared to the very popular Jupyter Notebooks!

I am very surprised, that everyone is talking about Jupyter Notebooks but IDEs are only mentioned very seldom. But maybe my preferences are also a bit different, because I grew up in a MS-DOS environment. :-)

2.3 Docker

For more advanced Data Science techniques, like the Big Data approaches on Apache Spark⁸ and Hadoop⁹, you will need Docker¹⁰. Download it from “Docker - Get Started”¹¹ and also create an account on the Docker website. The installation is easy.

What is Docker? Docker is “*an open-source project that automates the deployment of software applications inside containers by providing an additional layer of abstraction and automation of OS-level*

⁵ GW-Basic, <https://de.wikipedia.org/wiki/GW-BASIC>

⁶ Turbo-Pascal, https://de.wikipedia.org/wiki/Turbo_Pascal

⁷ MS-DOS, <https://de.wikipedia.org/wiki/MS-DOS>

⁸ Apache Spark, <https://spark.apache.org/>

⁹ Hadoop, <https://hadoop.apache.org/>

¹⁰ Docker, <https://www.docker.com/>

¹¹ Docker Get Started, <https://www.docker.com/get-started>

virtualization on Linux.”

After having installed it, you are able to pull my “machine-learning-pyspark” image to your computer and run my Jupyter-Notebooks. I will explain in section sec. 4 “Machine Learning with Python” how it works.

Please read also the Docker-Curriculum¹² for more information. It is a very well structured and nice tutorial which I can recommend for learning about docker images, docker containers and more.

¹² Docker-Curriculum, <https://docker-curriculum.com/>

3 Visualization of Different Datasets

In the data scientist environment the visualization is as important as the analysis itself and the work of a “Data Scientist” with data covers many areas. The Cross-Industry Standard Process for Data Mining (CRISP-Cycle)¹, which “Data Scientists” like to quote and follow in their work, describes what these are (see fig. 3.1).

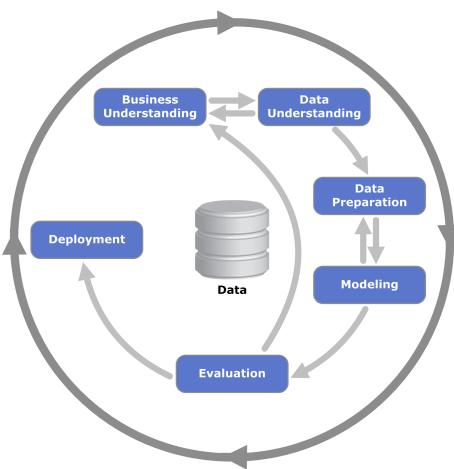


Figure 3.1: Cross-Industry Standard Process for Data Mining - CRISP-Cycle, Wikipedia, CC BY-SA 3.0

Here, the visualization of data at any point in the cycle plays a role:

Visualization helps to get a **better understanding of the data** (“Data understanding”), but also to get a **better understanding of the business model** (“Business understanding”). Visualization **supports data preparation** (“Data preparation”), for example, when it comes to identifying gaps or outliers or to get an indication by means of distribution diagrams which models might be suitable to solve a specific problem (e.g. classification, clustering, regression). Visualization helps in **modeling** (“modeling”) to illustrate whether the models are useful. Visualization is also particularly important in **validation** (“evaluation”) to show how meaningful the applied models are compared to the real world.

So you can see: Visualization techniques are useful at every development step (CRISP cycle). Often visualizations are as important for a “Data Scientist” as the analysis itself. Without a good visualization it is often difficult for a data scientist to communicate the result of his analysis to an outsider. This is

¹ Cross-Industry Standard Process for Data Mining (CRISP-Cycle), Wikipedia, CC-BY SA 3.0, https://en.wikipedia.org/wiki/Cross-industry_standard_process_for_data_mining

also due to the fact that the human eye can recognize patterns in visualizations very quickly, but on the other hand has difficulty understanding a table with numbers at a glance. In individual cases, a good visualization can even have the effect that a complicated model can be completely dispensed with and one only deals with patterns. There are numerous examples of this, which I will not go into here.

I worked on different datasets with the aim to visualize the data and in the **first part** of this chapter (see sec. 3.1) I will explain these examples. In the **second part** (see sec. 3.2) I will explain some Big Data visualizations problems and techniques to solve these. In the **third part** (see sec. 3.3) I will show you how to visualize and share the data with a “data app”. In the **fourth part** sec. 3.4 I will list some common professional tools, which offer visualization functionality and more.

3.1 Examples

In this chapter I work on different examples with the aim to visualize the data with e.g. line plots, histograms, bar charts. These examples consider different topics (consumer price, music, pedestrians,...), different format of the data (csv-file, API²) and cover various problems, that arise, when you work with data (e.g. converting, encoding formats,...). I used python and libraries like e.g Matplotlib or Seaborn, which are available for free. My examples are:

1. A public dataset of the “Consumer Price Index” from the official statistics website of the “Bayrisches Landesamt für Statistik”³.
2. A dataset of my own songs, which I listened to (66'955 songs since 2016) and downloaded from LastFM⁴.
3. An artificially treated dataset of “Marathon run-times”, where I showed how systematics in the data can be found by visualizing.
4. The number of pedestrians in inner cities when the Corona-exit-lock had been implemented.
5. The data from the Deutsche Bahn API to monitor status of their station elevators. Working with an API is a bit different, but interesting.
6. A very brief introduction into visualization of Big Data as this can be challenging with standard visualizing techniques.

The examples are available:

- as .py files for being used for example in an Spyder-IDE⁵ and
- as .ipynb files, which are Jupyter-Notebooks⁶.

² Application Programming Interface (API), Wikipedia, <https://en.wikipedia.org/wiki/API>

³ Bayerisches Landesamt für Statistik, www.statistikdaten.bayern.de

⁴ LastFM, www.last.fm

⁵ Spyder-IDE, <https://www.spyder-ide.org/>

⁶ Jupyter-Notebook, <https://jupyter.org/>

I explained in sec. 2.2, why I think, that working with an IDE is very important and has a lot of advantages compared to the Jupyter-Notebooks, which you find in nearly every Data Science book.

3.1.1 Consumer Price Index Example

In this example I will convert dates and visualize the data using the seaborn regression plot. First download the consumer prices⁷. The CSV-file has the following format (the columns are separated by semicolons):

	A	B	C	D	E
1	GENESIS-Tabelle: 61111-202z				
2	Verbraucherpreisindex (2015=100): Bayern, Verbraucherpreise,				
3	Monate, Jahre				
4	Verbraucherpreisindex				
5	Bayern				
6		Verbraucherpreisindex			
7		2015=100			
8	1970	Januar	29,6		
9	1970	Februar	29,7		
10	1970	März	29,8		
11	1970	April	29,8		
12	1970	Mai	29,9		
13	1970	Juni	30,0		
14	1970	Juli	30,0		
15	1970	August	30,1		
16	1970	September	30,1		

Figure 3.2: Consumer Price Index

Delete the first 7 lines (which are only some metadata) and save the file as 61111-202z-bearbeitet.csv (this is the easiest way to do it, but you can also program this task in Python). The new file has the following format:

```
1970 Januar 29,6
1970 Februar 29,7
1970 März 29,8
1970 April 29,8
1970 Mai 29,9
1970 Juni 30,0
1970 Juli 30,0
```

The preliminary work here is to convert the months (e.g. “Januar”, “Februar”) to a date. Therefore I import some libraries:

⁷ Consumerprice of “Bayerisches Landesamt für Statistik”, <https://www.statistikdaten.bayern.de/genesis/online?sequenz=statistikTabellen&selectionname=61111>. Code: 61111-202z

```
import pandas as pd
from matplotlib import pyplot, dates
from matplotlib.ticker import FuncFormatter
import seaborn as sns
from time import strftime
import locale
locale.setlocale(locale.LC_ALL, '')
```

I lived in Switzerland and bought my computer there. Unfortunately my keyboard and also my language setting is 'German_Switzerland.1252' and therefore importing the CSV file with the `pandas.read_csv` would not work, because in Switzerland and Germany have different separators (comma versus point): 29,6 (German) should become 29.6 (Switzerland). This conversion can be done immediately in the `pandas.read_csv` with converters, but first let's have a look into the importing without using a converter:

```
df=pd.read_csv('61111-202z-bearbeitet.csv',
               encoding="latin-1",
               names=['year', 'month', 'value'])
df.head()
```

As a consequence the result when importing directly with `pandas.read_csv` would be as follows, which is not what we want:

	year	month	value
0	1970;Januar;29	6.0	NaN
1	1970;Februar;29	7.0	NaN
2	1970;März;29	8.0	NaN
3	1970;April;29	8.0	NaN
4	1970;Mai;29	9.0	NaN

We can see, that the `pandas.read_csv` did not use the semicolon as a separator for the columns, but instead comma (which we saw in the number 29,6 and 29,7 and so on). It is a good exercise to learn how to convert different language settings, because you might have the same issue, when downloading datafiles from US and using it in your country.

We have to define converters, which I named `myMonthConverter` (converts "Januar" to 1, "Februar" to 2, ...) and `myValueConverter` (converts 29,6 to 29.6) as follows:

```
def myMonthConverter(s):
    return strftime(s, '%B').tm_mon

def myValueConverter(s):
    return s.replace(',', '.')

def fake_dates(x, pos):
    """ Custom formater to turn floats into e.g., 2016-05-08"""
    return dates.num2date(x).strftime('%Y-%m-%d')
```

Now read the csv file again, with defining the semicolon as separator and using Encoding=latin-1. It also uses the converter functions myMonthConverter and the myValueConverter from above.

```
df=pd.read_csv('61111-202z-bearbeitet.csv',
               sep=";",
               encoding="latin-1",
               names=['year', 'month', 'value'],
               converters={'month':myMonthConverter,
                           'value': myValueConverter})
df.head()
```

The result is as follows:

	year	month	value
0	1970	1	29.6
1	1970	2	29.7
2	1970	3	29.8
3	1970	4	29.8
4	1970	5	29.9

This looks much better. Now have a look at the column value. The type of the column value is not a number, but an object. Type df.dtypes to know which kind of data type we have:

```
year      int64
month     int64
value     object
dtype: object
```

Therefore I need to convert the column value from object into a Pandas numeric :

```
df['value'] = df['value'].apply(pd.to_numeric, errors='coerce')
```

Now we are done with reading the csv and converting the columns. Of course you can convert any kind of value by defining an appropriate function. Imagine, that you have for example a column, which contains the banking ratings and you want to convert them to numbers (e.g. AAA=1, AA=2, A=3 and so on). We know that there are different rating companies (Moodys&Fitch, Standard&Poors), with slightly different rating tables. Aligning both to a numeric rating scale is needed to work with the data. This kind of exercise is now easy, when you use a myconverter function as above.

As a next step I created new columns datenum and date, which I need for the graphics (the regression plot). The column datenum is a step, which I needed because of the Seaborn regression plot function sns.regplot. First the real date (e.g. "1970-01-01") need to be converted to a number (e.g. 719163.0), which is used in the x-Axis of the plot. Then the description of the x-Axis is transformed from 719163.0 back to the real date in a string-format. This seems to be a bit strange, but according to forums this is how it has to be done.

```
#Create new column 'date' based on 'year' and 'month' and convert to date
#Create new column 'datenum' as float for being used in the plot
df['date'] = df['year'].astype(str) + "-" + df['month'].astype(str) + "-1"
df['datenum'] = dates.datestr2num(df['date'])
df['date'] = df['date'].apply(pd.to_datetime,
                           errors='coerce')
df.dtypes
```

```
year           int64
month          int64
value          float64
date    datetime64[ns]
datenum         float64
dtype: object
```

```
df.head()
```

	year	month	value	date	datenum
0	1970	1	29.6	1970-01-01	719163.0
1	1970	2	29.7	1970-02-01	719194.0
2	1970	3	29.8	1970-03-01	719222.0
3	1970	4	29.8	1970-04-01	719253.0
4	1970	5	29.9	1970-05-01	719283.0

Finally I have everything I need for the regression Plot (`regplot`) and I am ready for the visualizing task. I use “Seaborn” for this and I recommend to have a look into the official Seaborn documentation for learning more about the Seaborn library:

```
#Color settings
sns.set(color_codes=True)

#Plot 'datenum' (=float) and 'value' (=float)
fig, ax = pyplot.subplots()
sns.regplot('datenum',
            'value',
            data=df,
            ax=ax)

#Create the x-axis which is 'datenum' converted to %Y-%m-%d
ax.xaxis.set_major_formatter(FuncFormatter(fake_dates))
ax.tick_params(labelrotation=90)
fig.tight_layout()
```

This is what we wanted: a Seaborn regression plot (`seaborn.regplot`), which required me to convert the x-axis from date to a number:

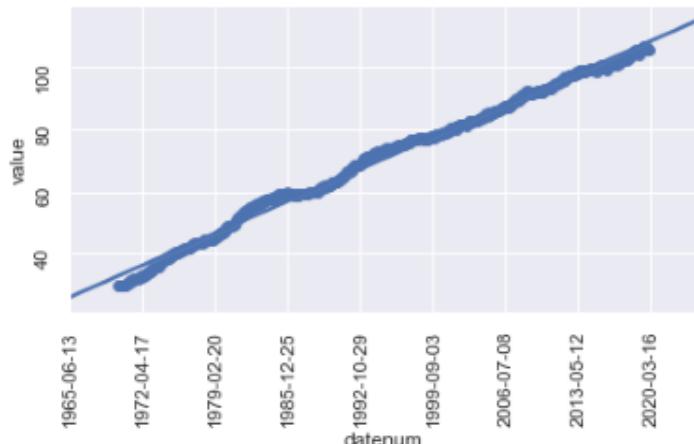


Figure 3.3: Consumer Price Index Figure

I would say, that the fig. 3.3 doesn't show anything noticeable or interesting (apart from a fairly steady rise in consumer prices over the whole period from 1970 until 2020). A bit disappointing so far.

Now let's examine the increments (absolute and relative):

```
import matplotlib.ticker as mtick
from pandas.plotting import register_matplotlib_converters
register_matplotlib_converters()

#Examine increments (absolute and relative)
df['increment_abs'] = df['value']-df['value'].shift(+1)
df['increment_rel'] = (1-df['value'].shift(+1))
                    /df['value'])*100
```

Replace the “NaN” values:

```
#Replace "NaN" by 0
df['increment_abs'].fillna(0, inplace=True)
df['increment_rel'].fillna(0, inplace=True)
```

And create a plot:

```
figin, axin = pyplot.subplots(2)
axin[0].plot(df['date'], df['increment_abs'])
axin[1].plot(df['date'], df['increment_rel'])

#Range of axis
```

```
axin[0].set_ylim([-1.1, +1.1])
axin[1].set_ylim([-1.1, +1.7])

#Title of axis
axin[0].set_title('absolute increment')
axin[1].set_title('relative increment')
for ax in figin.get_axes():
    ax.label_outer()

#Format y axis in percent
axin[1].yaxis.set_major_formatter(mtick.PercentFormatter(decimals=0))
```

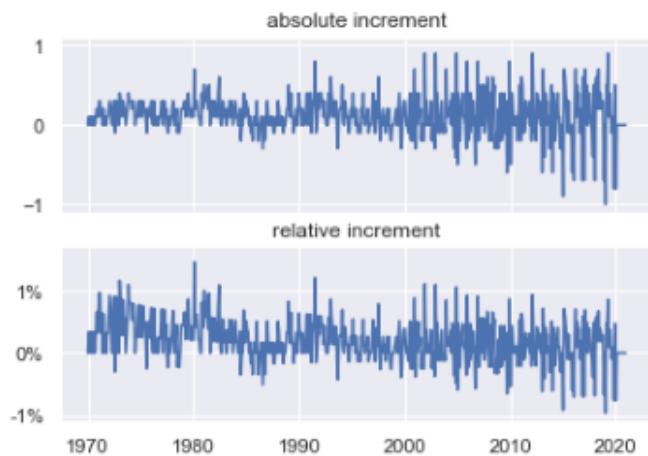
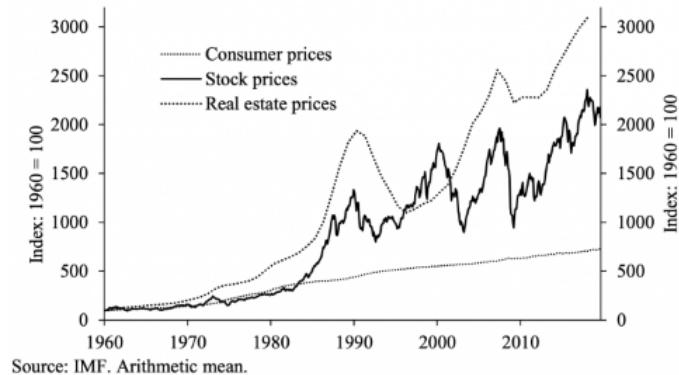


Figure 3.4: Consumer Price Figure - absolute and relative increments

The fig. 3.4 shows the increments (absolute and relative increments) and now we can see something interesting: perhaps one can say that the consumer prices grew more evenly between 1970 and 1995 and that the growth was almost entirely positive (abs/rel increments is above zero). On the other hand, the changes between 1995 and 2020 were somewhat more volatile and increases (positive changes of consumer prices) alternated with decreases (negative changes).

As this was interesting to me I tried to find an explanation or some evidence if we really could split up the whole period (from 1970 until 2020) in one going from 1970 until 1995 and in another one going from 1995 until 2020. It was funny for me, when I found the image fig. 3.5, which shows the consumer prices and stock rates. Wouldn't you say, that the stock prices were also more volatile between 1995 and 2020? Even more interesting: the volatility of the stock prices increased already in 1990 (five years ahead of the consumer prices).

Figure 7: Consumer, Stock and Real Estate Prices in US, Germany and Japan**Figure 3.5:** Consumer price index versus Stockprices, Real estate prices

Let's have a look into the tail of the data:

```
df.tail(15)
```

	year	month	value	date	datenum	increment_abs	increment_rel
597	2019	10	106.6	2019-10-01	737333.0	0.1	0.093809
598	2019	11	105.8	2019-11-01	737364.0	-0.8	-0.756144
599	2019	12	106.3	2019-12-01	737394.0	0.5	0.470367
600	2020	1	105.5	2020-01-01	737425.0	-0.8	-0.758294
601	2020	2	NaN	2020-02-01	737456.0	0.0	0.000000
602	2020	3	NaN	2020-03-01	737485.0	0.0	0.000000
603	2020	4	NaN	2020-04-01	737516.0	0.0	0.000000
604	2020	5	NaN	2020-05-01	737546.0	0.0	0.000000
605	2020	6	NaN	2020-06-01	737577.0	0.0	0.000000
606	2020	7	NaN	2020-07-01	737607.0	0.0	0.000000
607	2020	8	NaN	2020-08-01	737638.0	0.0	0.000000
608	2020	9	NaN	2020-09-01	737669.0	0.0	0.000000
609	2020	10	NaN	2020-10-01	737699.0	0.0	0.000000
610	2020	11	NaN	2020-11-01	737730.0	0.0	0.000000
611	2020	12	NaN	2020-12-01	737760.0	0.0	0.000000

There are some “NaN” values and I will give you right now a short glance into how machine-learning works come back to this later in sec. 4 again. I want to predict the values for the “NaN”, which we have from February 2020 on. In order to do this I have to eliminate some columns and split the whole dataset df_num in one which has the value, which I want to predict df_value and one, which contains the remaining columns df_prepared:

```
df_num = df.dropna(subset=["value"])
df_value = df_num["value"]
df_prepared = df_num.drop(["datenum", "date", "value",
                           "increment_abs", "increment_rel"],
                           axis = 1)
```

I am ready for fitting the LinearRegression:

```
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(df_prepared, df_value)
```

I can now use predict for example for the last 20 entries:

```
some_data = df_prepared.iloc[-20:]
some_values = df_value.iloc[-20:]
df_some_predictions = lin_reg.predict(some_data)
df_some_predictions
```



```
array([105.7954812 , 105.91550781, 106.03553442, 106.15556103, 106.27558764,
       ↵ 106.39561425, 106.51564086, 106.70133425, 106.82136086, 106.94138747,
       ↵ 107.06141408, 107.18144069, 107.3014673 , 107.42149391, 107.54152052,
       ↵ 107.66154713, 107.78157374, 107.90160035, 108.02162696, 108.20732036])
```

My linear regression has found the values for the last 20 entries. I can also do the same for the whole dataset as follows:

```
df_predictions = lin_reg.predict(df_prepared)
print("Predictions:", list(df_predictions))
```

The “mean squared error” is 1.775 can be calculated as follows:

```
from sklearn.metrics import mean_squared_error
import numpy as np
lin_mse = mean_squared_error(df_value, df_predictions)
lin_rmse = np.sqrt(lin_mse)
lin_rmse
```

Obviously I could have done this also in Excel, but as I am now in the Python framework, I can apply more tools on the data, which I will do in a next step. For example: as it seems that there is a connection between stock prices and consumer prices wouldn't it be nice to analyze if more "variables" (like the stock prices) could be found? And wouldn't it be interesting to create some sort of "predicting tool", which calculates the consumer prices index for me for a future date (remember, that the volatility of stock prices increased years before the consumer price index did, so the stock price could perhaps be a "predicting variable" for the consumer price index)? We already know, that there are some nice Python packages for doing this. This would be a task for a next step.

On my GitHub-profile you can download my Jupyter-Notebook⁸.

3.1.2 Last-FM Statistics of my Songs

I am listening quiet a lot to music, either with my app on my mobile phone or my home sound-system. Since 2016 I am using Last-FM⁹ to upload my music statistics (so called "scrobbling"). Last-FM is creating nice graphics for 2019 as shown in fig. 3.6. You can see that in 2019 I listened to 12,052 songs in total, which is 33 songs (=scrobbles) per day. The bar charts in the middle splits this up to a monthly view. A listening clock ("Höruhr"), shows when I was mainly listeing during the day. Not surprisingly the main part is in the evening around 18:00.

In this example I download my complete history of played songs since 2016 from Last-FM (66'955 songs in total) and re-built some of these nice statistics and figures, which last.fm provides. This are for example a bar-plot with monthly aggregates of total played songs. Or top 10 songs of the week and so on. Having the same plots at the end as last.fm has proves, that my results are correct. :-)

⁸ Consumer-Prices Jupyter-Notebook, <https://github.com/AndreasTraut/Visualization-of-Data-with-Python/blob/main/ConsumerPricesExample/ConsumerPrices.ipynb>

⁹ LastFM, www.last.fm

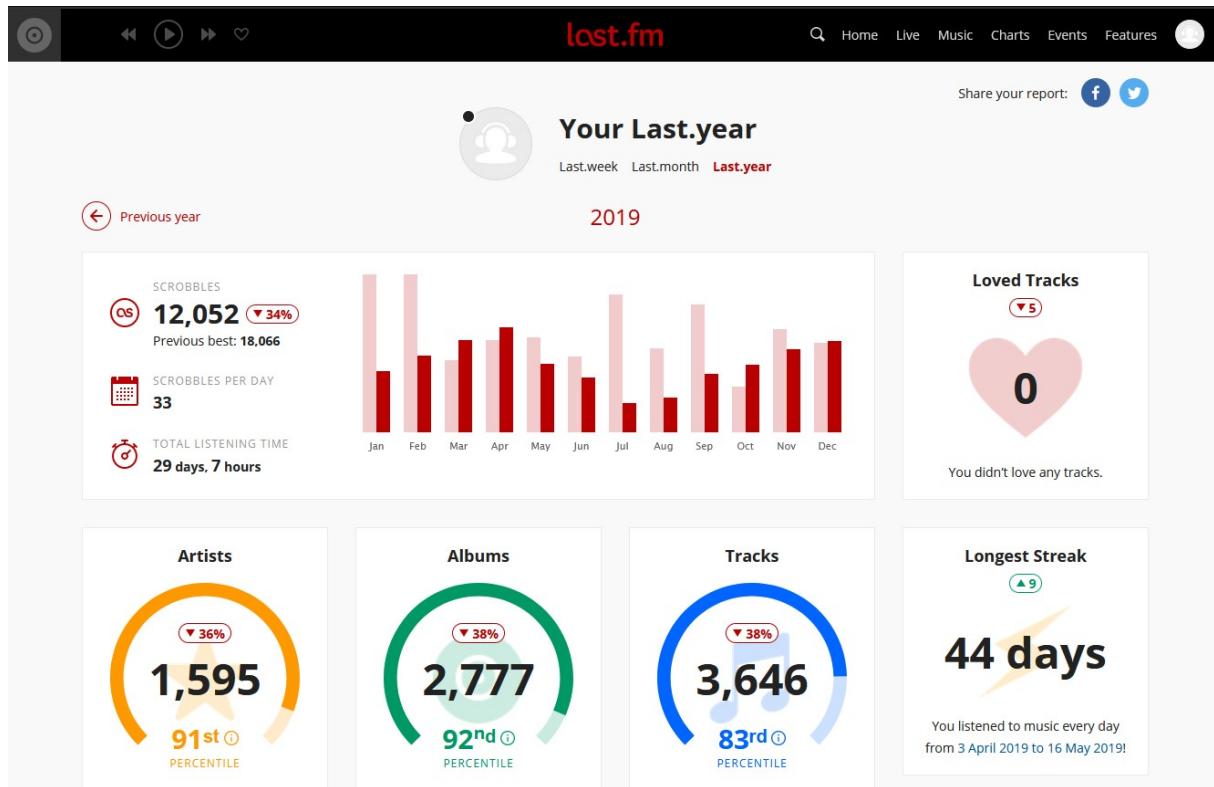


Figure 3.6: Last-FM Music Statistics - Overview year 2019

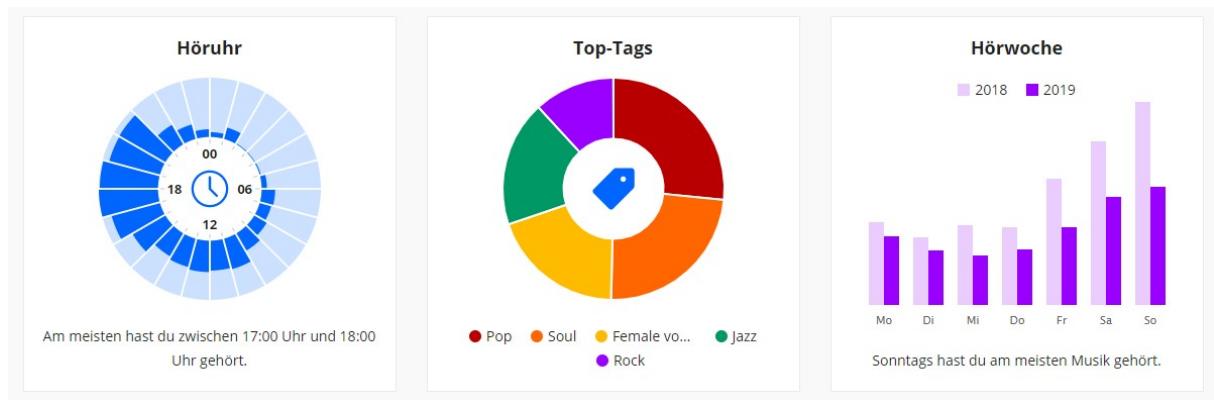


Figure 3.7: Last-FM Music Statistic - Listening clock 2019

The CSV file had the following shape:

A	B	C	D
1 Daniel Santacruz		Lento	06.02.2020 16:45
2 Mau y Ricky	Para Aventuras y Curiosidades	Mi Mala	06.02.2020 16:27
3 Nelson Freitas	Elevate	Something Good	06.02.2020 16:23
4 Jennifer Dias	Love U	Love U	06.02.2020 16:22
5 Nelson Freitas	Sempre Verão	Every Day All Day	06.02.2020 16:18
6 Daniel Santacruz	Lento	Lento	06.02.2020 16:16
7 Mogli	Wanderer (Expedition Happiness Soundtrack)	Road Holes	05.02.2020 15:49
8 Serena Ryder	Harmony (Deluxe)	For You	04.02.2020 17:36
9 Y'akoto	Perfect Timing	Perfect Timing	04.02.2020 17:32
10 Awa Ly	FIVE AND A FEATHER	LET ME LOVE YOU	04.02.2020 17:28

Figure 3.8: Last-FM Music Statistics - Format of the Database

Obviously the columns are ‘artist’, ‘album’, ‘song’, ‘timestamp’. First I wanted to reproduce the overall statistics, which is (as you can see from the screenshot above) 12’052 songs in total for 2019 and 33 songs per day.

```
import pandas as pd
import numpy as np
from matplotlib import pyplot
df = pd.read_csv('lastfm_data.csv',
                  names=['artist', 'album', 'song', 'timestamp'],
                  converters={'timestamp':pd.to_datetime})
```

First I extracted the year / month / date / weeofyear / hour / weekday from the timestamp:

```
#% Extracting year/month/... from timestamp and adding as new columns
dates = pd.DatetimeIndex(df['timestamp'])
df['year'] = dates.year
df['month'] = dates.month
df['weekofyear'] = dates.weekofyear
df['hour'] = dates.hour
df['weekday'] = dates.weekday #Monday=0
```

Next I wanted to have the overall statistics as for example “played songs per year” or “scrobbels per day”.

```
#% Overall statistics
print("\nPlayed songs per year:\n{}\n".format(df['year'].value_counts(sort=False)))
print("\nScrobbels per day:\n{}\n".format(df['year'].value_counts(sort=False)/365.))
```

This is what I found:

```
2018: 18'066 songs in total and 49.495890 songs-per-day.  
2017: 12'521 songs in total and 34.304110 songs-per-day.  
2016: 10'726 songs in total and 29.386301 songs-per day.
```

The screenshot shows a Jupyter Notebook interface with the title "jupyter lastfm Last Checkpoint: vor 22 Minuten (unsaved changes)". The menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, Help, Trusted, and Python 3. Below the menu is a toolbar with various icons. The code cell (In [56]) contains the following Python code:

```
%% Overall statistics
print("\nPlayed songs per year:\n{}".format(df['year'].value_counts(sort=False)))
print("\nScrobbles per day:\n{}".format(df['year'].value_counts(sort=False)/365.))
```

The output of the code is displayed below the cell:

```
Played songs per year:  
2015    12691  
2016    10726  
2017    12521  
2018    18066  
2019    12052  
2020     963  
Name: year, dtype: int64  
  
Scrobbles per day:  
2015    34.769863  
2016    29.386301  
2017    34.304110  
2018    49.495890  
2019    33.019178  
2020     2.638356  
Name: year, dtype: float64
```

Figure 3.9: Last-FM Music Statistics - Overall statistics

These are exactly the same numbers, as Last-FM showed me. So everything is fine so far. Now I even know that the accurate number is 33.019 songs per day! For the year 2018 I calculated 49.495890.

Now lets examine the “top artist”, “top album”, “top songs”:

```
print("\nTop artists:\n{}"
      .format(df['artist'].value_counts().head()))
print("\nTop album:\n{}"
      .format(df['album'].value_counts().head()))
print("\nTop songs:\n{}"
      .format(df['song'].value_counts().head(10)))
```

```
Top artists:  
Aretha Franklin      1284  
Caro Emerald          925  
Paloma Faith          895  
Dionne Bromfield     739  
Nikki Yanofsky        678  
Name: artist, dtype: int64
```

Top album:

Soul Queen	576
Good for the Soul	508
Do You Want the Truth or Something Beautiful?	451
Greatest Hits	405
Emerald Island EP	394

Name: album, dtype: int64

Top songs:

Without You	185
He's So Fine	158
Good for the Soul	152
It's A Beautiful Day	149
Hallelujah	149
This Guy's In Love With You	135
White Christmas	126
Cheek to Cheek	123
Yeah Right	117
Tangled Up	115

Name: song, dtype: int64

Next, I want to define a year / month / weekofyear and see some more detailed statistics:

```
#% Defining a year / month / weekofyear for examination
myYear = 2018
myMonth = 5
myWeekofYear = 21

#% Examine selected year
print("\nAll songs in year %s:\n"%(myYear),
      df.loc[df['year'] == myYear,
              ['artist', 'album', 'song']])
selection = df.loc[df['year'] == myYear,
                   ['artist', 'album', 'song', 'month']]
selectionPrev = df.loc[df['year'] == myYear-1,
                      ['artist', 'album', 'song', 'month']]
print("\nTop artists:\n{}"
      .format(selection['artist'].value_counts().head()))
print("\nTop songs:\n{}"
      .format(selection['song'].value_counts().head(10)))
```

As a next step I wanted to reproduce the bar chart (monthly aggregates of songs):

```
index = np.arange(12)
pltperMonth = pyplot.bar(index, perMonth, width=0.3,
                         label=myYear, color='red')
pltperMonthPrev = pyplot.bar(index - 0.3, perMonthPrev,
                             width=0.3, label=myYear-1,
                             color='peachpuff')
pyplot.title('Year {}. Scrobbels per month.'.format(myYear))
pyplot.xticks(index, ('Jan', 'Feb', 'Mär', 'Apr', 'Mai',
                      'Jun', 'Jul', 'Aug', 'Sep', 'Okt',
                      'Nov', 'Dez'))
pyplot.legend()
pyplot.show()
```

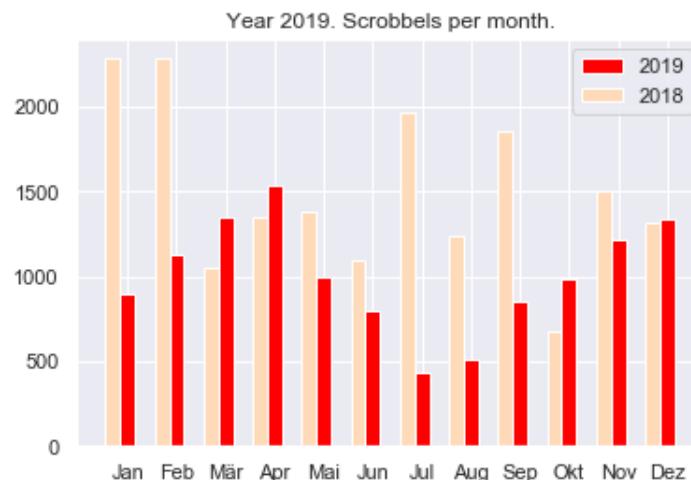


Figure 3.10: Last-FM Music Statistics - Reproduced Statistics of 2019

Nice: it looks also the same, but I can customize mine as I want. For example: I always missed the y-Axis in the Last-FM Chart, which I have now.

Last-FM graphics for 2018:



Figure 3.11: Last-FM Music Statistics - Overview year 2018

My graphics for 2018:

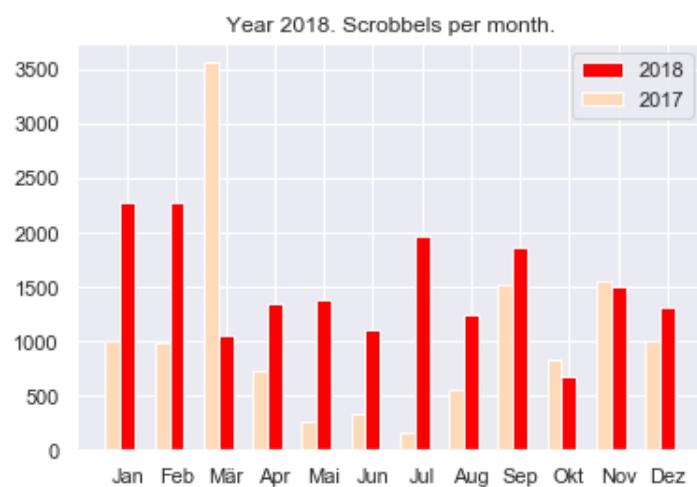


Figure 3.12: Last-FM Music Statistics - Reproduced Statistics of 2018

As you can see, Last-FM shows 49 songs per day for 2018. Remember, that I recalculated 49.495890 (as you can see in the Screenshot above), based on 365 days per year (when I take 365.25 days per year in order to reflect the leap years, I get 49.462). Applying the rounding rules both is rounded to 49 (not 50!). So Last-FM is correct.

Last-FM graphics for 2017:

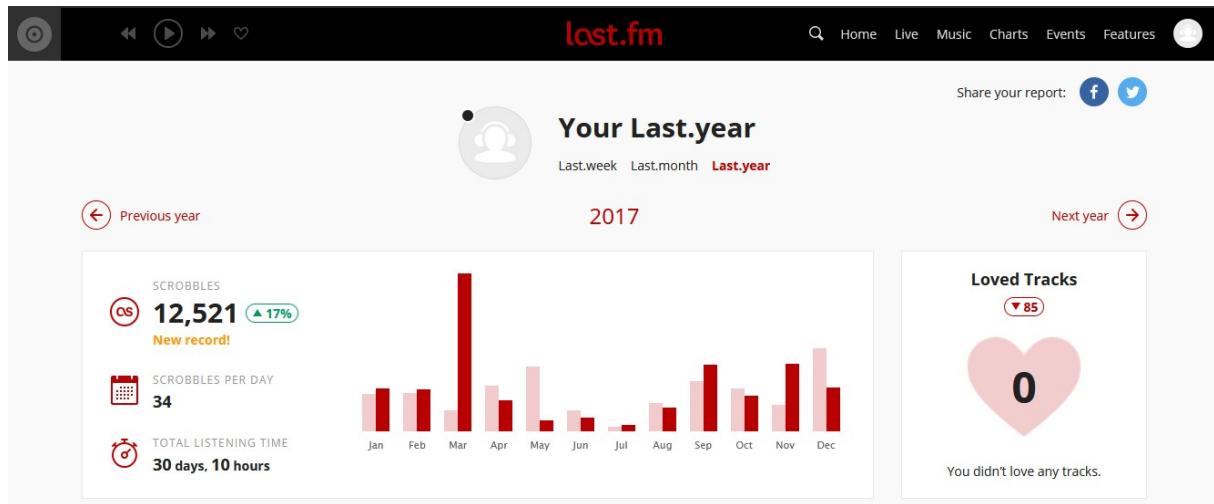


Figure 3.13: Last-FM Music Statistics - Overview year 2017

My graphics for 2017:

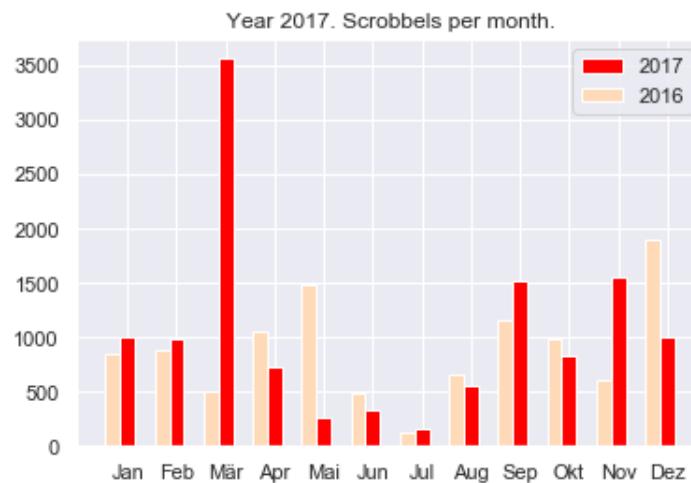


Figure 3.14: Last-FM Music Statistics - Reproduced Statistics of 2017

And as a last exercise I want to create a “Listening Clock”, which is a barplot showing the hours and number of songs I listened to. Obviously I listened mostly in the evening:

```
#%>% Listening Clock
isweekofyear = (df['weekofyear'] == myWeekofYear)
selection = df.loc[isyear & isweekofyear, ['hour']]
index = np.arange(24)
perHour = myselection['hour'].value_counts().sort_index()
pltperHour = pyplot.subplot(111)
```

```

pltperHour.bar(perHour.index,
               perHour,
               width=0.3,
               color='blue',
               alpha=0.5)
pltperHour.set_xticks(index)
pyplot.title('Year {}, Week-of-Year {}. Scrobbels per hour.'
             .format(myYear, myWeekofYear))
pyplot.show()

```

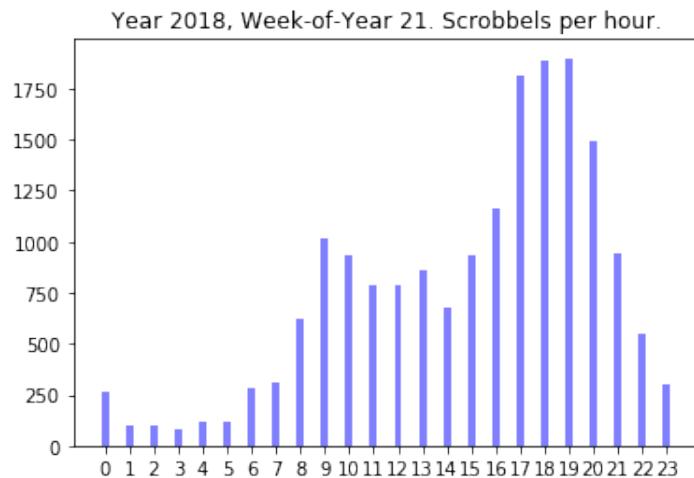


Figure 3.15: Last-FM Music Statistics - Listening clock 2018 (week 21)

On my GitHub profile you can find my Jupyter-Notebook for this example¹⁰

3.1.3 Marathon Runtimes: Finding Systematics

This example shows how different visualization techniques in Python (by using the libraries seaborn and matplotlib) can be used to find out whether there are dependencies, systematics or relationships in a dataset.

Imagine that you receive the following csv data record of 37'250 lines (it is an artificially treated dataset and only for exercise purposes). They show the age, gender and times of marathon runs as well as their nationality, size and weight. Are there any hidden relationships in the data records?

¹⁰ Last-FM Jupyter-Notebook, <https://github.com/AndreasTraut/Visualization-of-Data-with-Python/blob/main/LastFME/xample/lastfm.ipynb>

```
age, gender, split, final, nationality, size, weight
33, M, 01:05:38 02:08:51, DE, 183.41, 84.0
32, M, 01:06:26 02:09:28, DE, 178.61, 87.7
31, M, 01:06:49 02:10:42, IT, 171.94, 82.2
38, M, 01:06:16 02:13:45, IT, 172.29, 82.4
31, M, 01:06:32 02:13:59, IT, 178.59, 79.8
....
```

Download the csv file from my repository and examine the data. At the first glance you won't find anything unusual, but using the following visualization techniques in Python will lead to some conclusions. First we need to import the csv (see step 1) and convert the columns, which contain a time (hh:mm:ss) to seconds (see step 2).

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# %% 1 Read the data. The function "convert" will split the Data after ":" 
def convert(s):
    h, m, s = map(int, s.split(':'))
    return pd.Timedelta(hours=h, minutes=m, seconds=s)

data=pd.read_csv('marathon-data_extended.csv',
                 converters={'split':convert, 'final':convert})
print(data.dtypes)

# %% 2 Apply the converter "convert" to transform the hh:mm:ss.
data['split_sec'] = data['split'] / np.timedelta64(1, 's')
data['final_sec'] = data['final'] / np.timedelta64(1, 's')
```

Since we already suspect that there are connections in certain variables, we form corresponding quotients (see step 3) as e.g. "size to weight" quotient.

```
# %% 3 Add more colums.
data['size_to_weight'] = data['size'] / data['weight']
print(data.head())
```

We receive the following dataset:

```

age, gender, split, final, nationality, size, weight,
↳ split_sec, final_sec, size_to_weight
33, M, 01:05:38 02:08:51, DE, 183.41, 84.0,
↳ 3938.0, 7731.0, 2.183452
32, M, 01:06:26 02:09:28, DE, 178.61, 87.7, 3986.0,
↳ 7768.0, 2.036602
31, M, 01:06:49 02:10:42, IT, 171.94, 82.2, 4009.0,
↳ 7842.0, 2.091727
...

```

Next we will use a jointplot from the seaborn module (`sns.jointplot`, see step 4) and see the following:

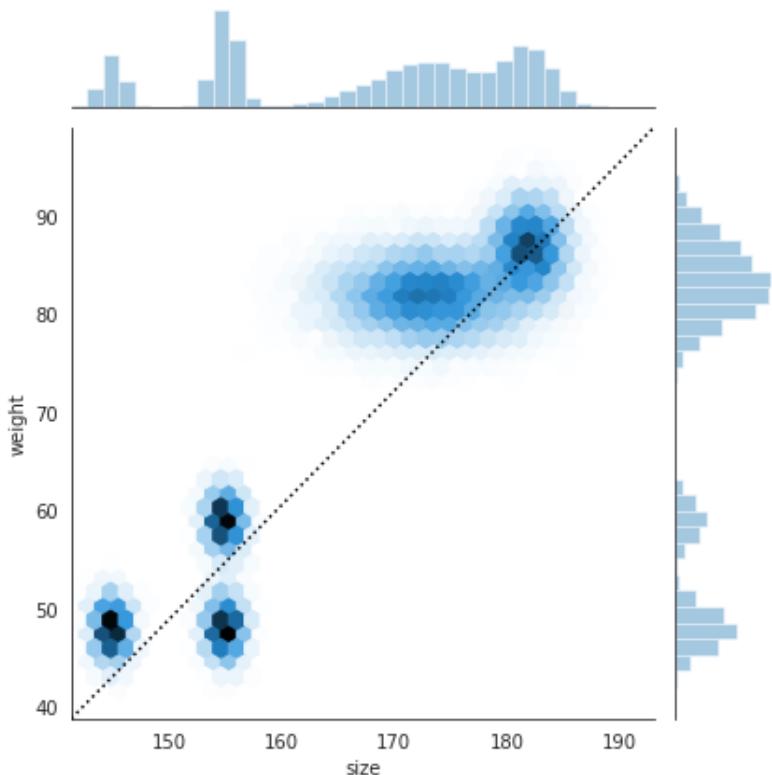


Figure 3.16: Marathon Example - Seaborn Jointplot

```

# %% 4 Joint-Plot with x=size and y=weight
with sns.axes_style('white'):
    g = sns.jointplot("size",
                      "weight",
                      data,
                      kind='hex')

```

```
g.ax_joint.plot(np.linspace(min(data['size']),
                             max(data['size'])),
                 np.linspace(min(data['weight']),
                             max(data['weight'])),
                 ':k')
```

Obviously there are some dependencies in the data records. So we will dig a bit deeper and use the `sns.distplot` (Histogram, see step 5) which will show the following:

```
#%>% 5 Histogram for 'size' and 'weight' (distplot=Distribution Plot)
sns.distplot(data['size'], kde=False);
plt.show()
sns.distplot(data['weight'], kde=False)
```

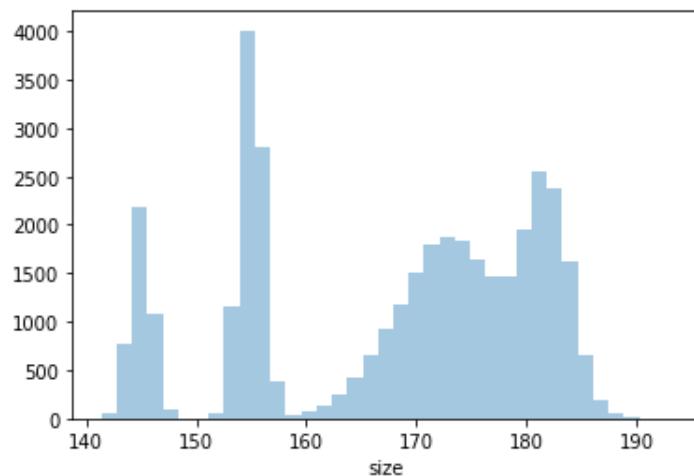


Figure 3.17: Marathon Example - Seaborn Histogram “size”

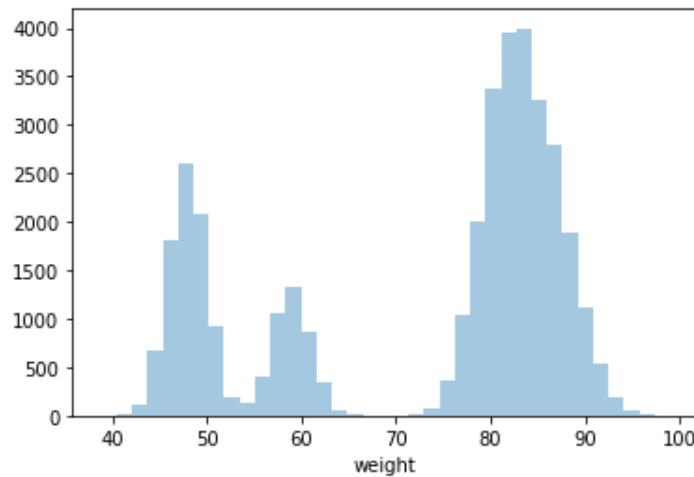


Figure 3.18: Marathon Example - Seaborn Histogram “weight”

As a next step we use the `sns.PairGrid` for examining if there are any correlations between the variables “nationality”, “size”, “final_sec” and “weight” (see step 6):

```
#%>% 6 PairGrids with variables 'nationality', 'size', 'final_sec', 'weight'  
#<-- colors for gender (hue) is GreenBlue (GnBu)  
g = sns.PairGrid(data,  
                  vars=['nationality', 'size', 'final_sec', 'weight'],  
                  hue='gender',  
                  palette='GnBu_r')  
g.map(plt.scatter, alpha=0.8)  
g.add_legend();
```

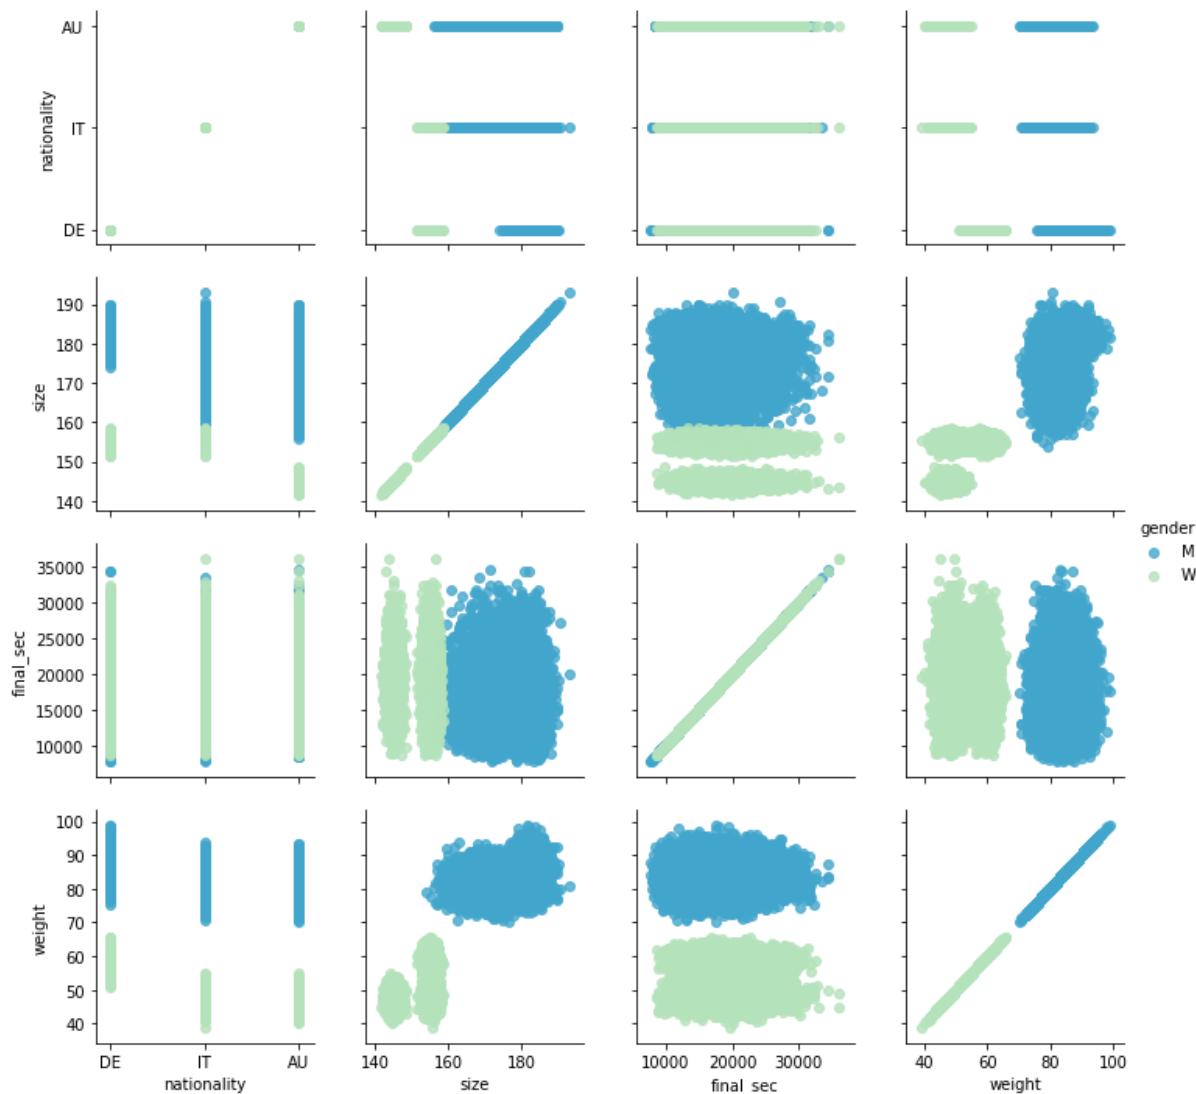


Figure 3.19: Marathon Example - Seaborn PairGrid

This visualization already reveals a lot of information: German people have a higher weight (women as well as men). Austrian women are the smallest people and so on. Let's use the Kernel density functions next for the variable "size to weight" (step 7) and "size" (step 8):

```
#%>% 7 KernelDensity (kde) for column "size_to_weight"
sns.kdeplot(data.size_to_weight[data.nationality=='DE'],
             label='Deutschland',
             shade=True)
sns.kdeplot(data.size_to_weight[data.nationality=='AU'],
             label='Österreich',
             shade=True)
```

```
plt.xlabel('size_to_weight');
plt.show()
```

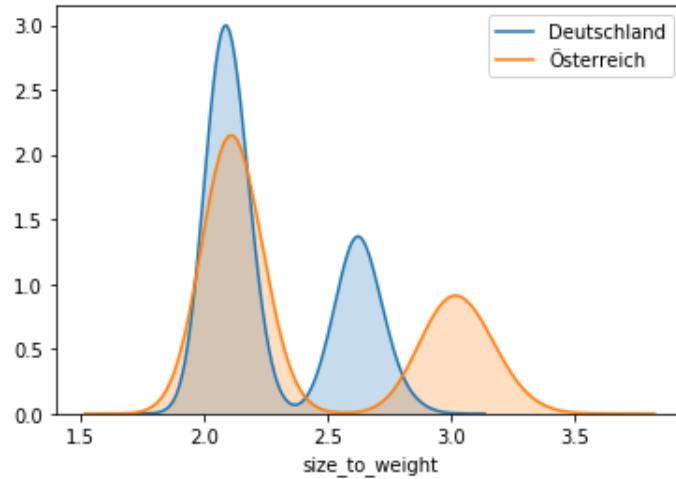


Figure 3.20: Marathon Example - Seaborn Kernel Density “size-to-weight”

```
#%>% 8 KernelDensity (kde) for column "size"
sns.kdeplot(data.weight[data.nationality=='DE'],
             label='Deutschland',
             shade=True)
sns.kdeplot(data.weight[data.nationality=='AU'],
             label='Österreich',
             shade=True)
plt.xlabel('size');
```

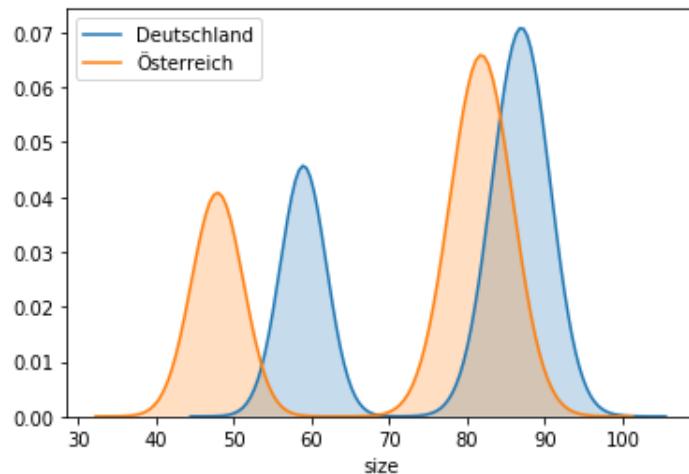


Figure 3.21: Marathon Example - Seaborn Kernel Density “size”

Now it would interesting to see some regression plots (step 9). In fig. 3.22 for men and in fig. 3.23 for women:

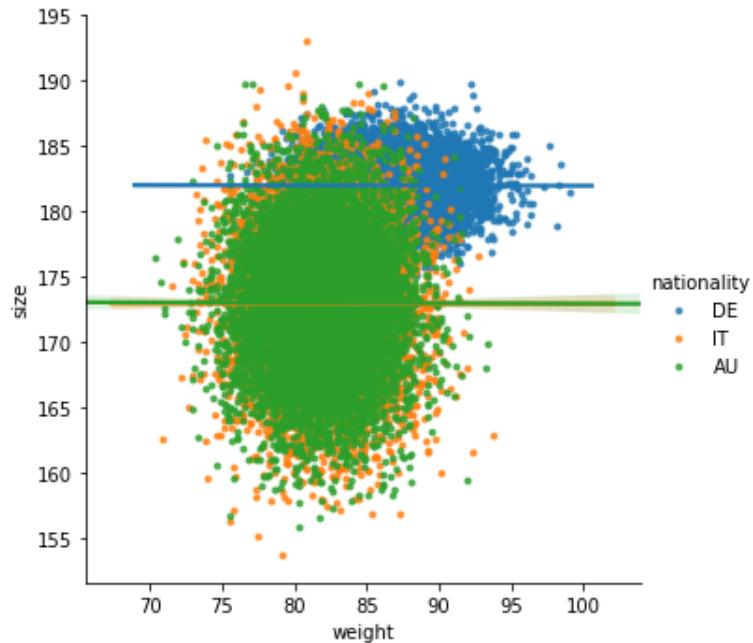


Figure 3.22: Marathon Example - Seaborn Regression Plots “men”

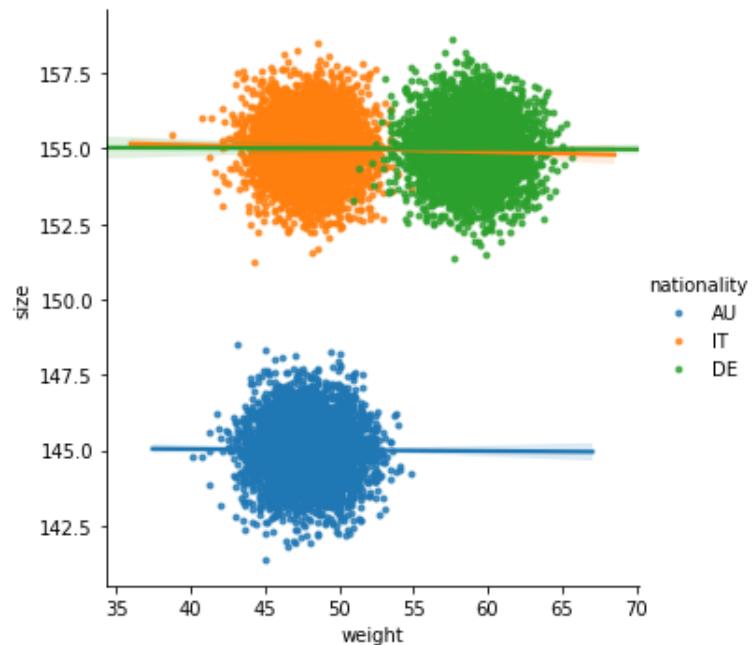


Figure 3.23: Marathon Example - Seaborn Regression Plots “women”

```
#%% 9 Regression Plot for "weight" and "size"
h = sns.lmplot('weight', 'size', hue='nationality',
                data=data[data.gender=="M"],
                markers=".")
h = sns.lmplot('weight', 'size', hue='nationality',
                data=data[data.gender=="W"],
                markers=".")
```

Here again we see, that Austrian women are smaller (see dots in blue). And finally we will use the Seaborn-Violinplots, `sns.violinplot` (step 10 and 11), which finally reveals all details, which have been hidden in this dataset:

```
#%% 10 Violinplot using "size" and "nationality"
men = (data.gender == 'M')
women = (data.gender == 'W')
with sns.axes_style(style=None):
    sns.violinplot("size", "nationality",
                    hue="gender",
                    data=data,
                    split=True,
                    inner="quartile",
                    palette=["lightblue", "lightpink"]);
plt.show()
```

For example: we see in fig. 3.24, that German men are taller (180 cm) with a more narrow distribution (standard deviation), than Italian and Austrian men. The distribution of Italian men and Austrian men seems to be identical (normal distribution with the same mean, but a bigger standard deviation). In contrast: Italian women are taller (155 cm) than Austrian women (145 cm).

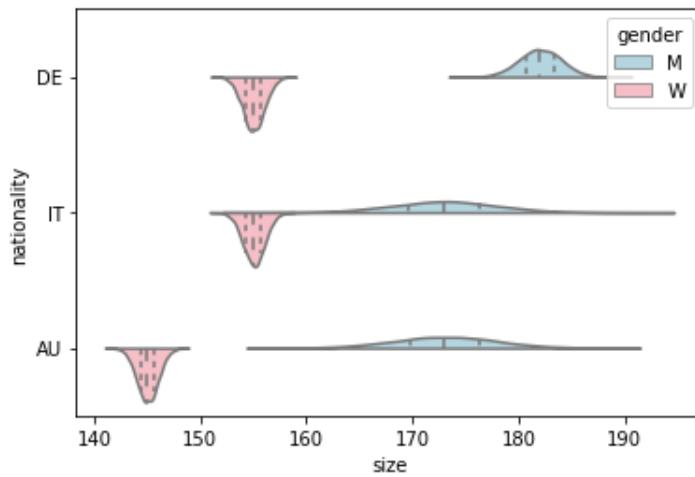


Figure 3.24: Marathon Example - Seaborn Violinplot “size”

```
#%>% 11 Violinplot using "weight" and "nationality"
with sns.axes_style(style=None):
    sns.violinplot("weight",
                   "nationality",
                   hue="gender",
                   data=data,
                   split=True,
                   inner="quartile",
                   palette=["lightblue", "lightpink"]);
```

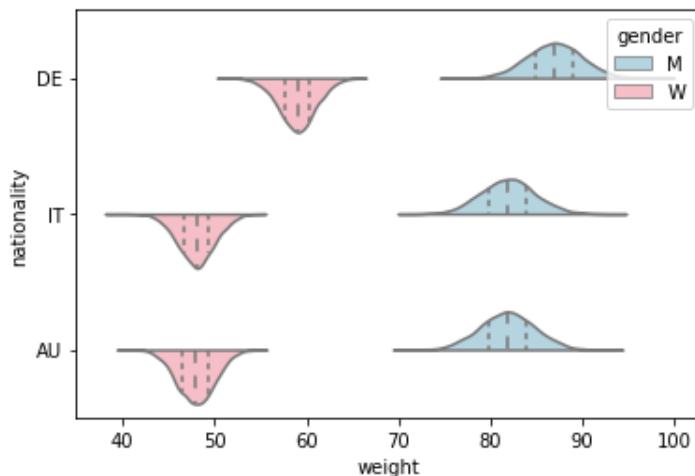


Figure 3.25: Marathon Example - Seaborn Violinplot “weight”

In fig. 3.25 we see, that the weight of Italian women and Austrian women seem (in contrast to their

size) to be identically distributed with about 48 kg in average. German women are heavier with about 58 kg in average. German men are the heaviest (with about 88 kg in average). A deeper examination of the distributions would need some background in mathematics, which we won't do here.

Obviously the underlying data has been treated artificially by me (I apologize for any negative sentiment I might have pushed to Austrian, Italian or German people).

The example above shows, how easy visualization techniques can be and how powerful Python is (combined with the libraries seaborn and matplotlib). Imagine doing the same in Excel: it would take a lot longer. A few lines of code are sufficient for revealing a lot of hidden information of a dataset. Without knowing too much about mathematics or statistics, the systematics in the underlying data are found. The same logic applies to any kind of data your company may hold in their hands (invoices, number of contracts, overtime hours, ...).

Data Scientist often forget, that all their visualizations (and also model), which they have built, need to be used by someone, who is probably not as skilled in all these technical requirements! Therefore it is important to find a solution, which is **easy to deploy** and **easy to use** for everyone (as well on a computer as also on a mobile phone), **stable** and **quickly customizable**. In Section sec. 3.3 I will show you how to share this data with an "Data App".

For the "Marathon runtimes" example I also wrote a testing file and included Travis¹¹ and Codecov¹². Travis and Codecov provide small icons, like the following:



Figure 3.26: Build Status passing

The advantage of doing this is: when I share my python code or Jupyter-Notebooks on code-sharing platforms, like GitHub¹³ other people will know, that my code has been tested and does not contain bugs. If you plan to share your code frequently then I recommend to have a look at least into Travis and Codecov (and there are a lot more services like these two, which might also be interesting).

On my GitHub profile you can find my Python-File for this example¹⁴

¹¹ Travis, <https://travis-ci.com/>

¹² Codecov, <https://codecov.io/>

¹³ GitHub, <https://github.com>

¹⁴ Marathon Python-Files, https://github.com/AndreasTraut/Visualization-of-Data-with-Python/blob/main/Example_Marathon_extended.py

3.1.4 Pedestrians during the first Corona-Lockdown

On Saturday 21.03.2020 exit locks had been implemented in order to protect the people from the Corona virus. I read an article about the company Hystreet¹⁵, who provides statistics about pedestrians in inner cities. Due to the implemented exit lock the number should decrease from 21.03.2020 onwards. Hystreet offers free downloadable csv files for private use. In this example I wanted to see these decreasing numbers. I downloaded the statistics for several week for Ulm, Münsterplatz. The following plot shows the number of pedestrians walking on Ulm, Münsterplatz. From 21.03.2020 onwards the number of people decreased:

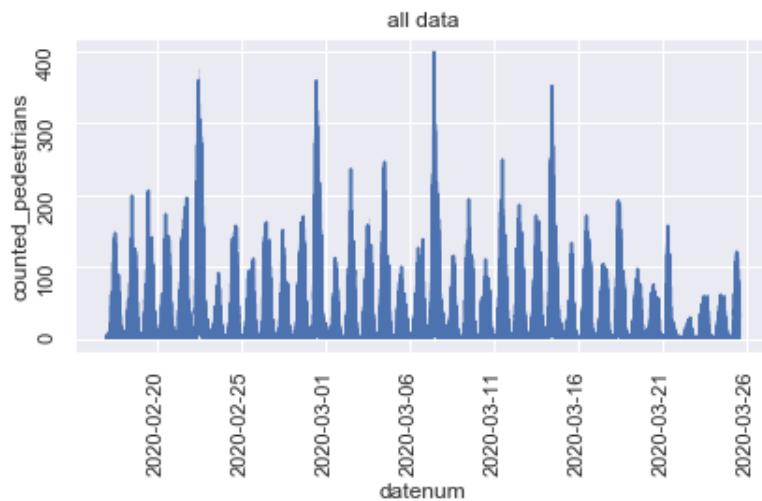


Figure 3.27: Pedestrians in Corona-Lockdown - Frequency

The numbers for Ulm, Münsterplatz in a barplot look like this:

¹⁵ Hystreet, <https://hystreet.com>

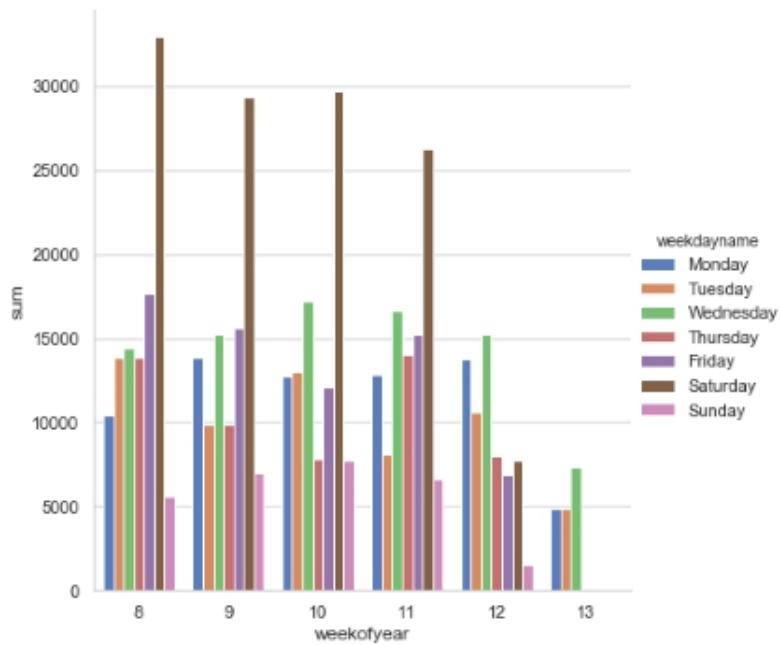


Figure 3.28: Pedestrians in Corona-Lockdown - Barplot “Ulm, Münsterplatz”

And here are some more graphics for München and Augsburg

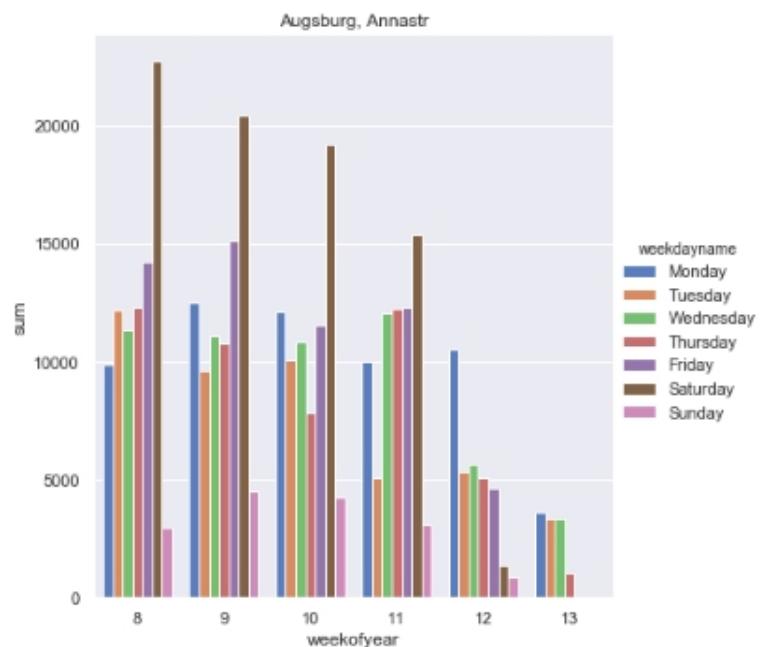


Figure 3.29: Pedestrians in Corona-Lockdown - Barplot “Augsburg”

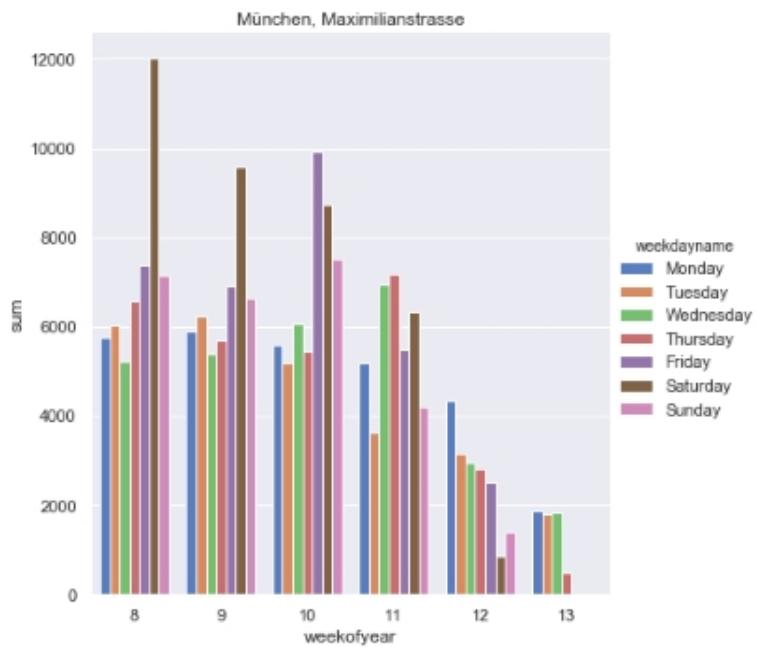


Figure 3.30: Pedestrians in Corona-Lockdown - Barplot “München, Maximilianstrasse”

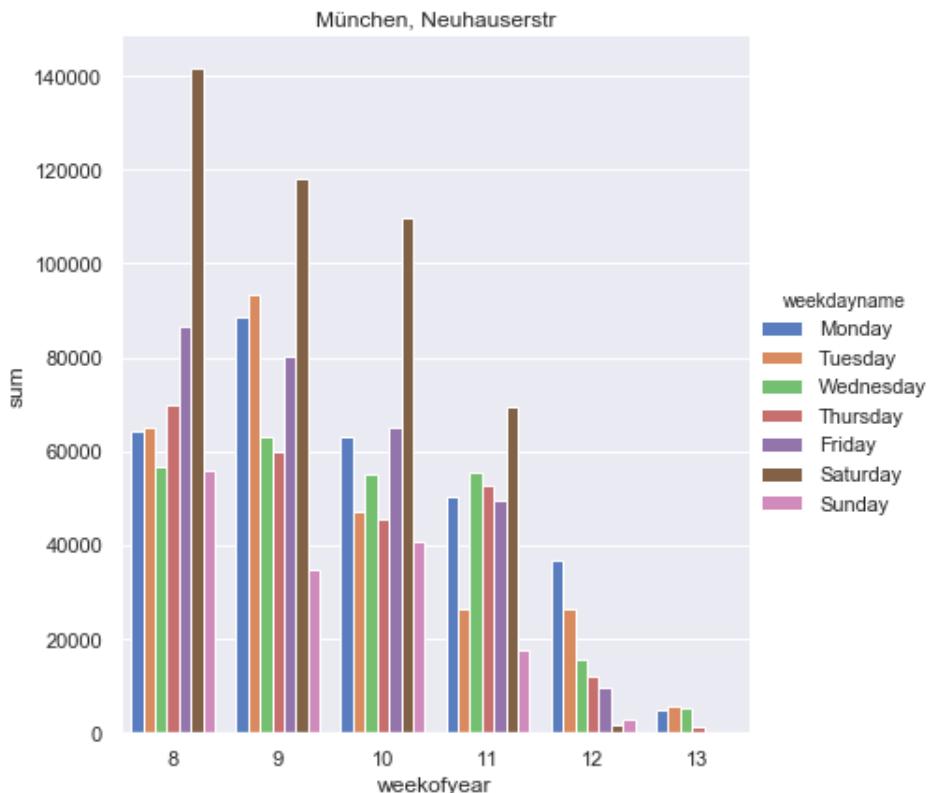


Figure 3.31: Pedestrians in Corona-Lockdown - Barplot “München, Neuhausstrasse”

Just to see, how it looks like for a longer time horizon: from August 2019 to March 2020. Here are the plots for Ulm, Münsterplatz and Munich, Maximilanstr:

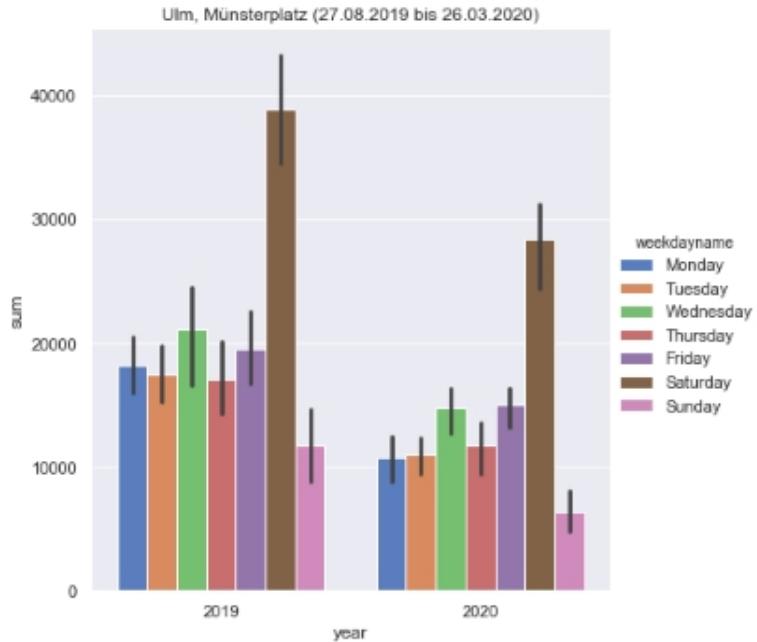


Figure 3.32: Pedestrians in Corona-Lockdown - Barplot “Ulm” 08/2019 to 03/2020

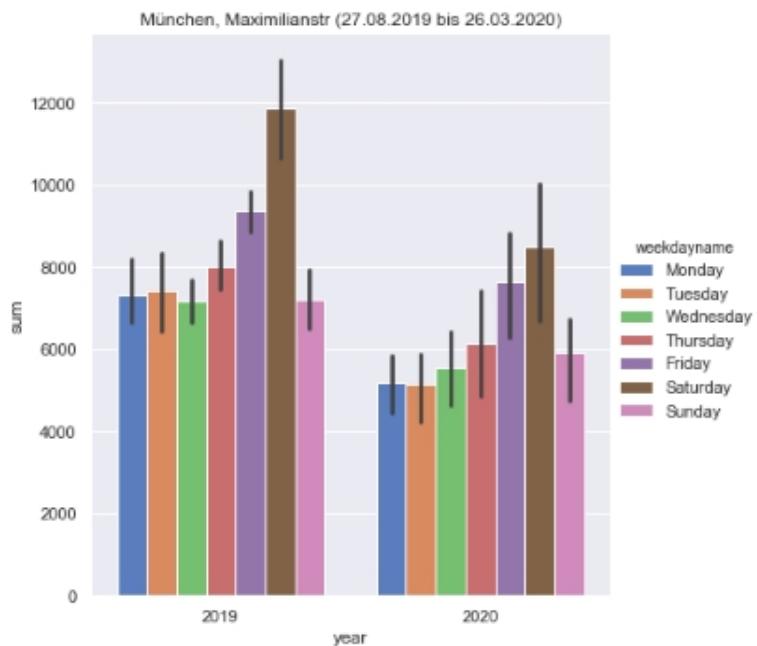


Figure 3.33: Pedestrians in Corona-Lockdown - Barplot “München, Maximilianstrasse” 08/2019 to 03/2020

As there are more datapoints in 2019 (126 from 27.08.2019 to 31.12.2019, compared to 85 from 01.01.2020 to 26.03.2020) the bars are higher in 2019. Here is an update of the above graphics as-of 07.04.2020 for Ulm, Münsterplatz:

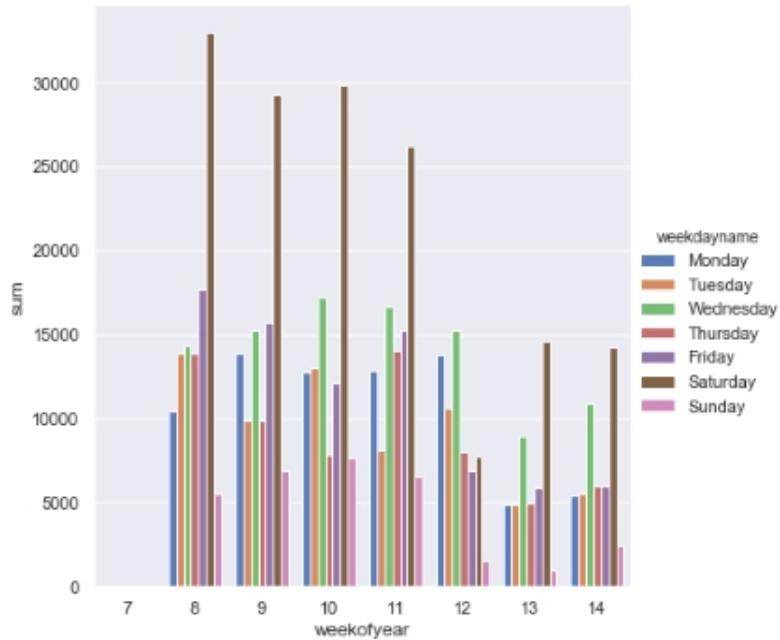


Figure 3.34: Pedestrians in Corona-Lockdown - Barplot “Ulm” April 2020

On my GitHub profile you can find my Python-File for this example¹⁶

3.1.5 Station Elevators of Deutsche Bahn: API Example

The visualization is sometimes a bit difficult, because the dataset is not yet available in the form you need to have them. Until now I worked with csv files, which is the easiest way to use a dataset. If you work with APIs instead of csv files, the data preparation and visualization will be a bit different. An API is an “Application Programming Interface”. It defines the interactions with a software component and the kind of calls or requests, that can be made. An API allows the users to use the interface independently of the implementation¹⁷.

I wanted to know which elevators from Deutsche Bahn are currently working and which ones are damaged. I knew, that I could extract this information from the Deutsche Bahn API “FaSta”¹⁸, but I would

¹⁶ Pedestrians during Corona-Lockdown Jupyter Notebook, <https://github.com/AndreasTraut/Visualization-of-Data-with-Python/blob/main/Pedestrians/Pedestrians.ipynb>

¹⁷ Application Programming Interface (API), HubSpire, <https://www.hubspire.com/resources/general/application-programming-interface/>

¹⁸ Deutsche Bahn API FaSta, https://developer.deutschebahn.com/store/apis/info?name=FaSta-Station_Facilities_Status&version=v2&provider=DBOpenData

need to implement some Python-Code (.py File) to extract the information I needed:

- the “station number” (the station number for Ulm it is 6323)
- “equipment number” (the number of the elevator) and
- the “status” (“available”/“monitoring disrupted”).

I subscribed to the Deutsche Bahn FaSta API and received a consumer key and consumer secret for a production environment and a sandbox environment (=testing environment). I am not allowed to share my keys here, but you can create some yourself, if you want:

API	Eigentümer	Status des Abonnements	Stufe	Aktionen
FaSta-Stat. - v2 API Status : PUBLISHED	DB Station&Service AG	Active	30_per_minute	<input type="button" value="Entfernen"/>

You need to insert your tokens in order to get access to the FaSta-API with the following python code:

```
myToken = 'xxxxxxxxxxxxxxxxxxxxxxxxxxxxx' #Produktion
myUrl = 'https://api.deutschebahn.com/fasta/v2'
head = {'Authorization': 'Bearer {}'.format(myToken)}
response = requests.get(myUrl, headers=head)
```

Each elevator has an equipment number and the four elevators in Ulm these are 10500702, 10500703, 10500704 and 10499292. I didn't find an documentation for these numbers and I found them by trial-and-error. Maybe the Deutsche Bahn didn't want the transparency over these numbers in order to hide the number of damaged elevators a bit.

To get access to the equipment number 10500702, use the following code:

```
myEquipmentnumber = myUrl + "/facilities/10500702"
response = requests.get(myEquipmentnumber, headers=head)
print("Response status code is: ", response.status_code)
result = response.content
```

You will receive the response status code 200 if everything works. Now let's read some more details from this equipment:

```
result_Eq = response.content
data = json.loads(result_Eq)
print("stationnumber: {}".format(data["stationnumber"]))
print("equipmentnumber: {}".format(data["equipmentnumber"]))
print("description: {}".format(data["description"]))
print("state: {}".format(data["stateExplanation"]))
```

```
stationnumber: 6323
equipmentnumber: 10500702
description: BAHNSTEIG 2/3 SÜD
state: available
```

This means, that the elevator with equipment number 10500702 at station number 6323 (which is Ulm) is available. Now I want to have a closer look into all elevators (including their GPS coordinates) at the station Ulm. After having understood and having extracted these meta data (station number, equipment number) I was able to visualize them: as you can see in the screenshot fig. 3.35, when I handed over a station number (e.g. 6323 for Ulm) to the Deutsche Bahn API with and received the number of elevators (here 4) and also the longitudes X=9.98278 and latitudes Y=48.39846 of this elevator.

The screenshot shows the Spyder Python IDE interface. On the left, the code editor displays a script named `apis_DeutscheBahn.py`. The code uses the `requests` library to interact with the Deutsche Bahn API. It retrieves station information for Ulm Hbf (station number 6323) and lists elevators (Aufzüge) available at that station. The right side of the interface shows the variable explorer, a table of variables with their names, types, sizes, and values. Below the variable explorer is the IPython console, which shows the execution of the script and the resulting output. A green bracket highlights the output area in the console.

```

7
8     import requests
9     import json
10
11     # myToken = '448b975c434636db038a8ba92cc4b70' #Sandbox
12     myToken = '448b975c434636db038a8ba92cc4b70' #Produktion
13     myUrl = 'https://api.deutschebahn.com/fasta/v2'
14     head = {'Authorization': 'Bearer {}'.format(myToken)}
15     response = requests.get(myUrl, headers=head)
16
17     #%%
18     myEquipmentnumber = myUrl + "/facilities/10500702"
19     response = requests.get(myEquipmentnumber, headers=head)
20     # print("Response status code is:", response.status_code)
21     result = response.content
22
23     result_Eq = response.content
24     data = json.loads(result_Eq)
25     print(data["stationnumber"])
26     print(data["equipmentnumber"])
27     # print(data["description"])
28     print(data["stateExplanation"])
29
30     #%%
31     myStationnumber= myUrl + "/stations/6323" #Ulm
32     # myStationnumbers= myUrl + "/stations/3629" #Leipheim
33     # myStationnumbers= myUrl + "/stations/6030"
34     response = requests.get(myStationnumber, headers=head)
35     # print("Response status code is:", response.status_code)
36     result = response.content
37
38     result_Stat = response.content
39     data = json.loads(result_Stat)
40     print("Bahnhof Nummer: ", data["stationnumber"], "Name: ", data["name"])
41     print("Anzahl Aufzüge: ", len(data["facilities"]))
42     i=0
43     while i<len(data["facilities"]):
44         t = data["facilities"][i]
45         print("Aufzug Nummer", t["equipmentnumber"], ": ",
46               t["stateExplanation"], " -> "
47               "Koordinaten (X=", t["geocoordX"], ", Y=", t["geocoordY"], ")")
48         i=i+1

```

In [7]: runcell(1, 'C:/Users/andre/Documents/Meine Python Dateien/API (Deutsche Bahn) Beispiel/apis_DeutscheBahn.py')
6323
10500702
available

In [8]: runcell(2, 'C:/Users/andre/Documents/Meine Python Dateien/API (Deutsche Bahn) Beispiel/apis_DeutscheBahn.py')
Bahnhof Nummer: 6323 Name: Ulm Hbf
Anzahl Aufzüge: 4
Aufzug Nummer 10500702 : available -> Koordinaten (X= 9.98278 , Y= 48.39846)
Aufzug Nummer 10500703 : monitoring disrupted -> Koordinaten (X= 9.98254 , Y= 48.39845)
Aufzug Nummer 10500704 : available -> Koordinaten (X= 9.98227 , Y= 48.39844)
Aufzug Nummer 10499292 : available -> Koordinaten (X= 9.98307 , Y= 48.39847)

In [9]:

Figure 3.35: Deutsche Bahn - Elevators 2

```

myStationnumber= myUrl + "/stations/6323" #Ulm
response = requests.get(myStationnumber, headers=head)
result = response.content
result_Stat = response.content
data = json.loads(result_Stat)
print("Bahnhof Nummer: ", data["stationnumber"], "Name: ", data["name"])
print("Anzahl Aufzüge: ", len(data["facilities"]))
i=0
while i<len(data["facilities"]):
    t = data["facilities"][i]
    print("Aufzug Nummer", t["equipmentnumber"], ": ",
          t["stateExplanation"], " -> "
          "Koordinaten (X=", t["geocoordX"], ", Y=", t["geocoordY"], ")")
    i=i+1

```

```

Bahnhof Nummer: 6323 Name: Ulm Hbf
Anzahl Aufzüge: 4
Aufzug Nummer 10500702 : available -> Koordinaten (X= 9.98278 , Y=
    ↵ 48.39846 )
Aufzug Nummer 10500703 : available -> Koordinaten (X= 9.98254 , Y=
    ↵ 48.39845 )

```

```
Aufzug Nummer 10500704 : available -> Koordinaten (X= 9.98227 , Y=
    ↳ 48.39844 )
Aufzug Nummer 10499292 : available -> Koordinaten (X= 9.98307 , Y=
    ↳ 48.39847 )
```

Taking these and using for example GPS-Coordinates¹⁹ you can easily visualize these longitudes and latitudes as shown in fig. 3.37:

The screenshot shows the Spyder Python IDE interface. On the left, the code editor displays a script named `api_DeutscheBahn.py`. The code uses the `requests` library to make API calls to Deutsche Bahn's elevator status service. It prints the equipment number, state, and coordinates for two elevators. In the center, a variable inspection window shows a dictionary named `data` containing the following key-value pairs:

Schlüssel	Typ	Große	Wert
equipmentnumber	int	1	10500702
geocoordX	float	1	9.98278
geocoordY	float	1	48.39846
state	str	1	ACTIVE
stateExplanation	str	1	available
stationnumber	int	1	6323
type	str	1	ELEVATOR

On the right, the IPython console shows the results of the print statements from the script. Arrows point from the highlighted code in the editor to the corresponding data in the variable viewer and the printed output in the console.

```
In [7]: runcell(1, 'C:/Users/andre/Documents/Meine Python Dateien/API (Deutsche Bahn) Beispiel/api_DeutscheBahn.py')
Out[7]: 6323
        10500702
        available
In [8]:
```

Figure 3.36: Deutsche Bahn - Elevators 1

Another possibility would have been to use Geopy²⁰ and type the following code (please verify the ToS for using this service on your own: it's limited!):

```
from geopy.geocoders import Nominatim
geolocator = Nominatim(user_agent="http")
location = geolocator.reverse("48.39846, 9.98278")
print(location.address)
```

The result would have been

¹⁹ GPS-Coordinates, www.gps-coordinates.net

²⁰ Geopy, <https://github.com/geopy/geopy>

"Steig 2 + 3, Bahnhofplatz, Fischerviertel, Weststadt, Ulm, Baden-Württemberg, 89073, Deutschland"

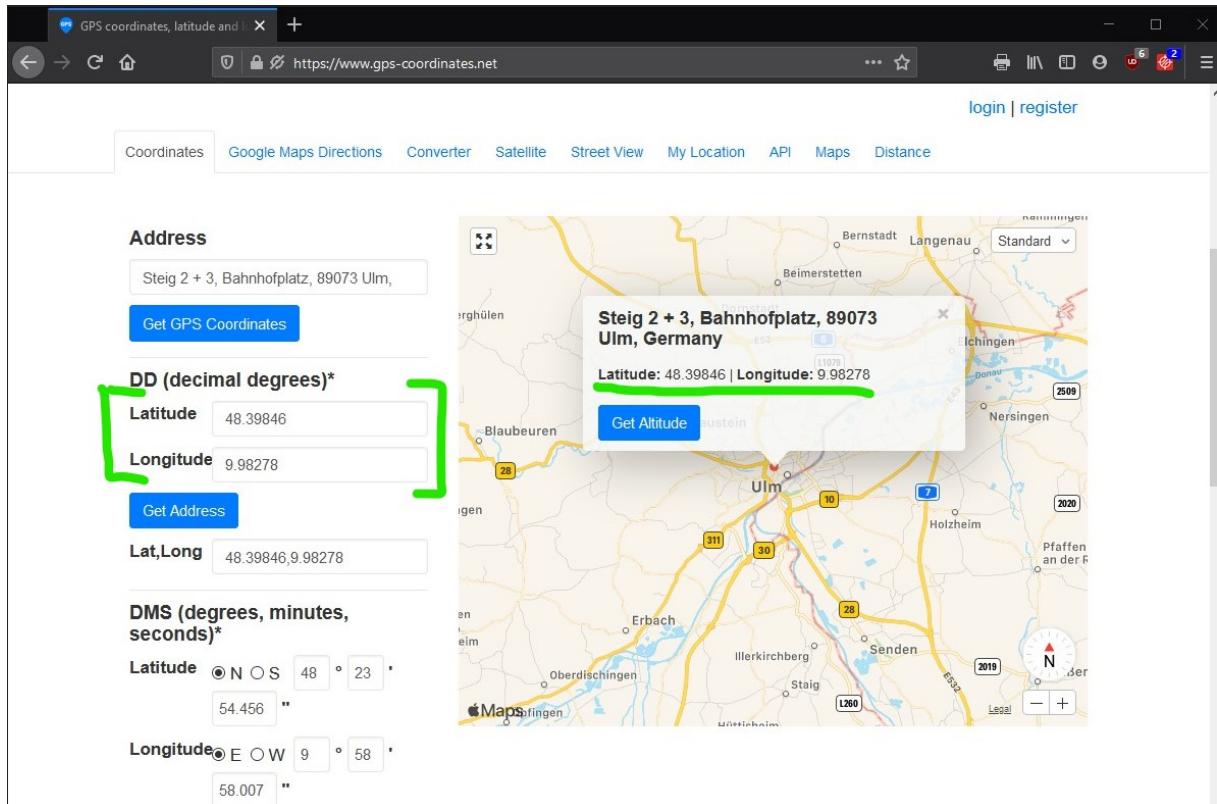


Figure 3.37: Deutsche Bahn - Latitude and Longitude

Building a longer history (not only one extraction of data, but many over a longer period) to show how many elevators are damaged in Germany during one year would also be possible. It would mean, that I have to loop each day over all “station numbers”, then over all “equipment numbers” and store the number of status=damaged. After one year I would have the history of damaged elevators. This is what a journalist did in order to write an article (which was available on www.br.de, but unfortunately disappeared).

3.1.6 Interactive Pivots

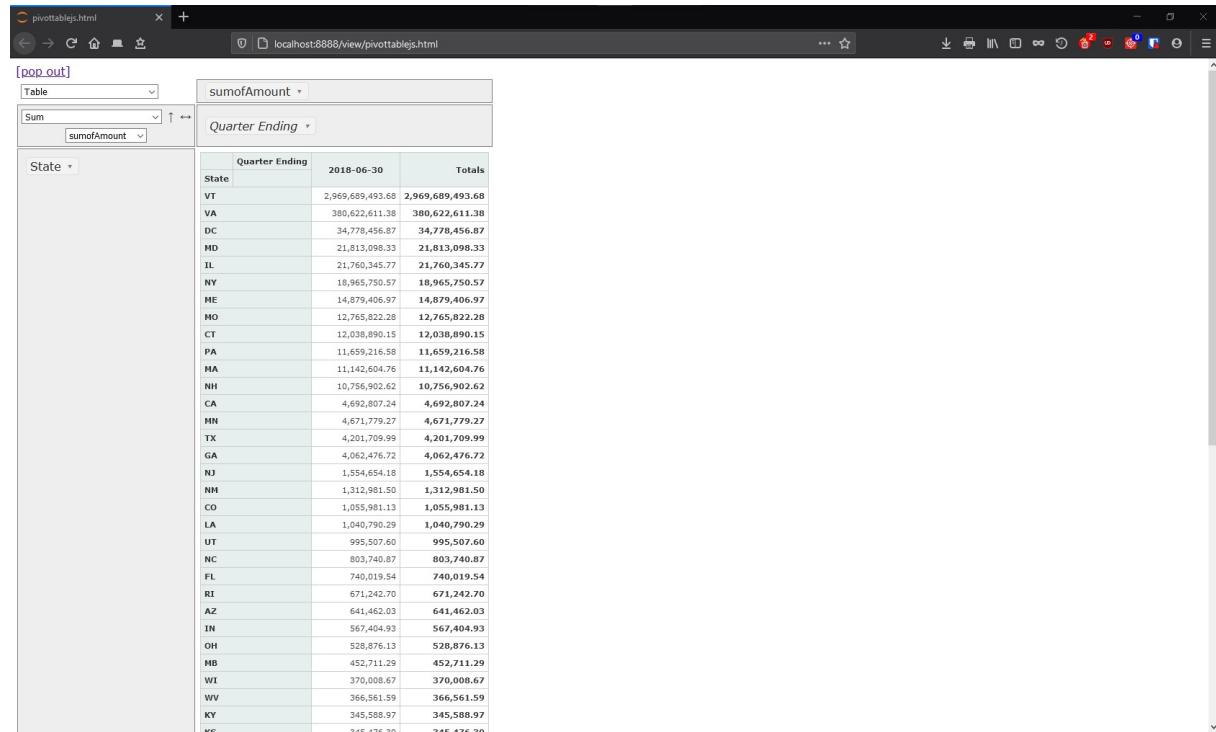
Another nice possibility to visualize pandas data frames is by using the “*pivottablejs*”²¹ implementation. I loved how easy this implementation is to use (basically one line of code) and how nice it creates pivot tables in a Browser (like Firefox, Chrome,...). I think everyone knows from Excel already, what a pivot table is and possibly loves its high intractability: grouping and aggregating is very easy. Data

²¹ Pivottable, <https://pivottable.js.org/examples/>

Scientists often forget, that their models and visualizations need to be used by people, who are not as familiar with the technical concepts as a Data Scientist. Therefore my recommendation for a Data Scientist is to always think about how the data and model should be transferred to the users: “*pivottablejs*” is one solution, which fulfills this criteria. I will describe another solution (“*Streamlit App*”, which is a lot more powerful than “*pivottablejs*”) in sec. 3.3.

```
# Install: !pip install pivottablejs
import pivottablejs
pivottablejs.pivot_ui(pandas_df)
```

Now you can drag and drop the columns and also use the filter select boxes:



The screenshot shows a web browser window with the URL `localhost:8888/view/pivottablejs.html`. The page displays an interactive pivot table. The top navigation bar shows the title "pivottables.html". On the left, there are filter dropdowns for "State" and "Quarter Ending", both currently set to "All". Below these is a "Sum" dropdown set to "sumofAmount". The main area contains a table with the following data:

	Quarter Ending	2018-06-30	Totals
State			
VT		2,969,689,493.68	2,969,689,493.68
VA		380,622,611.38	380,622,611.38
DC		34,778,456.87	34,778,456.87
MD		21,813,098.33	21,813,098.33
IL		21,760,345.77	21,760,345.77
NY		18,965,750.57	18,965,750.57
ME		14,879,406.97	14,879,406.97
MO		12,765,822.28	12,765,822.28
CT		12,038,890.15	12,038,890.15
PA		11,659,216.58	11,659,216.58
MA		11,142,604.76	11,142,604.76
NH		10,756,902.62	10,756,902.62
CA		4,692,807.24	4,692,807.24
MN		4,671,779.27	4,671,779.27
TX		4,201,709.99	4,201,709.99
GA		4,062,476.72	4,062,476.72
NJ		1,554,654.18	1,554,654.18
NM		1,312,981.50	1,312,981.50
CO		1,055,981.13	1,055,981.13
LA		1,040,790.29	1,040,790.29
UT		995,507.60	995,507.60
NC		803,740.87	803,740.87
FL		740,019.54	740,019.54
RI		671,242.70	671,242.70
AZ		641,462.03	641,462.03
IN		567,404.93	567,404.93
OH		528,876.13	528,876.13
MB		452,711.29	452,711.29
WI		370,008.67	370,008.67
WV		366,561.59	366,561.59
KY		345,588.97	345,588.97
VA		946,476.50	946,476.50

Figure 3.38: Interactive Pivots - Grouping and aggregating

Creating graphs is also possible:

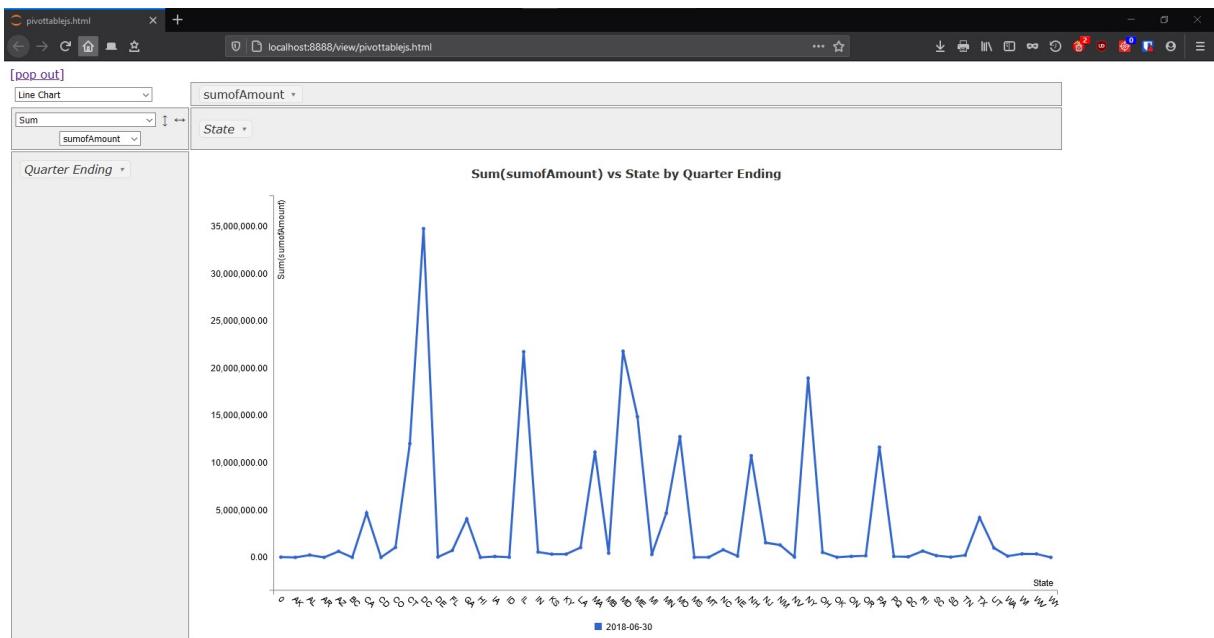


Figure 3.39: Interactive Pivots - Graph

3.2 Introduction to Big Data Visualization Techniques

3.2.1 Basic Problems in Big Data Visualizations

I will provide a short introduction into how Big Data visualization needs to be approached. Visualization of Big Data datasets can be challenging. There are different problems, which will become relevant if you want to do Big Data visualizations:

Can you load all the data into your memory? Probably not. Can you transfer the data from your database over your network (intranet / internet) to your terminal (a computer or mobile phone)? You might slow down the entire network in your company.

What about if you need slide controllers, filters, selections boxes for a user who wants to interact with the data and who wants to produce different visualizations? Then you would create a lot of back and forth transactions between the data base and the terminal. What would you do then?

What about the screen-resolution of your terminal? The screen resolution of a mobile phone is limited to some thousand pixels: visualizing Big Data on a screen with low resolution would end in an ugly black blob on the screen. A computer monitor might show a bit more detail, but can really assume that everyone has a monitor with 4K resolution?

3.2.2 Vermont Payments Example

First I downloaded the list of all state of Vermont payments to vendors ²² (Open Data Commons License), which is a 298 MB huge csv file with 1.6 million lines (exactly 1'648'466 lines). As said: visualization of such huge datasets can be difficult with common tools like Excel. For this example you can use Power-Query²³ or something similar. If you are already familiar with this tool, then you might solve this problem immediately and without additional efforts. But you loose for example the advantages of machine-learning algorithms from the Apache Spark Machine Learning Library²⁴, which is only one big difference to be mentioned.

First let's get ready for the Big Data environment and open my machine-learning Docker container. If you don't know how, then have a look into sec. [4.3.1](#).

```
import os
print("APACHE_SPARK_VERSION: ",
      os.environ["APACHE_SPARK_VERSION"])
print("HADOOP_VERSION: ",
      os.environ["HADOOP_VERSION"])
```

I am on APACHE_SPARK_VERSION: 3.0.0 and HADOOP_VERSION: 3.2. For initializing spark we need a SparkContext and a SparkSession:

```
import pyspark
from pyspark.sql import SparkSession
from pyspark.sql import functions as F
sc = pyspark.SparkContext(appName='Spark Modelling Context')
spark = SparkSession.builder \
    .appName('Spark Modelling Session') \
    .config('spark.executor.memory','5g') \
    .config('spark.executor.cores','4') \
    .getOrCreate()
```

Next step is to read the huge csv file:

```
import os
datapath = os.environ['PWD']
filename = datapath + "/data/Vermont_Vendor_Payments.csv"
```

²² Vermont payments to vendors, <https://data.vermont.gov/Finance/Vermont-Vendor-Payments/786x-sbp3>

²³ PowerQuery, <https://support.microsoft.com/de-de/office/einf%C3%BChrung-in-microsoft-power-query-f%C3%A4%C3%9F-Cr-excel-6e92e2f4-2079-4e1f-bad5-89f6269cd605>

²⁴ Spark Machine Learning Library, <https://spark.apache.org/mllib/>

```
#read in data from csv
data = spark.read.csv(path=filename, sep=',', encoding='utf-8', header=True,
    inferSchema=True)
```

First we want more details about the dataset:

```
data.describe().show(truncate=False, n=1, vertical=True)
```

```
-RECORD 0-----
summary           | count
Quarter Ending   | 1648418
Department        | 1648418
UnitNo            | 1648418
Vendor Number     | 1648418
Vendor             | 1648418
City               | 906137
State              | 1648418
DeptID Description| 1647881
DeptID             | 1648418
Amount             | 1648418
Account            | 1648418
AcctNo             | 1648418
Fund Description   | 1648416
Fund                | 1648417
only showing top 1 row
```

Once you are into the Apache Spark environment you can easily aggregate, sort, group by what ever you want. Take for example the columns “Department” and “Amount” (it should be obvious, what is in these columns, I guess). Then this line of code will show you the sum of column “Amount” grouped per department (sorted descending):

```
data.groupBy('Department') \
    .agg(F.sum('Amount')) \
    .cast('decimal(20,2') \
    .alias('sumofAmount')) \
    .sort('sumofAmount', ascending=False) \
    .show(truncate=False)
```

Department	sumofAmount
------------	-------------

Buildings & Gen Serv-Prop	254664862145.07
Vermont Health Access	7316059819.96
Natural Res Central Office	6115935633.73
Education Agency	5156573496.88
Education	3166972698.64
Transportation Agency	2795155337.06
Department of VT Health Access	2393175142.17
Finance & Management	2331413298.90
Agency of Transportation	1920833560.49
Children and Families	1850543830.45
null	1501137106.31
Office of VT Health Access	1466094349.44
Disabilities Aging Ind. Living	1371142726.00
Children and Family Services	1281368750.70
Mental Health	1259316412.60
Human Resources-Prop	1195825754.22
Human Resources-Gov'tal	1191172666.22
Treasurer's Office	882007501.28
Health	853538376.05
Aging and Independent Living	789393840.12

only showing top 20 rows

The following line of code will group by “Quarter Endings”, aggregate the “Amounts” (sum) and sort the result ascending:

```
spark_df = data.groupBy('Quarter Ending') \
    .agg(F.sum('Amount')) \
    .cast('decimal(20,2)') \
    .alias('sumofAmount')) \
    .sort('Quarter Ending', ascending=True)
spark_df.show(truncate=False)
```

Quarter Ending	sumofAmount
03/31/2010	28598462713.41
03/31/2011	15900537688.65
03/31/2012	17106963001.97
03/31/2013	7799135775.53
03/31/2014	2078065068.29

```

| 03/31/2015 | 3191476205.73 |
| 03/31/2016 | 2135004787.94 |
| 03/31/2017 | 3281966008.12 |
| 03/31/2018 | 4471045973.70 |
| 03/31/2019 | 3294727802.75 |
| 06/30/2010 | 14937551502.74 |
| 06/30/2011 | 11468901690.66 |
| 06/30/2012 | 18502465970.34 |
| 06/30/2013 | 5774960317.50 |
| 06/30/2014 | 4655109691.64 |
| 06/30/2015 | 5824948519.54 |
| 06/30/2016 | 3499716477.98 |
| 06/30/2017 | 4680238788.00 |
| 06/30/2018 | 3552314458.12 |
| 06/30/2019 | 3598682370.54 |
+-----+

```

only showing top 20 rows

If you want to focus on the top 6 states (which are VT, MA, NH, PA, VA, GA), then the following code will filter on these states, aggregate the “Amount” (sum) and sort by “Quarter Ending” (ascending):

```

spark_df = data.filter((F.col('State') == 'VT') |
                      (F.col('State') == 'MA') |
                      (F.col('State') == 'NH') |
                      (F.col('State') == 'PA') |
                      (F.col('State') == 'VA') |
                      (F.col('State') == 'GA')) \
    .groupBy('State', 'Quarter Ending') \
    .agg(F.sum('Amount')) \
    .cast('decimal(20,2)') \
    .alias('sumofAmount')) \
    .sort('Quarter Ending', ascending=True)
spark_df.show()

```

State	Quarter Ending	sumofAmount
VT	03/31/2010	26122164537.74
PA	03/31/2010	5857825.21
NH	03/31/2010	6802959.20
MA	03/31/2010	2328199767.23
GA	03/31/2010	5204426.97

VA	03/31/2010	1392512.09
GA	03/31/2011	6103226.21
VA	03/31/2011	336391.83
NH	03/31/2011	7040302.64
PA	03/31/2011	8241318.26
MA	03/31/2011	1168511063.45
VT	03/31/2011	14564606674.20
PA	03/31/2012	6260265.01
GA	03/31/2012	4402372.46
VT	03/31/2012	14613390813.89
NH	03/31/2012	15779425.80
VA	03/31/2012	392115.68
MA	03/31/2012	2333892646.13
VA	03/31/2013	680498.51
PA	03/31/2013	6297576.34

only showing top 20 rows

You can continue working in Apache Spark with grouping, aggregating and sorting your data until you have a result table, which might be interesting for you. This result can easily be moved into Pandas²⁵ or Excel where creating of bar-plots is very easy:

```
import pandas as pd
import matplotlib.pyplot as plt
pandas_df = spark_df.toPandas()
```

We have now left the “Big Data” environment Spark and are entering the “Pandas”, where creating plots is easier, but only possible on a “small” number of data points. Here we only a small amount of data points remains. For example:

²⁵ Pandas, <https://pandas.pydata.org/>

	Department	sumofAmount
0	Buildings & Gen Serv-Prop	254664862145.07
1	Vermont Health Access	7316059819.96
2	Natural Res Central Office	6115935633.73
3	Education Agency	5156573496.88
4	Education	3166972698.64
5	Transportation Agency	2795155337.06
6	Department of VT Health Access	2393175142.17
7	Finance & Management	2331413298.90
8	Agency of Transportation	1920833560.49
9	Children and Families	1850543830.45

Figure 3.40: Big Data Visualization - Data Format

Similarly you may want to plot a chart of the aggregated “Amount” for a “Quarter Ending”:

```
spark_df = data.groupBy('Quarter Ending') \
    .agg(F.sum('Amount')) \
    .cast('decimal(20,2)') \
    .alias('sumofAmount')) \
    .sort('Quarter Ending', ascending=True)
```

The result is a series with timestamps and the “sumofAmount”, which can be plotted with Pandas very easily.

```
spark_df = data.groupBy('Quarter Ending') \
    .agg(F.sum('Amount')) \
    .cast('decimal(20,2)') \
    .alias('sumofAmount')) \
    .sort('Quarter Ending', ascending=True)
pandas_df = spark_df.toPandas()
pandas_df['Quarter Ending'] = pandas_df['Quarter
    ↵ Ending'].astype('datetime64')
pandas_df['sumofAmount'] = pandas_df['sumofAmount'].astype('float')
pandas_df.plot(x='Quarter Ending',
    y='sumofAmount')
```

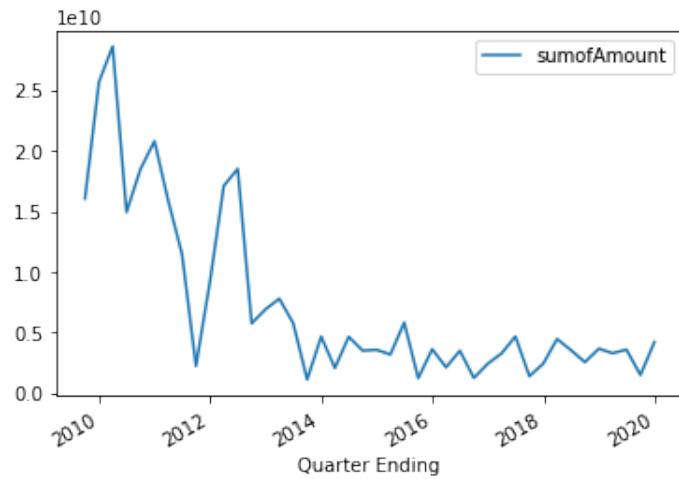


Figure 3.41: Big Data Visualization - Line Chart

```
pandas_df['Quarter Ending'] = pandas_df['Quarter
    ↵ Ending'].astype('datetime64')
pandas_df['sumofAmount'] = pandas_df['sumofAmount'].astype('float')
ax = pandas_df.groupby(['Quarter Ending',
    ↵ 'State']).sum().unstack().plot(kind='bar', stacked=True)
ax.xaxis.set_major_locator(plt.AutoLocator())
plt.show()
```

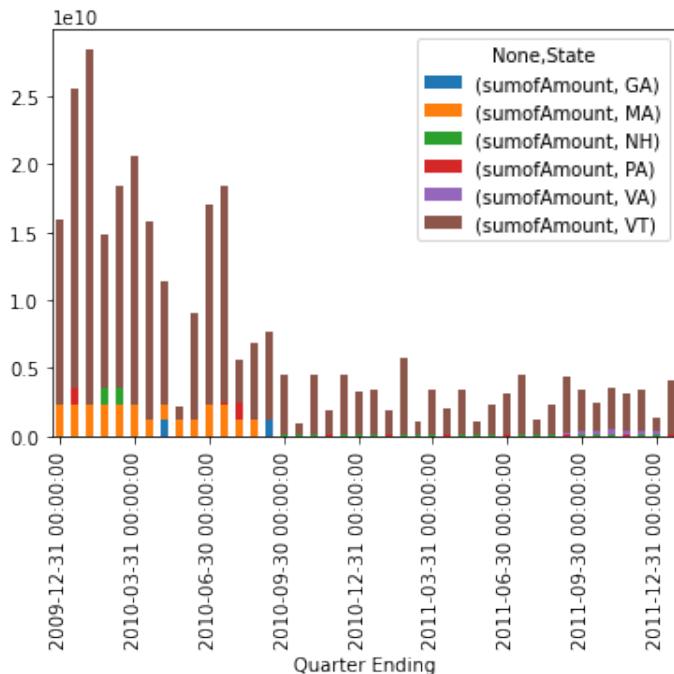


Figure 3.42: Big Data Visualization - Stacked Barplot

I think this example shows, how easy the visualization of Big Data datasets can be done, if you use more advanced tools instead of Excel.

On my Docker machine-learning repository²⁶ you will find the Jupyter-Notebook, with this example.

3.3 Visualize Data with Data Apps

Now, in the second part of this chapter, I will show you how to visualize and share the data with a “data app”. I used Streamlit²⁷, which is surprisingly easy if you want to connect your data with python code directly to a very intuitive and easy to use application.

Data Scientist often forget, that all their visualizations (and also model), which they have built, need to be used by someone, who is probably not as skilled in all these technical requirements! Therefore it is important to find a solution, which is **easy to deploy** and **easy to use** for everyone (as well on a computer as also on a mobile phone), **stable** and **quickly customizable**.

There are different solutions: when you are using the programming language R then the combination

²⁶ My docker machine-learning repository, <https://hub.docker.com/repository/docker/andreasraut/machine-learning-pyspark>

²⁷ Streamlit, <https://www.streamlit.io/>

of Tidyverse²⁸ and Shiny-App²⁹ will be an interesting option for you. But to me the “R / Tidyverse / Shiny” bundle seems a be the “old-standard” or even a bit “old-fashioned” as an article on Data-Revenue³⁰ reveals a strong increasing popularity for Streamlit):

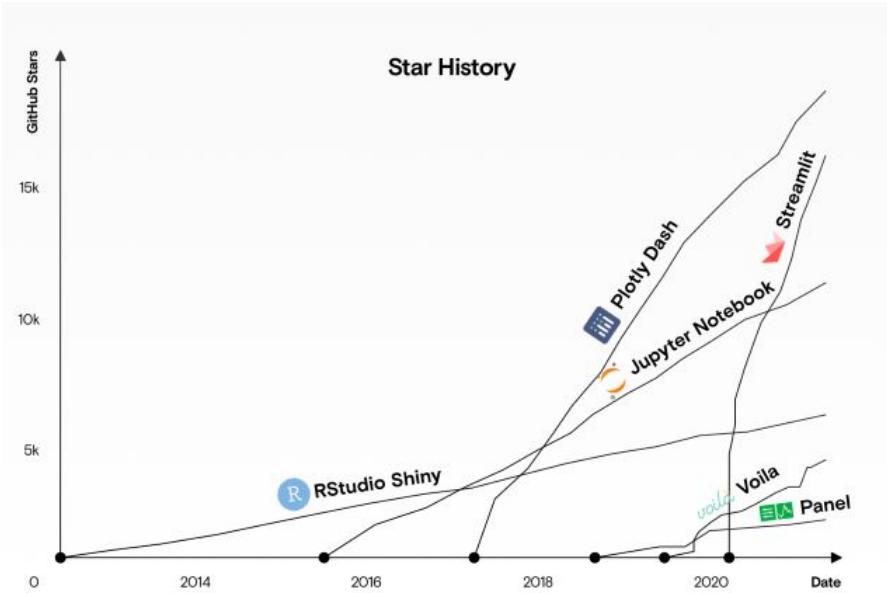


Figure 3.43: Data App - Comparison of Streamlit, RStudio-Shiny and others, Quelle: [DataRevenue-Blog](#)

I tested Streamlit³¹ and think: it is fantastic, because I didn't have to spend time on building a webpage or learn HTML, CSS or Wordpress. Everything is in Python and once the setup is done (which is easy) all I have to do for updating the whole data app is to save the Python file (no compiling needed). I believe that Python in combination with Streamlit is a very strong combination which will beat the “R / Tidyverse / Shiny” alternative! Here are some examples:

I used the data of the “Marathon runtimes” example and as you can see I only had to change some very minor things in the python code (like `import streamlit as st` and `write st.pyplot(g)` instead of `plt.show()`) in order to create a “data app”. You can upload another Excel-csv file by pressing the “Browse files” button, which will then be visualized. Using the checkboxes below will open more graphics (like histograms, kernel density, violin plots,...). See my “data app”³² and play around yourself. I uploaded the results of these two datasets (the left and right side of fig. 3.44) in the results folder.

²⁸ Tidyverse, <https://www.tidyverse.org/>

²⁹ Shiny-App, <https://shiny.rstudio.com/>

³⁰ Data-Revenue, <https://www.datarevenue.com/de-blog/streamlit-vs-dash-vs-shiny-vs-voila-vs-flask-vs-jupyter>

³¹ Streamlit, <https://www.streamlit.io/>

³² Marathon Streamlit “data app”, https://share.streamlit.io/andreastraut/visualize-results-in-apps/main/app_Example_Marathon_extended.py

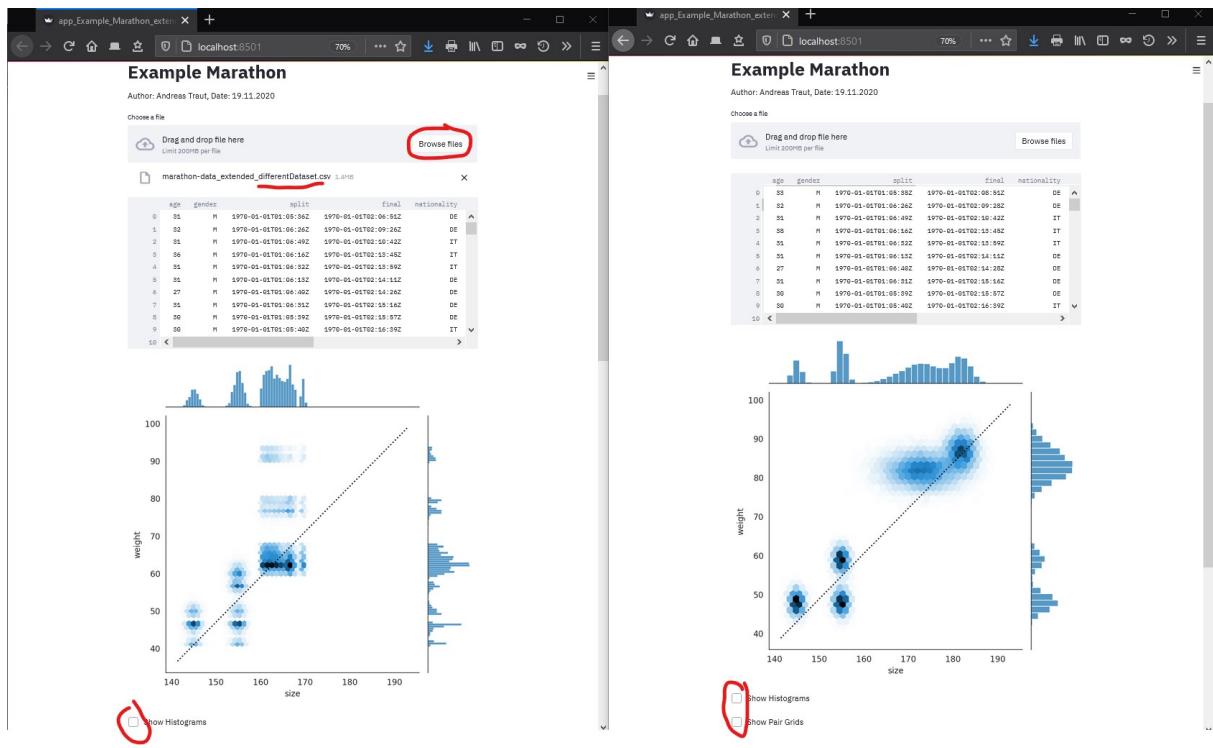


Figure 3.44: Data App - Marathon Example

And here is a second example: as everyone is talking about Corona/Covid and epidemiological models I thought that implementing the SEIR-Model would be an interesting example. Believe me: I read the Wikipedia SEIR-article³³ and implemented a Streamlit app³⁴ in less than half an hour. This is why I love Streamlit: highly efficient, lovely design and easy to deploy.

³³ Wikipedia SEIR model, <https://de.wikipedia.org/wiki/SEIR-Modell>

³⁴ SEIR modell app, https://share.streamlit.io/andreasraut/visualize-results-in-apps/main/app_SEIR_model

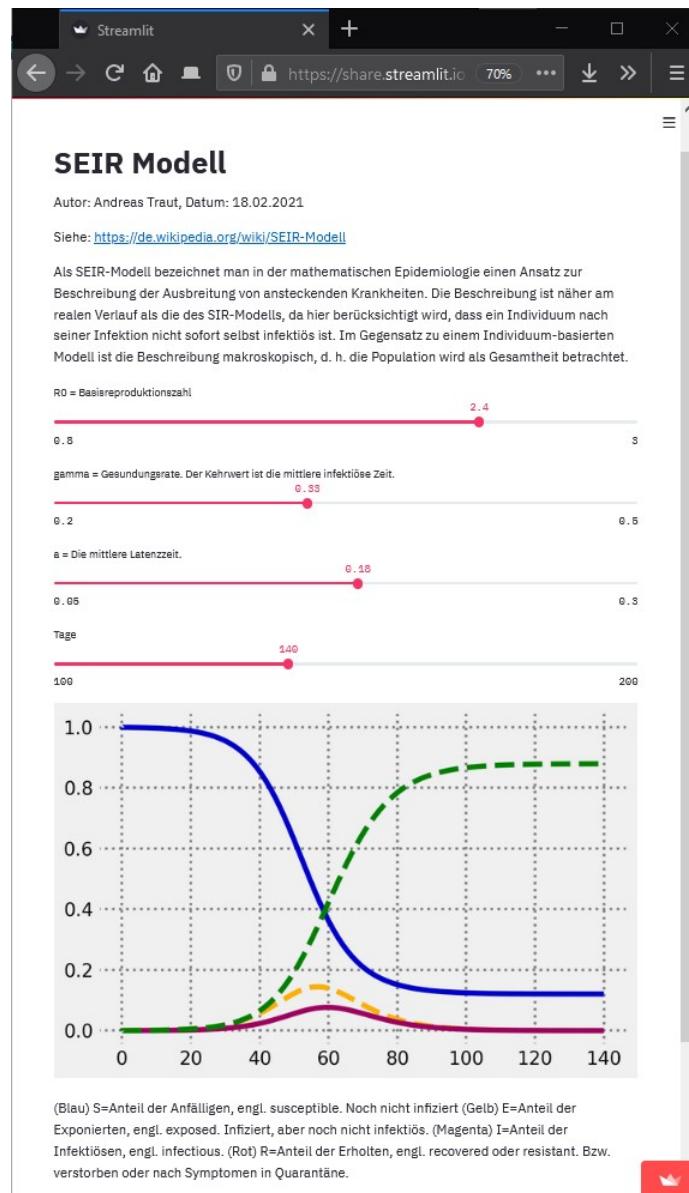


Figure 3.45: Data App - SEIR Model Example

3.4 Professional Tools

I will list some common professional tools, which offer visualization functionality and more. These tools cost some money, but I recommend to have a look into these: many companies use these or similar tools. There are different tools, which provide fantastic possibilities for visualization of data.

Additionally these tools provide a lot of functionality concerning other topics, like

- “**data integration**” (how can different data sources be connected?) or

- “**reporting**” (how can beautiful dashboards, which show all relevant graphics, be created?).

Obviously the best tools are not for free. I will only list up some examples here in my book and I recommend to read the official documentation on their websites for getting a feeling about what these tools can do and maybe also their limitations.

3.4.1 Power BI

See here: <https://powerbi.microsoft.com/de-de/>

Power BI from Microsoft is a very popular and probably the leading visualization tool for Big Data.

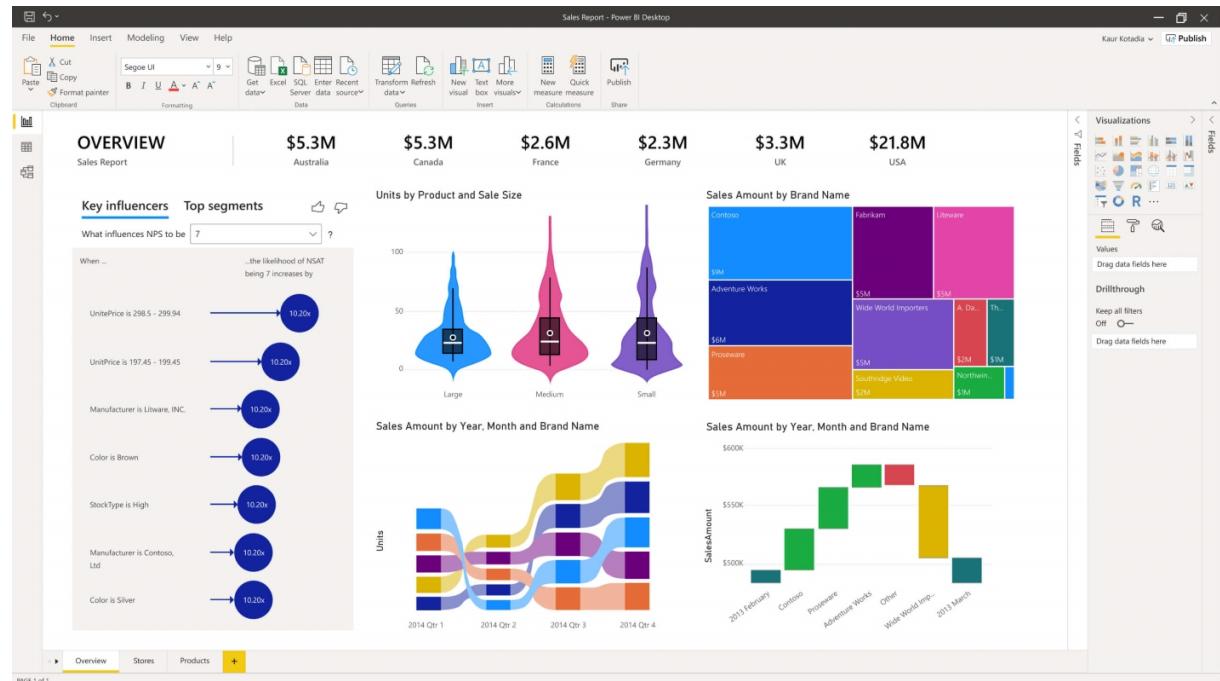


Figure 3.46: Professional BI - Power BI

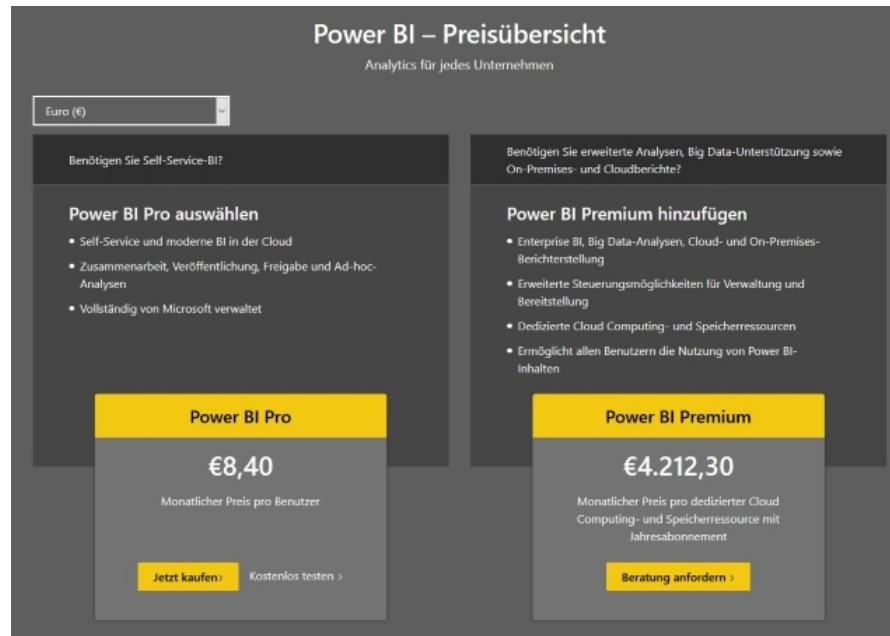


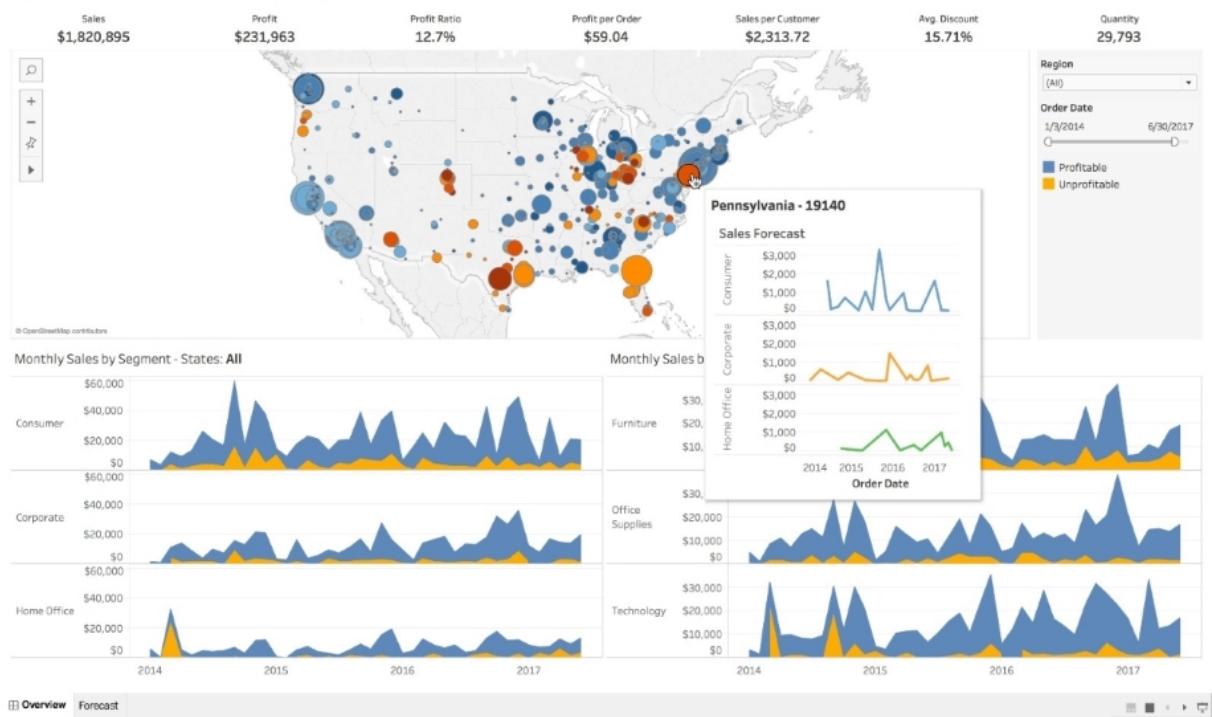
Figure 3.47: Professional BI - Power BI Prices

3.4.2 Tableau

See here: <https://www.tableau.com/de-de>

Tableau from Salesforce (which is an Oracle subsidiary) is a very interesting alternative.

Executive Overview - Profitability (All)


Figure 3.48: Professional BI - Tableau 1

From Zero to Senior Data Science

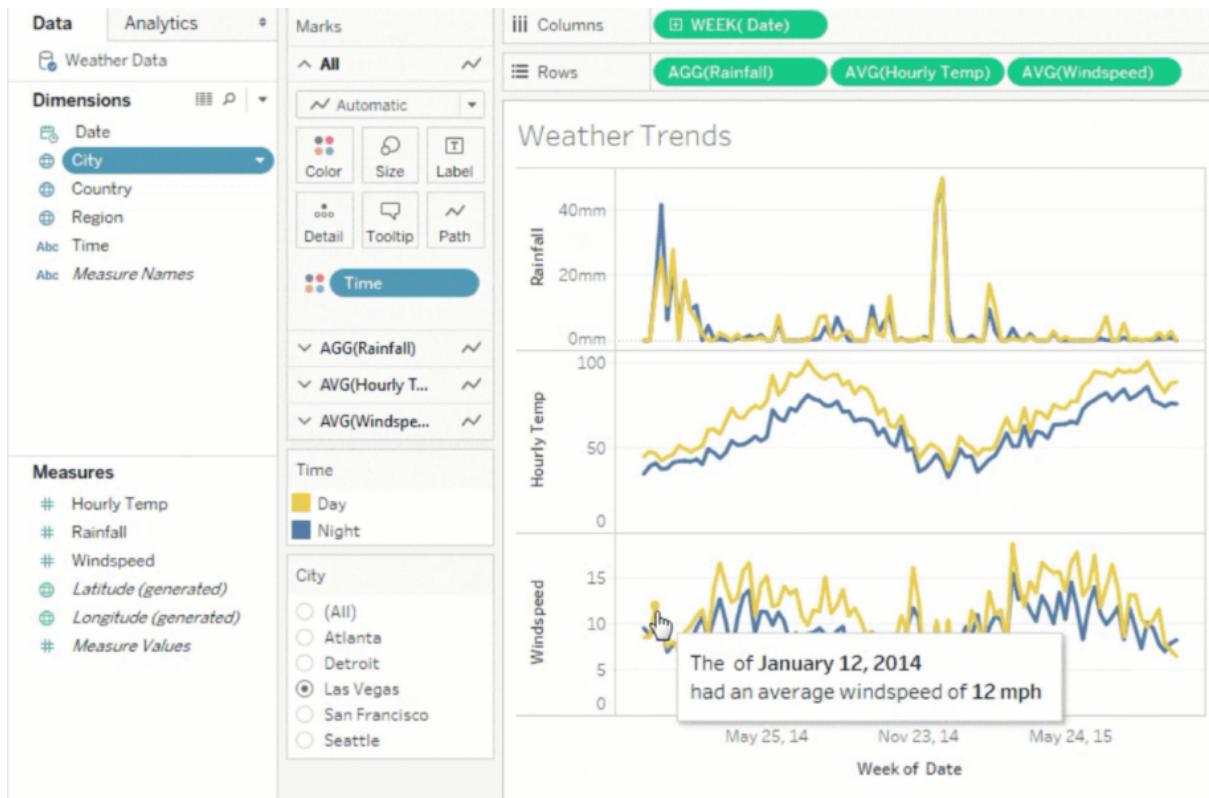


Figure 3.49: Professional BI - Tableau 2

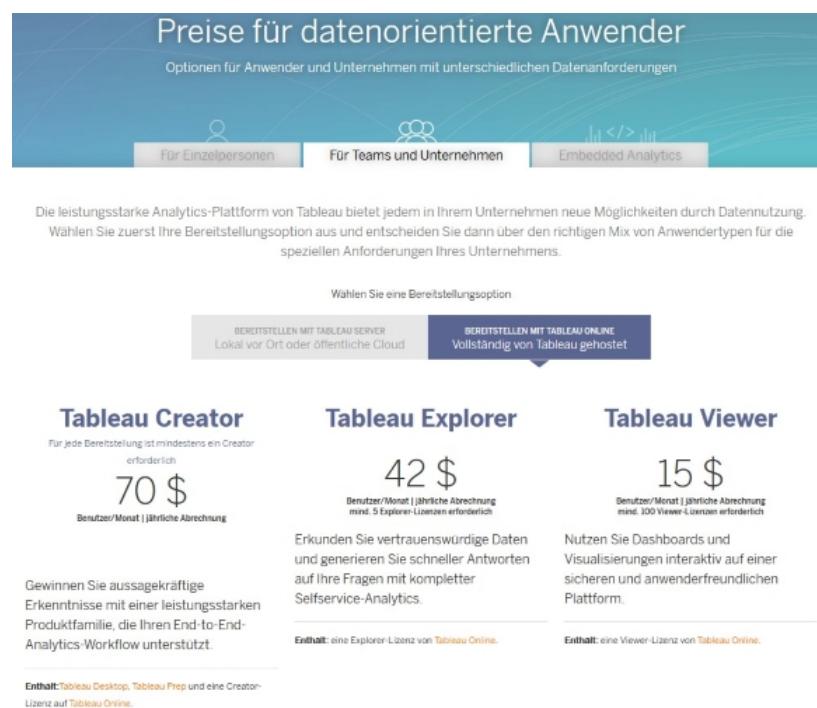


Figure 3.50: Professional BI - Tableau Prices

3.4.3 QLink

See here: <https://www.qlik.com/de-de/>

QLink is a fast growing tool for business intelligence and data visualization.

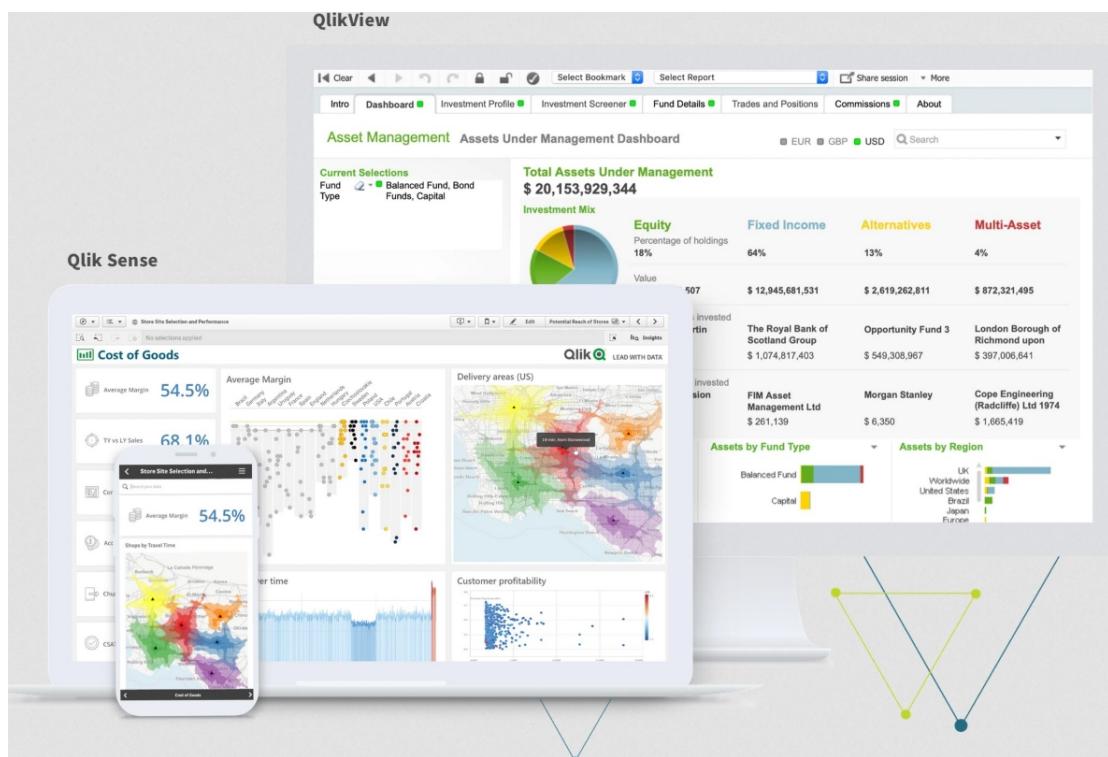


Figure 3.51: Professional BI - QLink

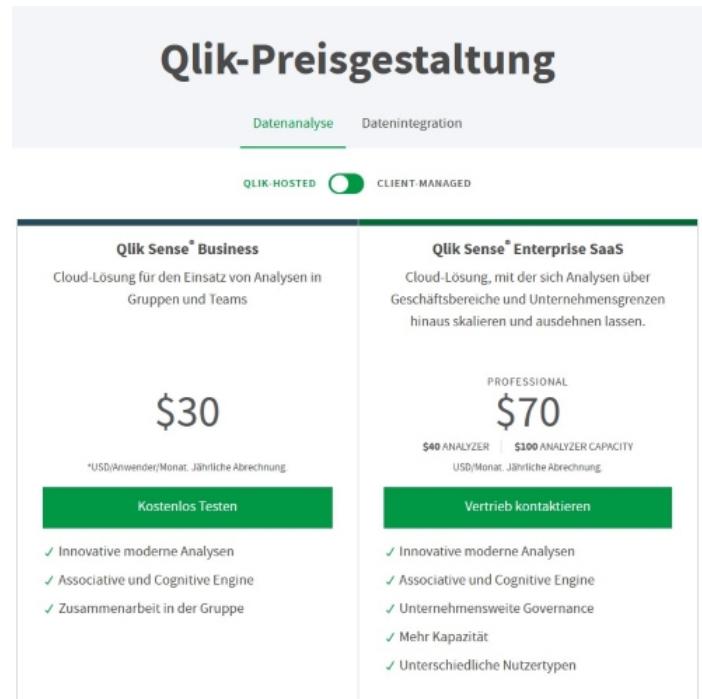


Figure 3.52: Professional BI - QLink Prices

4 Machine Learning with Python: Comparison

Small Data vs Big Data

4.1 Movies Database Example

A good starting point for finding useful datasets is Kaggle¹. I downloaded the movies dataset². The dataset from Kaggle contains the following columns:

Rank	Title	Year	Score	Metascore	Genre	Vote	
Director	Runtime	**Revenue**		Description	RevCat		

In this example I want to predict the **Revenue** based on the other information, which I have for each movie (e.g. every movie has a year, a scoring, a title ...). There are some “NaN”-values in the column “Revenue” and instead of filling them with an assumption (e.g. median-value) as I did in another Jupiter-Notebook³, I wanted to predict these values. You might guess the conclusion already: predicting the revenue based on the available information as shown above (the columns) might not work. But essential to me is more to follow a well established standard-process of data-cleaning, data-preparing, model-training and error-calculation in this example in order to learn how to apply this process to better datasets, than the movies-dataset, later.

Therefore, here is how I approached the problem step-by-step:

4.1.1 Import the Data

```
def load_data(path=PATH):
    csv_path = os.path.join(path, "movies.csv")
    return pd.read_csv(csv_path)
```

¹ Kaggle, www.kaggle.com

² Movies Dataset from Kaggle, <https://www.kaggle.com/isaactaylorofficial/imdb-10000-most-voted-feature-films-041118>

³ Movies Stratified Sample Extended Jupyter-Notebook, <https://github.com/AndreasTraut/Machine-Learning-with-Python/blob/master/Movies%20Machine%20Learning%20-%20StratifiedSample.ipynb>

```
movies = load_data()
movies.head()
```

Rank	Title	Year	Score	Metascore	Genre	Vote	Director	Runtime	Revenue	Description
1	The Shawshank Redemption	1994	94.3	80.0	Drama	2011509	Frank Darabont	142	28.34	Two imprisoned men bond over a number of years...
2	The Dark Knight	2008	93.0	84.0	Action, Crime, Drama	198020	Christopher Nolan	152	534.86	When the menace known as the Joker emerges fro...
3	Inception	2010	8.8	74.0	Action, Adventure, Sci-Fi	176020	Christopher Nolan	148	292.58	A thief who steals corporate secrets through t...
4	Fight Club	1999	8.8	66.0	Drama	1609459	David Fincher	139	37.03	An insomniac office worker and a devil-may-care...
5	Pulp Fiction	1994	8.9	94.0	Crime, Drama	1570194	Quentin Tarantino	154	107.93	The lives of two mob hitmen, a boxer, a gangst...

The datatypes of the columns are:

```
movies.dtypes
```

Rank	int64
Title	object
Year	int64
Score	float64
Metascore	float64
Genre	object
Vote	int64
Director	object
Runtime	int64
Revenue	float64
Description	object
dtype:	object

And with `movies.info()` we get also information about the count of non-null entries:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Rank              10000 non-null   int64  
 1   Title             10000 non-null   object  
 2   Year              10000 non-null   int64  
 3   Score             10000 non-null   float64 
 4   Metascore         6781 non-null   float64 
 5   Genre              10000 non-null   object  
 6   Vote              10000 non-null   int64  
 7   Director           9999 non-null   object  
 8   Runtime            10000 non-null   int64  
 9   Revenue            7473 non-null   float64 
 10  Description        10000 non-null   object  
dtypes: float64(3), int64(4), object(4)
memory usage: 859.5+ KB
```

4.1.2 Separate “NaN”-Values

In column “Revenue” there are 7473 “non-null” values, and 2527 “null” values. I separated the rows with “NaN”-values in column “Revenue”. These are the 2527 datarows, where column “Revenue” is null (see fig. 4.1):

These are the datarows, where column “Revenue” is null:

Revenue											Description
Rank	Title	Year	Score	Metascore	Genre	Vote	Director	Runtime			
82	83	A Clockwork Orange	1971	8.3	80.0	Crime, Drama, Sci-Fi	662768	Stanley Kubrick	136	NaN	In the future, a sadistic gang leader is impr...
513	514	To Kill a Mockingbird	1962	8.3	87.0	Crime, Drama	262064	Robert Mulligan	129	NaN	Atticus Finch, a lawyer in the Depression-era ...
581	582	Death Proof	2007	7.0	NaN	Action, Thriller	236539	Quentin Tarantino	113	NaN	Two separate sets of voluptuous women are stal...
620	621	My Neighbour Totoro	1988	8.2	86.0	Animation, Family, Fantasy	226126	Hayao Miyazaki	86	NaN	When two girls move to the country to be near ...
685	686	Hachi: A Dog's Tale	2009	8.1	NaN	Drama, Family	212349	Lasse Hallström	93	NaN	A college professor's bond with the abandoned ...

In [10]:	movies_RevenueNaN = movies[movies["Revenue"].isnull()]
Out[10]:	
In [11]:	len_movies_RevenueNaN = len(movies_RevenueNaN)
Out[11]:	2527

Figure 4.1: Movies Database - NaN Values

```
movies_RevenueNaN = movies[movies["Revenue"].isnull()]
len_movies_RevenueNaN = len(movies_RevenueNaN)
len_movies_RevenueNaN
```

And these are the 7473 columns where “Revenue” is not null:

```
movies_NotNull = movies[movies["Revenue"].notnull()]
len_movies_NotNull = len(movies_NotNull)
len_movies_NotNull
```

I want to see this in a plot:

```
fig, axs = plt.subplots()
axs.pie([len_movies_RevenueNaN, len_movies_NotNull],
        labels=['Revenue=NaN', 'Revenue NotNull'],
        colors = ['green', 'yellow']
       )
plt.show
```

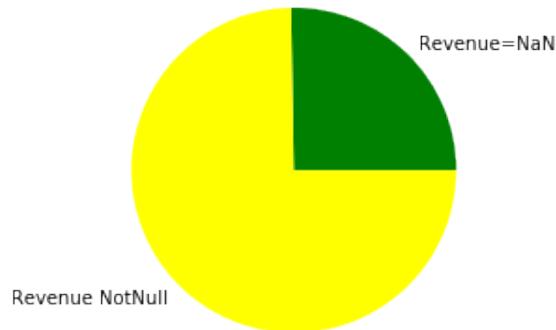


Figure 4.2: Movies Database - Plot Null and NotNull

4.1.3 Visualization of the Data

A first approach should always be create some visualization of the data in order to better understand them.

```
movies_NotNull['Revenue'].hist()
```

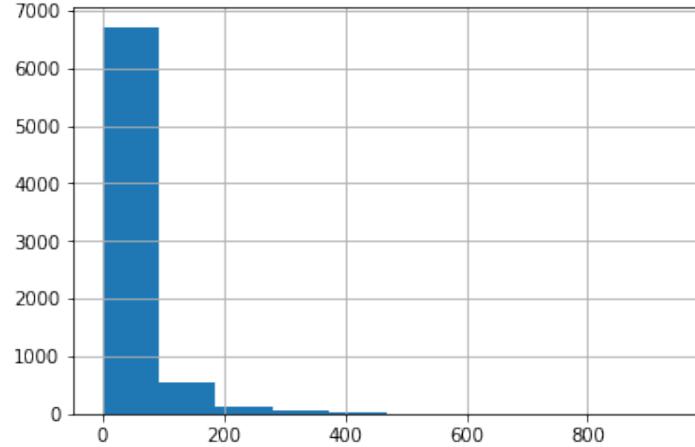


Figure 4.3: Movies Database - Histogram

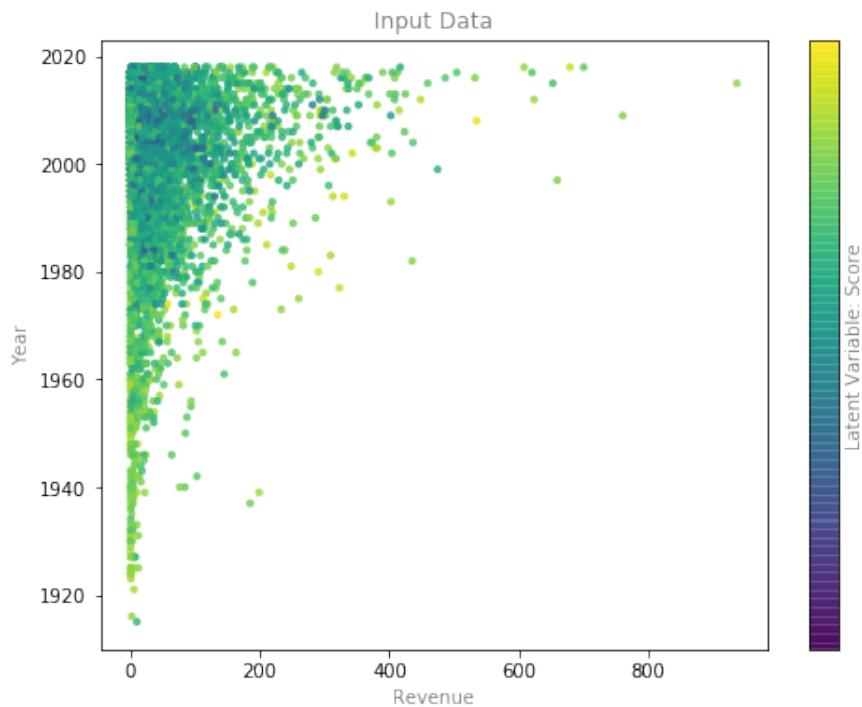


Figure 4.4: Movies Database - Scatterplot Revenue Year

```
fig, ax = plt.subplots(figsize=(8, 6))
point_style = dict(cmap='Paired', s=50)
pts = ax.scatter(movies_NotNull['Revenue'],
```

```
movies_NotNull['Year'],
c=movies_NotNull['Score'],
s=10,
alpha=0.8)
cb = fig.colorbar(pts, ax=ax)
# format plot
format_plot(ax, 'Input Data', 'Revenue', 'Year')
cb.set_ticks([])
cb.set_label('Latent Variable: Score', color='gray')
fig.savefig('images/movies/movies_revenue_year_score.png')
```

The result of this code is the plot fig. 4.4.

```
fig, ax = plt.subplots(figsize=(8, 6))
point_style = dict(cmap='Paired', s=50)
pts = ax.scatter(movies_NotNull['Year'],
                 movies_NotNull['Score'],
                 c=movies_NotNull['Revenue'],
                 s=10,
                 alpha=0.4)
cb = fig.colorbar(pts, ax=ax)
# format plot
format_plot(ax, 'Input Data', 'Year', 'Score')
cb.set_ticks([])
cb.set_label('Latent Variable: Score', color='gray')
fig.savefig('images/movies/movies_year_score_revenue.png')
```

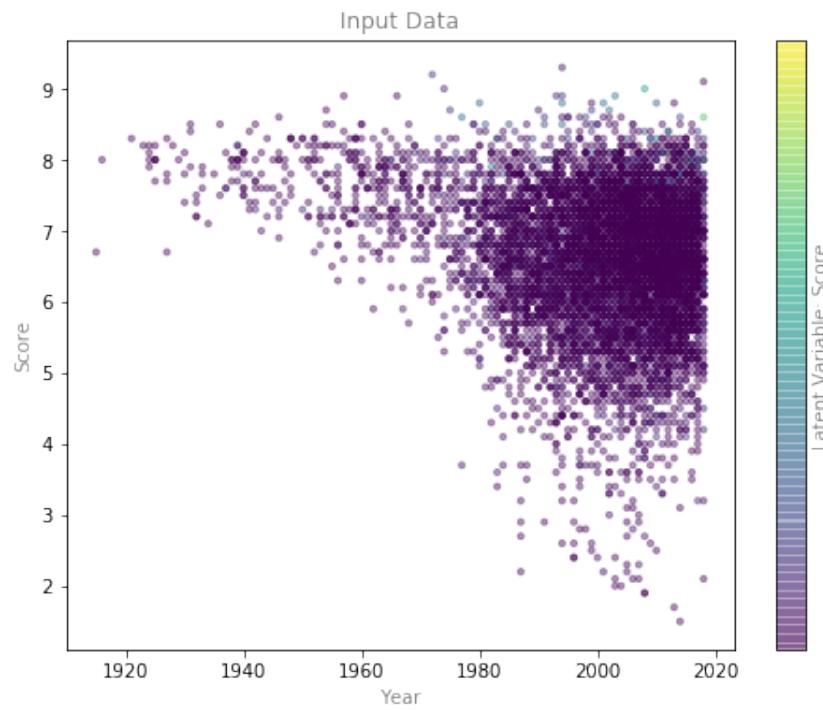


Figure 4.5: Movies Database - Scatterplot Year Score

4.1.4 Draw a Stratified Sample

I drew a stratified sample (based on “Revenue”) on this remaining dataset and I received a training dataset and testing dataset. First copy the dataset `movies_NotNull`, which we already prepared above:

```
movies_NotNullC = movies_NotNull[:].copy(deep=True)
```

Then let's create bins and count the values being in these bins (which is basically a histogram or the distribution):

```
movies_NotNullC['RevCat']=pd.cut(movies_NotNullC['Revenue'],
                                bins=[-1,100,200,300,np.inf],
                                labels=[1, 2, 3, 4])
movies_NotNullC['RevCat'].value_counts()
```

```
1    6753
2    522
3    120
```

```
4      78
Name: RevCat, dtype: int64
```

In order to get the proportions, I only have to divide by `len(movies_NotNullC)` which is 7473:

```
movies_NotNullC["RevCat"].value_counts() / len(movies_NotNullC)
```

```
1    0.903653
2    0.069851
3    0.016058
4    0.010438
Name: RevCat, dtype: float64
```

Now I split the dataset into a training and a testing dataset using stratified sampling in order to have the same distributions in these datasets:

```
split = StratifiedShuffleSplit(n_splits=1,
                               test_size=0.9,
                               random_state=42)
for train_index, test_index in split.split(movies_NotNullC,
                                            movies_NotNullC["RevCat"]):
    strat_train_set = movies_NotNullC.iloc[train_index]
    strat_test_set = movies_NotNullC.iloc[test_index]
```

Calculate the proportions of the bins on the stratified sample by dividing `len(strat_train_set)` which is 747. The expectation is, that these proportions are the same as above (otherwise we would have an error somewhere):

```
strat_train_set["RevCat"].value_counts() / len(strat_train_set)
```

```
1    0.903614
2    0.069612
3    0.016064
4    0.010710
Name: RevCat, dtype: float64
```

Looks good so far: the numbers look very similar to the ones above. This means: the proportions in the stratified sample (which has only 1/10-th of the data compare of the whole dataset) correspond to

the proportions in the whole dataset. That is what we wanted to have. Let's create a function `revenue_cat_proportions` to better compare these numbers: the overall data, the stratified sample and also a random sample. We can re-use this function later again:

```
def revenue_cat_proportions(data):
    return data["RevCat"].value_counts() / len(data)

train_set, test_set = train_test_split(movies_NotNullC,
                                       test_size=0.9,
                                       random_state=42)

compare_props = pd.DataFrame({
    "Overall": revenue_cat_proportions(movies_NotNullC),
    "Stratified": revenue_cat_proportions(strat_test_set),
    "Random": revenue_cat_proportions(test_set),
}).sort_index()

compare_props["Rand. %error"] = 100 * compare_props["Random"] /
    compare_props["Overall"] - 100
compare_props["Strat. %error"] = 100 * compare_props["Stratified"] /
    compare_props["Overall"] - 100

print(compare_props)

for set_ in (strat_train_set, strat_test_set):
    set_.drop("RevCat", axis=1, inplace=True)
```

	Overall	Stratified	Random	Rand. %error	Strat. %error
1	0.903653	0.903657	0.903063	-0.065336	0.000476
2	0.069851	0.069878	0.069432	-0.600432	0.038109
3	0.016058	0.016057	0.016652	3.699078	-0.004460
4	0.010438	0.010407	0.010853	3.983966	-0.289348

4.1.5 Split of Dataset into Training-Data and Test-Data

```
movies_train = strat_train_set.drop('Revenue', axis=1)
movies_train_labels = strat_train_set['Revenue'].copy()
len_movies_train = len(movies_train)

movies_test = strat_test_set.drop('Revenue', axis=1)
```

```
movies_test_labels = strat_test_set['Revenue'].copy()
len_movies_test = len(movies_test)
```

The whole dataset of 10000 rows has been split up into

- a training dataset (“movies_train”),
- a testing dataset (“movies_test”) and
- a dataset, where “Revenue”=“NaN”

Here are the counts:

```
len_movies_train = 747
len_movies_test = 6726

len_movies_train + len_movies_test = 7473
len_movies_NotNull = 7473
len_movies_RevenueNaN = 2527

len_movies_train + len_movies_test + len_movies_RevenueNaN = 10000
```

Let's visualize these counts:

```
fracs = [len_movies_train, len_movies_test, len_movies_RevenueNaN]
labels = ['training data (movies_train)', 'testing data (movies_test)',
          'Revenue=NaN']
fig, axs = plt.subplots()
axs.pie(fracs, labels=labels)
plt.show
```

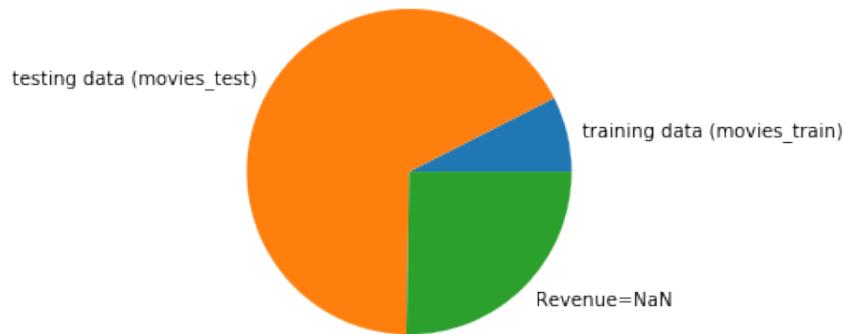


Figure 4.6: Movies Database - Stratified Sample (Train, Test, NaNs)

4.1.6 Create a Pipeline

Now I want to implement a pipeline for doing all the data preparation work quicker when I am testing it later. Only numerical-columns will be taken. The other will be thrown away (for the moment, I might change this later):

Created a pipeline to fill the “NaN”-value in other columns (e.g. “Metascore”, “Score”).

```
imputer = SimpleImputer(strategy="median")
movies_train_num = movies_train.select_dtypes(include=[np.number])
imputer.fit(movies_train_num)
X = imputer.transform(movies_train_num) # Transform the training set:
movies_tr = pd.DataFrame(X,
                          columns=movies_train_num.columns,
                          index=movies_train.index)
movies_tr.head()
```

	Rank	Year	Score	Metascore	Vote	Runtime
733	734.0	2007.0	6.3	68.0	204005.0	86.0
5851	5852.0	2003.0	5.2	21.0	16061.0	102.0
3816	3817.0	2003.0	5.3	47.0	33688.0	116.0
5384	5385.0	1980.0	6.7	58.0	18517.0	103.0
1058	1059.0	1998.0	6.7	63.0	152346.0	136.0

```
num_pipeline = Pipeline([('imputer', SimpleImputer(strategy='median'))])
num_attribs = ['Rank', 'Year', 'Score', 'Metascore', 'Vote', 'Runtime']
#list(movies_train_num)
full_pipeline = ColumnTransformer([('num', num_pipeline, num_attribs)])
```

Apply now the full pipeline on the training-dataset “movies_train”. But before we do this, we have a look into the column “Metascore” and the “NaN”-values in the training dataset “movies_train”:

```
tmp = movies_train[movies_train["Metascore"].isnull()]
tmp.head(2)
```

Rank	Title	Year	Score	Metascore	Genre	Vote	Director	Runtime	Description
53845385	The Final Countdown	1980	6.7	NaN	Action, Sci-Fi	18517	Don Taylor	103	A modern aircraft carrier is thrown back in ti...
59145915	The Chase	1994	45.8	NaN	Action, Adventure, Comedy	15791	Adam Rifkin	89	Escaped convict Jack Hammond takes a woman hos...

Now apply the Pipeline:

```
movies_train_prepared = full_pipeline.fit_transform(movies_train)
```

Now let's count the "nan" values in the new prepared dataset "movies_train_prepared". I have to transform it back to a Pandas-Dataframe format first:

```
tmp_num = movies_train.select_dtypes(include=[np.number])
tmp_prep = pd.DataFrame(movies_train_prepared,
                        columns=tmp_num.columns,
                        index=movies_train.index)
tmp = tmp_prep[tmp_prep["Metascore"].isnull()]
tmp
```

Rank	Year	Score	Metascore	Vote	Runtime

Zero, as we wanted! All "nan"-values in "movies_train_prepared" have been removed by the "median" value (this was how the pipeline was built). Great, this was part of the job, the Pipeline should have done. The other part was to eliminate some columns. We now have only 6 remaining instead of 10 columns.

4.1.7 Fit the Model with "DecisionTreeRegressor"

I used the training dataset and fittet it with the DecisionTreeRegressor model

```
tree_reg = DecisionTreeRegressor(random_state=42)
tree_reg.fit(movies_train_prepared, movies_train_labels)
```

```
movies_predictions = tree_reg.predict(movies_train_prepared)
```

How big is the error for all training-datasets?

```
trainmse = mean_squared_error(movies_train_labels, movies_predictions)
trainrmse = np.sqrt(trainmse)
```

4.1.8 Cross-Validation

I verified with a cross-validation, how good this model/parameters are

```
scores = cross_val_score(tree_reg,
                         movies_train_prepared,
                         movies_train_labels,
                         scoring="neg_mean_squared_error",
                         cv=10)
rmse_scores = np.sqrt(-scores)
def display_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())
display_scores(rmse_scores)
```

```
Scores: [ 69.66166984  62.59977724  62.6450061  112.98391756  47.99491086
         ← 52.16787811  42.65104005  83.21059338  41.09199207  63.84158945]
Mean: 63.88483746577791
Standard deviation: 20.447326452253154
```

4.1.9 Test the model

We take on arbitrary row in the testing dataset “movies_test”. Take for example row number 322:

```
some_data = movies_test.iloc[0:20]
some_data_label = movies_test_labels.iloc[0:20]
some_data.head(2)
```

Rank	Title	Year	Score	Metascore	Genre	Vote	Director	Runtime	Description
31233124	Victor Frankenstein	2015	6.0	36.0	Drama, Horror, Sci-Fi	45259	Paul McGuigan	110	Told from Igor's perspective, we see the trouble...
30613062	Battleship Potemkin	1925	8.0	Nan	Drama, History	46636	Sergei M. Eisenstein	75	In the midst of the Russian Revolution of 1905...

As we didn't apply the pipeline on the testing dataset (we only did on the training dataset "movies_train"), there might still some "nan" values in columns "Metascore".

Compare the true value ("some_data_labels") and the predicted value ("some_data_predictions") side-by-side. Left side is the true value from the original movies dataset. Right side is the predicted value based on the tree model:

```
some_movies = movies.iloc[some_data.index[0:len(some_data)]]
some_data_prepared = full_pipeline.fit_transform(some_data)
some_data_predictions = tree_reg.predict(some_data_prepared)
side_by_side = [(true, pred)
    for true, pred in
    zip(list(some_data_label),
        list(some_data_predictions))]
side_by_side
```

[(5.78, 10.91),
(0.05, 0.44),
(0.3, 2.68),
(159.6, 11.99),
(33.63, 2.19),
(44.9, 26.83),
(38.52, 22.52),
(33.04, 2.19),
(3.2, 11.99),
(37.49, 16.38),
(17.88, 64.19),
(41.19, 191.45),
(16.19, 12.19),
(0.05, 3.02),
(16.68, 64.19),
(35.11, 11.54),

```
(36.0, 40.22),  
(3.61, 19.64),  
(0.59, 0.05),  
(0.99, 5.48)]
```

The mean-squared-error is as follows:

```
mse = mean_squared_error(some_data_label, some_data_predictions)  
rmse = np.sqrt(mse)  
rmse
```

```
51.3429319867886
```

Taking now the whole testing dataset:

```
movies_test_prepared = full_pipeline.fit_transform(movies_test)  
movies_test_predictions = tree_reg.predict(movies_test_prepared)  
lin_mse = mean_squared_error(movies_test_labels, movies_test_predictions)  
lin_rmse = np.sqrt(lin_mse)
```

The result is `lin_rmse = 54.68`. We want to calculate also the mean and standard deviation: `movies_test_labels.mean()` which results in 36.20 and `movies_test_labels.std()` which is 60.60.

A side-by-side comparison off the testing dataset. Left side is the true value from the original movies dataset. Right side is the predicted value based on the tree model:

```
side_by_side = [(true, pred)  
    for true, pred in  
    zip(list(movies_test_labels),  
        list(movies_test_predictions))]
```

Plotting the true labels and the predicted labels on the testing dataset:

```
fig, ax = plt.subplots(figsize=(8, 8))  
pts = ax.scatter(movies_test_predictions,  
                 movies_test_labels,  
                 s=10,  
                 alpha=0.8)
```

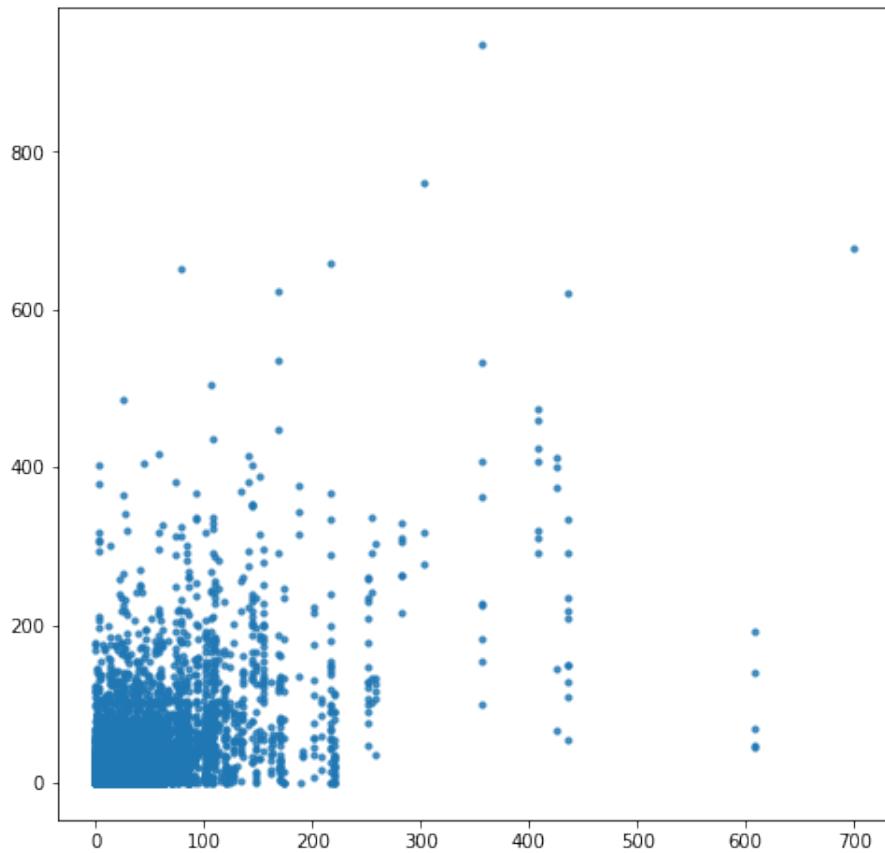


Figure 4.7: Movies Database - Plot True vs Predicted Value Side-by-Side Testdataset

Looks confusing. I would have expected something a bit more similar to the following plot. Plotting the same for the training dataset:

```
movies_train_prepared = full_pipeline.fit_transform(movies_train)
movies_train_predictions = tree_reg.predict(movies_train_prepared)
fig, ax = plt.subplots(figsize=(8, 6))
pts = ax.scatter(movies_train_predictions,
                 movies_train_labels,
                 s=10,
                 alpha=0.8)
```

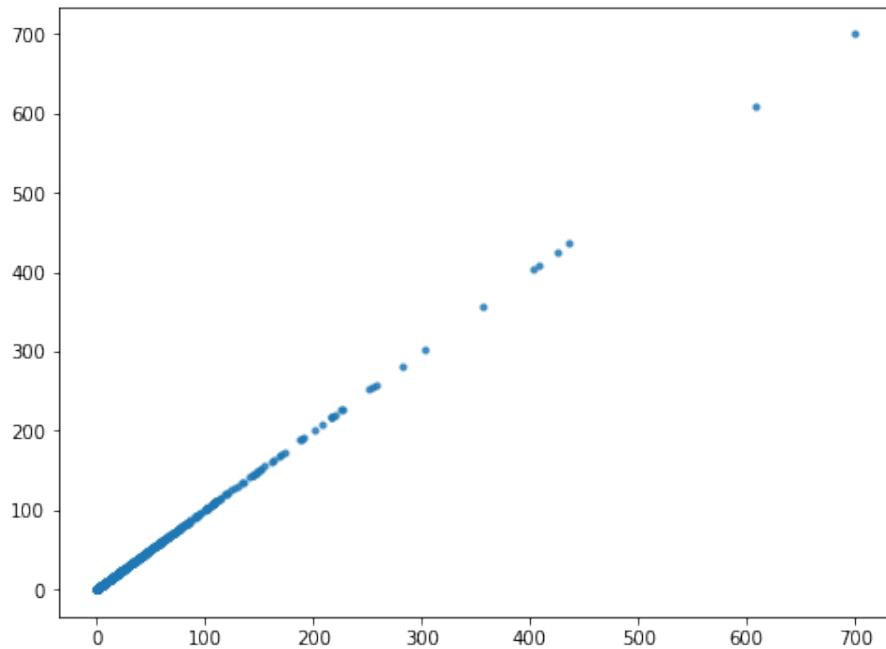


Figure 4.8: Movies Database - Plot True vs Predicted Value Side-by-Side Trainingdataset

Now I calculate the “Revenue” where it has “NaN”-values:

```
movies_RevenueNaN_prepared = full_pipeline.fit_transform(movies_RevenueNaN)
movies_RevenueNaN_predictions = tree_reg.predict(movies_RevenueNaN_prepared)
```

These are the predictions:

```
movies_RevenueNaN_predictions
```

```
array([74.1 , 11.99, 83.08, ..., 2.98, 43.49, 47.29])
```

I will insert the prediction into the dataset:

```
movies_RevenueNaN.loc[:, "Revenue"] = movies_RevenueNaN_predictions
```

4.1.10 Conclusion

The conclusion of this machine learning example is obvious: it is rather not possible to predict the “Revenue” based on the available numerical information (the most useful numerical features were

“year”, “score”, ... and the other categorical like “genre” don’t seem to have much more added value in my opinion). Lot’s of information, which is in the dataset has not yet been used, e.g. “Genre”, “Director”. These information could have an positive impact on the correctness of the predictions. But as “Genre” has 486 different values it is a bit more complicated to treat them as “categorial” values:

```
movies['Genre'].value_counts()
```

Comedy, Drama, Romance	494
Drama	482
Comedy, Drama	407
Drama, Romance	365
Comedy	357
...	
Action, Fantasy, War	1
War	1
Musical, Romance, War	1
Comedy, Romance, Family	1
Crime, Drama, Western	1
Name: Genre, Length: 486, dtype: int64	

As already mentionned right in the beginning of this Jupyter-Notebook the “OneHotEncoder” could be used. But before we should work on these 486 categorial values: could we simplify it, e.g. extract “Drama” and use it as a separate criteria? This will a topic in another Jupyter-Notebook.

On my GitHub profile you can find my Jupyter-Notebook for this example⁴

4.2 “Small Data”: Machine Learning using Scikit-Learn

The *second example* is a .py file for being used in an IDE (integrated developer environment), like the Spyder-IDE from the Anaconda distribution (see sec. 2 for hints on installing) and apply the *Scikit-Learn Python Machine Learning Library*⁵ (you may call this example a “Small Data” example if you want). I will show you a typical structure for a machine-learning example and put it into a mind-map (see fig. 4.9). The same structure will be applied on the third example, the “Big Data” example (see sec. 4.3).

So let’s start with the “Scikit-Learn”: the Mind-Map fig. 4.9 shows you what we need to do, starting from “1. import and create index”, then *discovering the data, preparing and cleaning the data, creating pipelines* and so on to “7. select and train model”. After that we will do the *cross-validation, save the*

⁴ Movies Example Jupyter Notebook, <https://github.com/AndreasTraut/Machine-Learning-with-Python/blob/master/Movies%20Machine%20Learning%20-%20Predict%20NaNs.ipynb>

⁵ Scikit-Learn Python Machine Learning Library, <https://scikit-learn.org/stable/>

model until we reach “11. evaluate final results”. You will find the same structure in the .py file and it should be a guide to work out your own problems with the same structure.

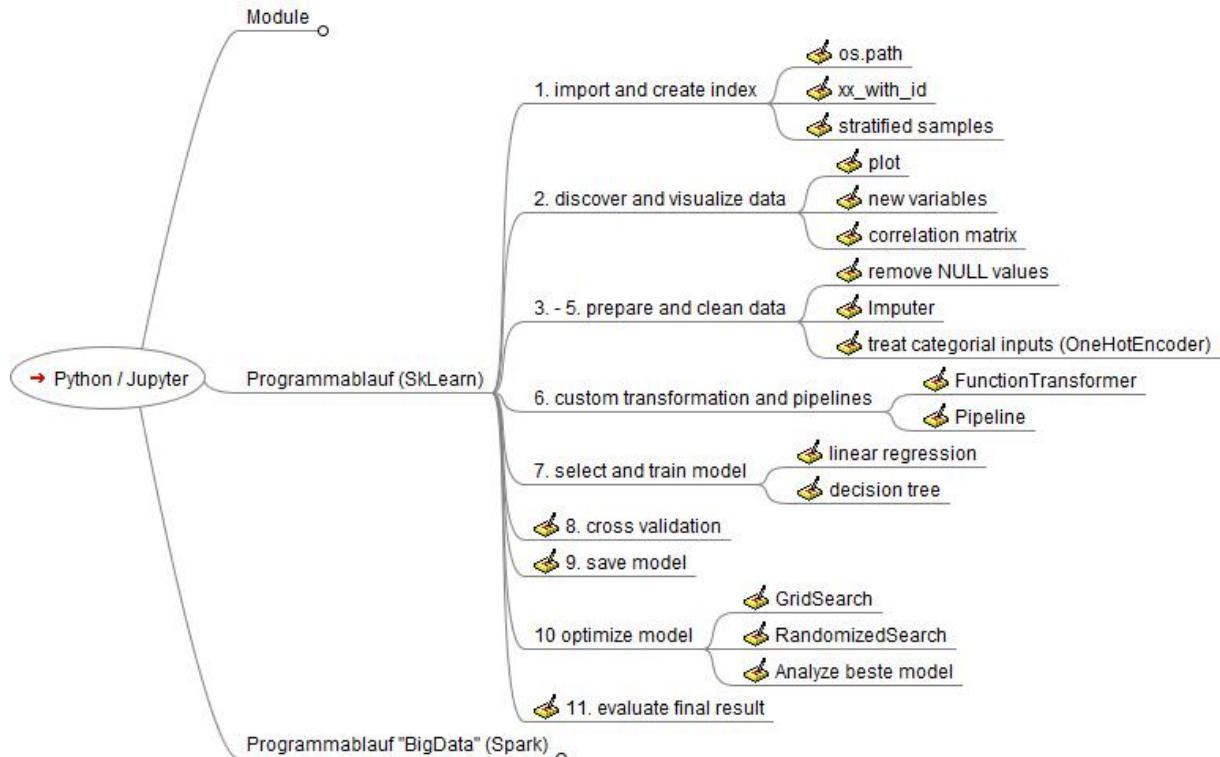


Figure 4.9: Mind Map - Scikit-Learn “Small Data”

```

#%## ##### #####
# 1. create index
  # 1.1 Alternative 1: generate id with static data
  # 1.2 Alternative 2: generate stratified sampling
  # 1.3 verify if stratified example is good
# 2. Discover and visualize the data to gain insights
# 3. prepare for Machine Learning
  # 3.1 find all NULL-values
  # 3.2 remove all NULL-values
# 4. Use "Imputer" to clean NaNs
# 5. treat "categorial" inputs
# 6. custom transformer and pipelines
  # 6.1 custom transformer
  # 6.2 pipelines
# 7. select and train model
  # 7.1 LinearRegression model
  # 7.2 DecisionTreeRegressor model
# 8. crossvalidation
  
```

```
# 8.1 for DecisionTreeRegressor  
# 8.2 for LinearRegression  
# 8.3 for RandomForestRegressor  
# 8.4 for ExtraTreesRegressor  
# 9. Save Model  
# 10. Optimize Model  
    # 10.1 GridSearchCV  
        # 10.1.1 GridSearchCV on RandomForestRegressor  
        # 10.1.2 GridSearchCV on LinearRegressor  
    # 10.2 Randomized Search  
    # 10.3 Analyze best models  
# 11. Evaluate final model on test dataset  
##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### ##### #####
```

I aligned this “Small Data” structure to the Apache Spark “Big Data” structure (see Mind Map fig. 4.14) in order to learn from each of this two approaches. Finally I will put these two Mind Maps into one big (see fig. 4.21) which you can take as a guide to navigate through all of your machine-learning problems.

Common Imports

These are the common imports which we need. For the moment it is not necessary to understand all of this. If you want to know more about what these imports do (e.g. `sklearn.model_selection`), then use my references and read to the official documentation (in case of `sklearn` you have to go to the Scikit-Learn website). Getting familiar with the official documentation and learning how to quickly find, what you need for your specific problem is always a good idea. Most of these official documentation are well structured and once you understood, how to navigate through them you won’t need Google or Stackoverflow⁶ (which is a question and answer forum for programmers) to solve your problems.

```
#To support both python 2 and python 3  
from __future__ import division, print_function, unicode_literals  
  
# Common imports  
import pandas as pd  
import numpy as np  
import matplotlib as mpl  
import matplotlib.pyplot as plt  
import os  
import tarfile  
from six.moves import urllib
```

⁶ Stackoverflow, <https://stackoverflow.com>

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import FunctionTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
import hashlib
from sklearn.model_selection import StratifiedShuffleSplit
from pandas.plotting import scatter_matrix
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor
import joblib
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint
from scipy import stats
```

Read the csv-file

It is always a good idea to create variables, which can be reused later. This makes the code easier to read. Here the path to the CSV file has been set to the variable HOUSING_PATH.

```
HOUSING_PATH = os.path.join("datasets", "housing")
def load_housing_data(housing_path=HOUSING_PATH):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)
housing = load_housing_data()
```

4.2.1 Create Index (1)

4.2.1.1 Alternative 1: Generate ID with Static Data (1.1)

```
housing_with_id["id"] = housing["longitude"] * 1000 + housing["latitude"]
train_set, test_set = split_train_test_by_id(housing_with_id, 0.2, "id")
```

4.2.1.2 Alternative 2: Generate Stratified Sampling (1.2)

```

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
strat_test_set["income_cat"].value_counts() / len(strat_test_set)
housing["income_cat"].value_counts() / len(housing)

```

4.2.1.3 Verify if Stratified Example is good (1.3)

```

def income_cat_proportions(data):
    return data["income_cat"].value_counts() / len(data)
train_set, test_set = train_test_split(housing, test_size=0.2,
                                       random_state=42)
compare_props = pd.DataFrame({
    "Overall": income_cat_proportions(housing),
    "Stratified": income_cat_proportions(strat_test_set),
    "Random": income_cat_proportions(test_set),
}).sort_index()

compare_props["Rand. %error"] = 100 * compare_props["Random"] /
    compare_props["Overall"] - 100
compare_props["Strat. %error"] = 100 * compare_props["Stratified"] /
    compare_props["Overall"] - 100
print(compare_props)

```

4.2.2 Discover and Visualize the Data to Gain Insights (2)

Create some visualizations as I described in sec. 3. For example:

```

housing.plot(kind="scatter",
             x="longitude",
             y="latitude",
             alpha=0.4,
             s=housing["population"]/100,
             label="population",
             figsize=(10,7),
             c="median_house_value",
             cmap=plt.get_cmap("jet"),
             colorbar=True,

```

```
    sharex=False,
    title="housing_prices_scatterplot")
plt.legend()
save_fig("housing_prices_scatterplot")
```

4.2.3 Prepare for Machine Learning (3)

```
housing = strat_train_set.drop("median_house_value",
                               axis=1) # drop labels for training set
housing_labels = strat_train_set["median_house_value"].copy()
```

4.2.3.1 Find all NULL-Values (3.1)

```
print(housing.info())
print("Are there nans in column total_bedrooms?\n",
      housing["total_bedrooms"].isnull().any())
print("Show rows with nan:\n",
      housing[housing["total_bedrooms"].isnull()])
```

4.2.3.2 Remove all NULL-Values (3.2)

There are different options for handling NULL values, which occur in a column: we can delete the entire row or we can delete the entire column or we can fill these with an assumption, like for example the median (which is often a good assumption for filling NULL values, but not always, as we have seen in the movies database example).

```
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
# sample_incomplete_rows.dropna(subset=["total_bedrooms"])      # option 1
# sample_incomplete_rows.drop("total_bedrooms", axis=1)        # option 2
median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) #
↪ option 3 # In[56]:
print("sample_incomplete_rows\n",
      sample_incomplete_rows)
```

4.2.4 Use “Imputer” to Clean NaNs (4)

Remove all text attributes because median can only be calculated on numerical attributes

```
imputer = SimpleImputer(strategy="median")
housing_num = housing.select_dtypes(include=[np.number]) #or: housing_num =
    ↵ housing.drop('ocean_proximity', axis=1)
imputer.fit(housing_num)
print("imputer.strategy\n",
      imputer.strategy)
print("imputer.statistics_\n",
      imputer.statistics_)
print("housing_num.median\n",
      housing_num.median().values) # Check that this is the same as
    ↵ manually computing the median of each attribute:
print("housing_num.mean\n",
      housing_num.mean().values) # Check that this is the same as manually
    ↵ computing the median of each attribute.
```

Transform the training set:

```
X = imputer.transform(housing_num) # :
housing_tr = pd.DataFrame(X, columns=housing_num.columns,
    ↵ index=housing.index)
housing_tr.loc[sample_incomplete_rows.index.values]
```

4.2.5 Treat “Categorical” Inputs (5)

Use the OneHotEncoder for the categorial values:

```
cat_encoder = OneHotEncoder()
housing_cat_1hot = cat_encoder.fit_transform(housing_cat)
```

4.2.6 Custom Transformer and Pipelines (6)

4.2.6.1 Custom Transformer (6.1)

```
def add_extra_features(X, add_bedrooms_per_room=True):
    rooms_per_household = X[:, rooms_ix] / X[:, household_ix]
```

```
population_per_household = X[:, population_ix] / X[:, household_ix]
if add_bedrooms_per_room:
    bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
    return np.c_[X, rooms_per_household,
                population_per_household,
                bedrooms_per_room]
else:
    return np.c_[X, rooms_per_household,
                population_per_household]

attr_adder = FunctionTransformer(add_extra_features,
                                 validate=False,
                                 kw_args={"add_bedrooms_per_room": False})
housing_extra_attribs = attr_adder.fit_transform(housing.values)

housing_extra_attribs = pd.DataFrame(
    housing_extra_attribs,
    columns=list(housing.columns)+["rooms_per_household",
        "population_per_household"],
    index=housing.index)
print("housing_extra_attribs.head()\n", housing_extra_attribs.head())
```

4.2.6.2 Pipelines (6.2)

```
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', FunctionTransformer(add_extra_features,
                                          validate=False)),
    ('std_scaler', StandardScaler()),
])
housing_num_tr = num_pipeline.fit_transform(housing_num)
print("housing_num_tr\n", housing_num_tr)

try:
    from sklearn.compose import ColumnTransformer
except ImportError:
    from future_encoders import ColumnTransformer # Scikit-Learn < 0.20

num_attribs = list(housing_num)
cat_attribs = ["ocean_proximity"]
```

```
full_pipeline = ColumnTransformer([
    ("num", num_pipeline, num_attribs),
    ("cat", OneHotEncoder(), cat_attribs),
])
housing_prepared = full_pipeline.fit_transform(housing)
print("housing_prepared\n", housing_prepared)
```

4.2.7 Select and Train Model (7)

4.2.7.1 LinearRegression Model (7.1)

```
lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)
# let's try the full preprocessing pipeline on a few training instances

some_data = housing.iloc[:1]
some_labels = housing_labels.iloc[:1]
some_data_prepared = full_pipeline.transform(some_data)
print("Predictions:\n", lin_reg.predict(some_data_prepared))
print("Labels:\n", list(some_labels)) # Compare against the actual values:

housing_predictions = lin_reg.predict(housing_prepared)
lin_mse = mean_squared_error(housing_labels, housing_predictions)
lin_rmse = np.sqrt(lin_mse)
print("lin_rmse\n", lin_rmse)
```

4.2.7.2 DecisionTreeRegressor Model (7.2)

```
tree_reg = DecisionTreeRegressor(random_state=42)
tree_reg.fit(housing_prepared, housing_labels)
housing_predictions = tree_reg.predict(housing_prepared)
tree_mse = mean_squared_error(housing_labels, housing_predictions)
tree_rmse = np.sqrt(tree_mse)
print("tree_rmse\n", tree_rmse)
```

4.2.8 Cross-Validation (8)

4.2.8.1 For DecisionTreeRegressor (8.1)

```
scores = cross_val_score(tree_reg,
                         housing_prepared,
                         housing_labels,
                         scoring="neg_mean_squared_error",
                         cv=10)
tree_rmse_scores = np.sqrt(-scores)
def display_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())
```

4.2.8.2 For LinearRegression (8.2)

```
lin_scores = cross_val_score(lin_reg,
                             housing_prepared,
                             housing_labels,
                             scoring="neg_mean_squared_error",
                             cv=10)
lin_rmse_scores = np.sqrt(-lin_scores)
display_scores(lin_rmse_scores)
```

4.2.8.3 For RandomForestRegressor (8.3)

```
forest_reg = RandomForestRegressor(n_estimators=10,
                                   random_state=42)
forest_reg.fit(housing_prepared,
               housing_labels)

housing_predictions = forest_reg.predict(housing_prepared)
forest_mse = mean_squared_error(housing_labels,
                                 housing_predictions)
forest_rmse = np.sqrt(forest_mse)
print(forest_rmse)
```

```
forest_scores = cross_val_score(forest_reg,
                                 housing_prepared,
                                 housing_labels,
                                 scoring="neg_mean_squared_error",
                                 cv=10)
forest_rmse_scores = np.sqrt(-forest_scores)
display_scores(forest_rmse_scores)
```

4.2.8.4 For ExtraTreesRegressor (8.4)

```
extratree_reg = ExtraTreesRegressor(n_estimators=10,
                                    random_state=42)
extratree_reg.fit(housing_prepared,
                  housing_labels)
housing_predictions = extratree_reg.predict(housing_prepared)
extratree_mse = mean_squared_error(housing_labels,
                                   housing_predictions)
extratree_rmse = np.sqrt(extratree_mse)
print(extratree_rmse)
extratree_scores = cross_val_score(extratree_reg,
                                   housing_prepared,
                                   housing_labels,
                                   scoring = "neg_mean_squared_error",
                                   cv=10)
extratree_rmse_scores = np.sqrt(-extratree_scores)
display_scores(extratree_rmse_scores)
```

4.2.9 Save Model (9)

```
joblib.dump(forest_reg, "forest_reg.pkl")
# and later...
my_model_loaded = joblib.load("forest_reg.pkl")
```

4.2.10 Optimize Model (10)

4.2.10.1 GridSearchCV (10.1)

```
param_grid = [
    # try 12 (3x4) combinations of hyperparameters
    {'n_estimators': [30, 40, 50],
     'max_features': [2, 4, 6, 8, 10]},
    # then try 6 (2x3) combinations with bootstrap set as False
    {'bootstrap': [False],
     'n_estimators': [3, 10],
     'max_features': [2, 3, 4]},
]
forest_reg = RandomForestRegressor(random_state=42)
# train across 5 folds, that's a total of (12+6)*5=90 rounds of training
grid_search = GridSearchCV(forest_reg,
                           param_grid,
                           cv=5,
                           scoring='neg_mean_squared_error',
                           return_train_score=True)
grid_search.fit(housing_prepared, housing_labels)
print(grid_search.best_params_)
print(grid_search.best_estimator_)
cvres = grid_search.cv_results_
for mean_score, params in zip(cvres["mean_test_score"],
                               cvres["params"]):
    print(np.sqrt(-mean_score), params)
```

4.2.10.1.1 GridSearchCV on RandomForestRegressor (10.1.1)

```
param_grid = [
    # try 12 (3x4) combinations of hyperparameters
    {'fit_intercept': [True],
     'n_jobs': [2, 4, 6, 8, 10]},
    # then try 6 (2x3) combinations with bootstrap set as False
    {'normalize': [False],
     'n_jobs': [3, 10]},
]
lin_reg = LinearRegression()

# train across 5 folds, that's a total of (12+6)*5=90 rounds of training
lin_grid_search = GridSearchCV(lin_reg,
                               param_grid,
```

```

        cv=5,
        scoring='neg_mean_squared_error',
        return_train_score=True)
lin_grid_search.fit(housing_prepared,
                    housing_labels)

# print(lin_grid_search.best_params_)

print(lin_grid_search.best_estimator_)
cvres = lin_grid_search.cv_results_
for mean_score, params in zip(cvres["mean_test_score"],
                               cvres["params"]):
    print(np.sqrt(-mean_score), params)

547     # from sklearn.model_selection import GridSearchCV
548     param_grid = [
549         # try 12 (3x4) combinations of hyperparameters
550         {'fit_intercept': [True], 'n_jobs': [2, 4, 6, 8, 10]},
551         # then try 6 (2x3) combinations with bootstrap set as False
552         {'normalize': [False], 'n_jobs': [3, 10]},
553     ]
554
555     lin_reg = LinearRegression()
556     # train across 5 folds, that's a total of (12+6)*5=90 rounds of training
557     lin_grid_search = GridSearchCV(lin_reg, param_grid, cv=5,
558                                     scoring='neg_mean_squared_error',
559                                     return_train_score=True)
560     lin_grid_search.fit(housing_prepared, housing_labels)
561     # print(lin_grid_search.best_params_)
562     print(lin_grid_search.best_estimator_)
563     cvres = lin_grid_search.cv_results_
564     for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
565         print(np.sqrt(-mean_score), params)

```

Figure 4.10: Small Data - Optimize Model GridSearchCV on LinearRegressor (10.1.2)

4.2.10.1.2 GridSearchCV on LinearRegressor (10.1.2)

4.2.10.2 Randomized Search (10.2)

```

param_distributions = {
    'n_estimators': randint(low=1, high=200),
    'max_features': randint(low=1, high=8),
}

forest_reg = RandomForestRegressor(random_state=42)

```

```
rnd_search = RandomizedSearchCV(forest_reg,
                                 param_distributions=param_distrib,
                                 n_iter=10, cv=5,
                                 scoring='neg_mean_squared_error',
                                 random_state=42)
rnd_search.fit(housing_prepared,
               housing_labels)
cvres = rnd_search.cv_results_
for mean_score, params in zip(cvres["mean_test_score"],
                               cvres["params"]):
    print(np.sqrt(-mean_score), params)
```

```
573     # from sklearn.model_selection import RandomizedSearchCV
574     # from scipy.stats import randint
575     param_distrib = {
576         'n_estimators': randint(low=1, high=200),
577         'max_features': randint(low=1, high=8),
578     }
579
580     forest_reg = RandomForestRegressor(random_state=42)
581     rnd_search = RandomizedSearchCV(forest_reg,
582                                     param_distributions=param_distrib,
583                                     n_iter=10, cv=5,
584                                     scoring='neg_mean_squared_error',
585                                     random_state=42)
586     rnd_search.fit(housing_prepared, housing_labels)
587     cvres = rnd_search.cv_results_
588     for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
589         print(np.sqrt(-mean_score), params)
```

Figure 4.11: Small Data - Optimize Model Randomized Search (10.2)

4.2.10.3 Analyze best models (10.3)

```
feature_importances = grid_search.best_estimator_.feature_importances_
feature_importances
extra_attribs = ["rooms_per_hhold", "pop_per_hhold", "bedrooms_per_room"]
#cat_encoder = cat_pipeline.named_steps["cat_encoder"] # old solution
cat_encoder = full_pipeline.named_transformers_["cat"]
cat_one_hot_attribs = list(cat_encoder.categories_[0])
attributes = num_attribs + extra_attribs + cat_one_hot_attribs
sorted(zip(feature_importances, attributes), reverse=True)
```

```

595     feature_importances = grid_search.best_estimator_.feature_importances_
596     feature_importances
597     extra_attribs = ["rooms_per_hhold", "pop_per_hhold", "bedrooms_per_room"]
598     cat_encoder = full_pipeline.named_transformers_["cat"]
599     cat_one_hot_attribs = list(cat_encoder.categories_[0])
600     attributes = num_attribs + extra_attribs + cat_one_hot_attribs
601     sorted(zip(feature_importances, attributes), reverse=True)

```

Figure 4.12: Small Data - Optimize Model Analyze best models (10.3)

4.2.11 Evaluate final model on test dataset (11)

```

final_model = grid_search.best_estimator_

X_test = strat_test_set.drop("median_house_value", axis=1)
y_test = strat_test_set["median_house_value"].copy()

X_test_prepared = full_pipeline.transform(X_test)
final_predictions = final_model.predict(X_test_prepared)

final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)

print ("final_predictions\n", final_predictions )
print ("final_rmse \n", final_rmse )

confidence = 0.95
squared_errors = (final_predictions - y_test) ** 2
mean = squared_errors.mean()
m = len(squared_errors)

# from scipy import stats

print("95% confidence interval: ",
      np.sqrt(stats.t.interval(confidence,
                               m - 1,
                               loc=np.mean(squared_errors),
                               scale=stats.sem(squared_errors))))
)

```

```
610 final_model = grid_search.best_estimator_
611
612 X_test = strat_test_set.drop("median_house_value", axis=1)
613 y_test = strat_test_set["median_house_value"].copy()
614
615 X_test_prepared = full_pipeline.transform(X_test)
616 final_predictions = final_model.predict(X_test_prepared)
617
618 final_mse = mean_squared_error(y_test, final_predictions)
619 final_rmse = np.sqrt(final_mse)
620
621 print ("final_predictions\n", final_predictions )
622 print ("final_rmse \n", final_rmse )
623
624 confidence = 0.95
625 squared_errors = (final_predictions - y_test) ** 2
626 mean = squared_errors.mean()
627 m = len(squared_errors)
628
629 # from scipy import stats
630 ▼ print("95% confidence interval: ",
631 ▼     np.sqrt(stats.t.interval(confidence, m - 1,
632 loc=np.mean(squared_errors),
633 scale=stats.sem(squared_errors)))
634 )
```

Figure 4.13: Small Data - Evaluate final model on test dataset (11)

4.3 “Big Data”: Machine Learning using Spark ML Library

This will be an example for a “Big-Data”⁷ environment and uses the *Apache Spark MLlib*⁸ scalable machine learning library. Various tutorials, documentation, “code-fragments” and guidelines can be found in the internet **for free** (at least for your private and non-commercial use). The best is in my opinion the official documentation ⁹. A few more helpful sources are the following GitHub repositories:

- tirthajyoti/Spark-with-Python¹⁰ (MIT Licence)
- Apress/learn-pyspark¹¹ (Freeware License)
- mahmoudparsian/pyspark-tutorial¹² (Apache License v2.0)

Concerning the “Big Data” topic I want to add the following: I passed a certification as “*Data Scientist Specialized in Big Data Analytics*” at Fraunhofer Institute. I must say: Understanding the concept

⁷ Big-Data, Wikipedia, https://en.wikipedia.org/wiki/Big_Data

⁸ Apache Spark MLlib, <https://spark.apache.org/mllib/>

⁹ Official Apache Spark ML documentation, <https://spark.apache.org/docs/latest/ml-guide.html>

¹⁰ GitHub repository “tirthajyoti/Spark-with-Python”, <https://github.com/tirthajyoti/Spark-with-Python> (MIT Licence)

¹¹ GitHub repository “Apress/learn-pyspark”, <https://github.com/Apress/learn-pyspark> (Freeware License)

¹² GitHub repository “mahmoudparsian/pyspark-tutorial”, <https://github.com/mahmoudparsian/pyspark-tutorial> (Apache License v2.0)

of “Big-Data” and how to differentiate “standard” machine learning from a “scalable” Big Data environment is not easy. I recommend a separate training! Some steps are a bit similar to “Scikit-Learn” (e.g. data-cleaning, preprocessing), but the technical environment for running the code is different and also the code itself is different.

For this reason I added a separate chapter, see sec. 5 which covers the topics “Map-Reduce”¹³ (one of the powerful programming models for Big Data) and “K-Means-Clustering in Spark”.

Let’s start with the “Big Data” structure, which I put into a mind map (you can download it from my GitHub repository). I aligned the structure to the Scikit-Learn mind map (see fig. 4.9) in order to learn from each of this two approaches.

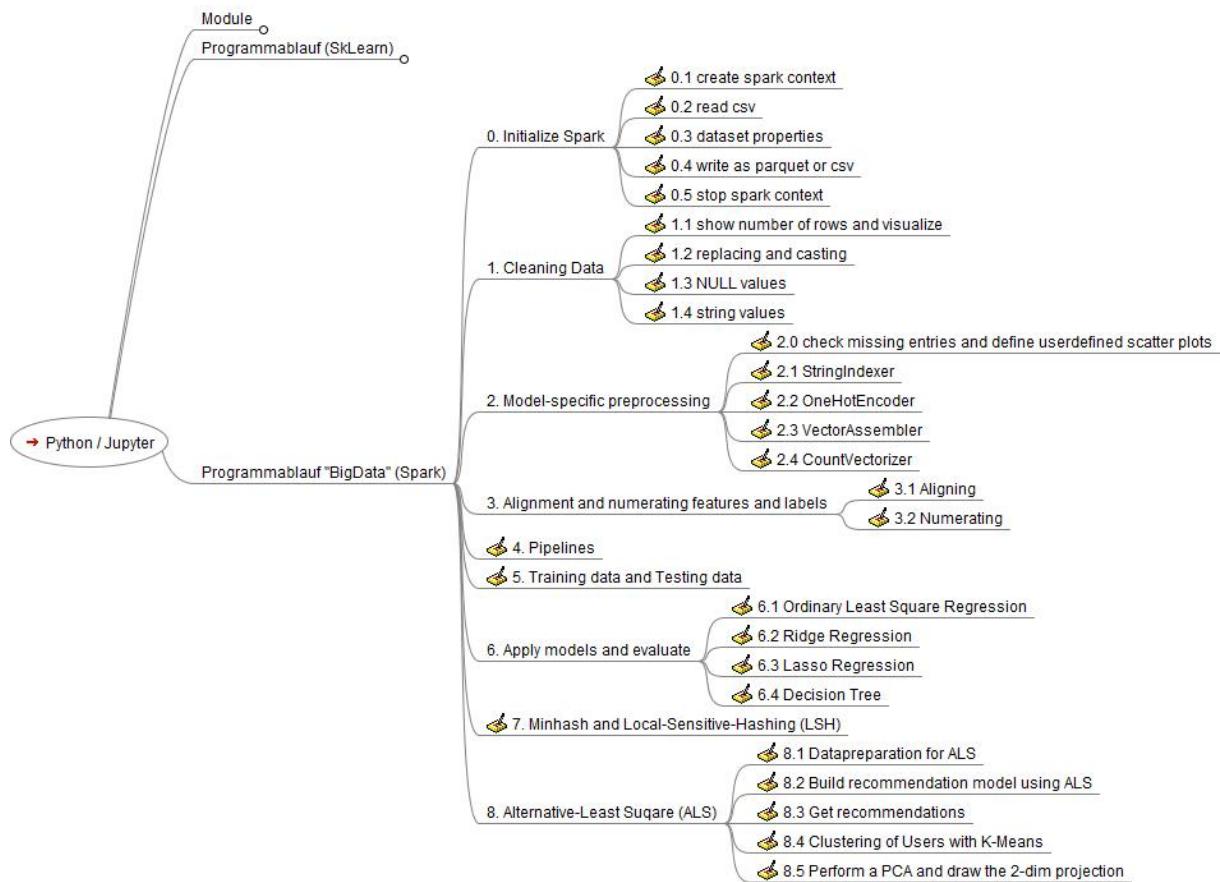


Figure 4.14: Mind Map - Apache Spark ML “Big Data”

4.3.1 Get Ready

There are different ways to get ready and started with the Apache Spark and Hadoop environment:

¹³ Map Reduce, Wikipedia, <https://en.wikipedia.org/wiki/MapReduce>

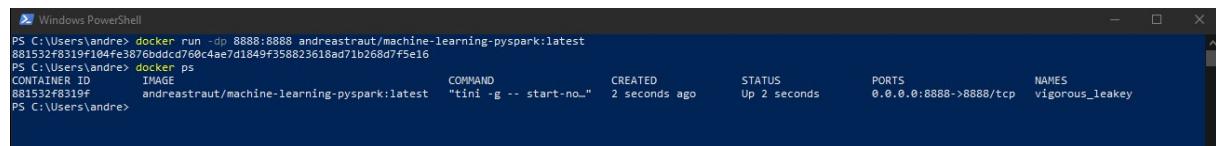
- I guess, that you can install it on your own computer (which I found very difficult because of lack of user-friendly and easy understandable documentation).
- Or you can dive into a Cloud environment, like e.g. Microsoft Azure or Amazon EWS or Google Cloud and try to get a virtual machine up and running for your purposes. Have a look at my documentation¹⁴, where I shared my experiences, which I had with Microsoft Azure: getting started with a Cloud service is not easy!
- Or you can use Docker¹⁵, which I did. What is Docker? Docker is “*an open-source project that automates the deployment of software applications inside containers by providing an additional layer of abstraction and automation of OS-level virtualization on Linux.*” I recommend learning from the Docker-Curriculum¹⁶ what docker container and docker images are and how it works. Despite the difficult subject matter it is a nice tutorial.

For the following explanation I decided to use Docker. I found a container, which had *Apache Spark Version 3.0.0* and *Hadoop 3.2* installed and built my machine-learning code (using pyspark) on top of this container. I shared my code and developments on Docker-Hub in the my docker machine-learning repository¹⁷. After having installed the Docker application you will need to pull my “machine-learning-pyspark” image to your computer. Open a command shell (type cmd on a Windows computer) and enter the following command:

```
docker pull andreastraut/machine-learning-pyspark
```

Then type the following:

```
docker run -dp 8888:8888 andreastraut/machine-learning-pyspark:latest
```



CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
881532f8319f	andreastraut/machine-learning-pyspark:latest	"tini -g -- start-no..."	2 seconds ago	Up 2 seconds	0.0.0.0:8888->8888/tcp	vigorous_leakey

Figure 4.15: Big Data - Run Docker

You will see in your Docker Dashboard that a container is running:

¹⁴ My experiences with Microsoft Azure, <https://github.com/AndreasTraut/Experiences-with-MicrosoftAzure>

¹⁵ Docker, <https://www.docker.com/>

¹⁶ Docker-Curriculum, <https://docker-curriculum.com/>

¹⁷ My docker machine-learning repository, <https://hub.docker.com/repository/docker/andreastraut/machine-learning-pyspark>



Figure 4.16: Big Data - Docker Dashboard

After having opened your browser (e.g. Chrome, Edge or Firefox Browser), navigate to `localhost:8888` (8888 is the port, which will be opened).



Figure 4.17: Big Data - Docker Localhost

The folder “data” contains the datasets. If you would like to do further analysis or produce alternate visualizations of the Airbnb-data, you can download them from Inside-AirBnB¹⁸. It is available below under consideration of the *Creative Commons CC0 1.0 Universal (CC0 1.0) Public Domain Dedication Licence*¹⁹. The data for the Vermont-Vendor-Payments²⁰ is available under condieration of the Open Data Commons Open Database License²¹.

¹⁸ Inside-AirBnB, <http://insideairbnb.com/get-the-data.html>

¹⁹ Creative Commons 1.0 Universal Public Domain Dedication Licence, <http://creativecommons.org/publicdomain/zero/1.0/>

²⁰ “Vermont_Vendor_Payments.csv”, <https://data.vermont.gov/Finance/Vermont-Vendor-Payments/786x-sbp3>

²¹ Open Data Commons Open Database License , <http://opendatacommons.org/licenses/odbl/1.0/>

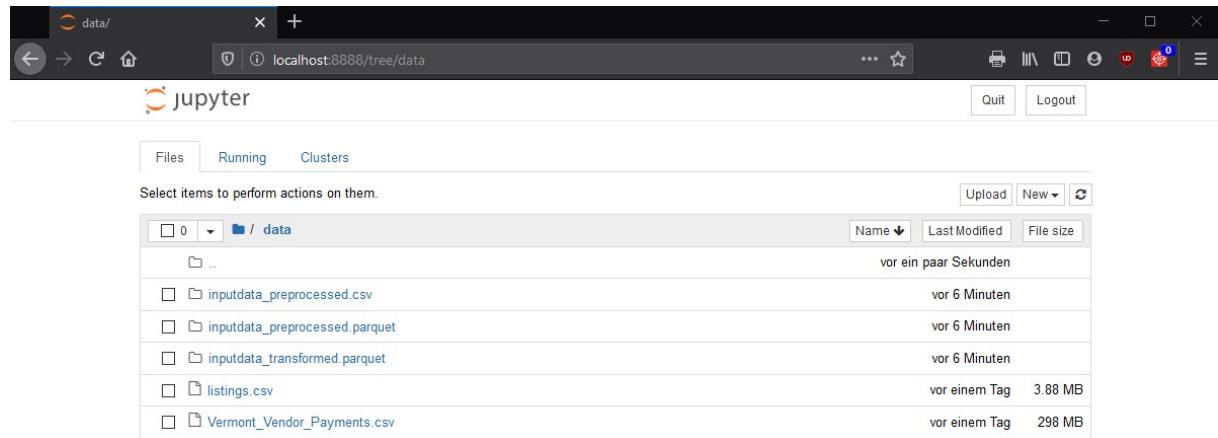


Figure 4.18: Big Data - Docker Localhost Datafiles

When you open the Jupyter-Notebook, you will see, that Apache Spark Version 3.0.0 and Hadoop Version 3.2 is installed:

```
In [1]: import os
print("APACHE_SPARK_VERSION: ", os.environ["APACHE_SPARK_VERSION"])
print("HADOOP_VERSION: ", os.environ["HADOOP_VERSION"])
print(os.environ)

APACHE_SPARK_VERSION: 3.0.0
HADOOP_VERSION: 3.2
environment({'SHELL': '/bin/bash', 'HOSTNAME': 'efcab749826e', 'LANGUAGE': 'en_US.UTF-8', 'SPARK_OPTS': '--driver-java-options=-Xms1024M --driver-java-options=-Xmx4096M --driver-java-options=-Dlog4j.logLevel=intc', 'NB_UID': '1000', 'FWD': '/home/jovyan', 'MINICONDA_MD5': 'd63adf39f2c220950a063e0529d4ff74', 'HOME': '/home/jovyan', 'LANG': 'en_US.UTF-8', 'NB_GID': '100', 'XDG_CACHE_HOME': '/home/jovyan/.cache', 'APACHE_SPARK_VERSION': '3.0.0', 'PYTHONPATH': '/usr/local/spark/python:/usr/local/spark/python/lib/py4j-0.10.9-src.zip', 'HADOOP_HOME': '3.2', 'SHLVL': '0', 'CONDA_DIR': '/opt/conda', 'MINICONDA_VERSION': '4.8.3', 'SPARK_HOME': '/usr/local/spark', 'CONDA_VERSION': '4.8.3', 'NB_USER': 'jovyan', 'LC_ALL': 'en_US.UTF-8', 'PATH': '/opt/conda/bin:/user/local/sbin:/usr/local/bin:/user/sbin:/bin:/usr/local/spark/bin', 'DEBIAN_FRONTEND': 'noninteractive', 'KERNEL_LAUNCH_TIMEOUT': '40', 'JPY_PARENT_PID': '6', 'TERM': 'xterm-color', 'CLICOLOR': '1', 'PAGER': 'cat', 'GIT_PAGER': 'cat', 'MPLBACKEND': 'module://ipykernel.pylab.backend_inline'})
```

```
In [2]: !conda list
# packages in environment at /opt/conda:
#
# Name           Version        Build  Channel
_libgcc_mutex   0.1            conda_forge    conda-forge
_openmp_mutex   4.5            0_gnu       conda-forge
abseil-cpp     20200225.2    hel1b5a44_0  conda-forge
alembic         1.4.2          pyh9f0ad1d_0  conda-forge
arrow-cpp       0.17.1        py38h1234567_5  cpu      conda-forge
```

Figure 4.19: Big Data - Docker Jupyter Notebook

4.3.2 Initialize Spark (0)

I recommend to read the Apache Spark “Get Started Guide”²². Initializing a Spark sessions and read a CSV file (see also the official Spark documentation²³).

```
import pyspark
from pyspark.sql import SparkSession
from pyspark.sql import functions as F
```

4.3.2.1 Create Spark Context and Spark Session (0.1)

```
sc = pyspark.SparkContext(appName='Spark Modelling Context')

spark = SparkSession.builder \
    .appName('Spark Modelling Session') \
    .config('spark.executor.memory','5g') \
    .config('spark.executor.cores','4') \
    .getOrCreate()
```

4.3.2.2 Read CSV (0.2)

```
import os
datapath = os.environ['PWD']
filename = datapath + "/data/listings.csv"
#read in data from csv
data = spark.read.csv(path=filename,
                      sep=',',
                      encoding='utf-8',
                      header=True,
                      inferSchema=True)
data.cache()
```

```
DataFrame[id: string, name: string, host_id: string, host_name: string,
         neighbourhood_group: string, neighbourhood: string, latitude: string,
         longitude: string, room_type: string, price: string, minimum_nights:
         int, number_of_reviews: string, last_review: string, reviews_per_month:
         string, calculated_host_listings_count: double, availability_365: int]
```

²² Apache Spark “Get Started Guide”, <http://spark.apache.org/docs/latest/quick-start.html>

²³ Official Spark documentation, <https://spark.apache.org/docs/latest/sql-getting-started.html#starting-point-sparksession>

4.3.2.3 Dataset Properties and some Select, Group and Aggregate Methods (0.3)

With the following command you will get some properties (like count, mean, stddev, min, max).

```
data.describe().show(truncate=10)
```

Similarly to the following command you can also get the distinct count value of other columns:

```
data.select(F.countDistinct('host_name')).show()
```

```
+-----+
| count(DISTINCT host_name) |
+-----+
|          6483 |
+-----+
```

Grouping works as follows:

```
data.groupBy('host_name').count().sort('count', ascending=False).show(n=5,
↪ truncate=False)
```

```
+-----+
| host_name | count |
+-----+
| Anna      | 207   |
| Julia     | 180   |
| Daniel    | 166   |
| Michael   | 161   |
| David     | 145   |
+-----+
only showing top 5 rows
```

4.3.2.4 Write as Parquet or CSV (0.4)

If you want to persist (=save) your intermediate you can do it as follows:

```
data.select(*data.columns[:-  
    ↵ 1]).write.format("parquet").save("data/inputdata_preprocessed.parquet",  
    ↵ mode='overwrite')  
data.select(*data.columns[:-  
    ↵ 1]).write.csv('data/inputdata_preprocessed.csv', mode='overwrite',  
    ↵ header=True)
```

Reading works as follows:

```
filename = "data/inputdata_preprocessed.parquet"  
data = spark.read.parquet(filename)
```

4.3.2.5 Stop a Spark Session and Spark Context (0.5)

```
spark.stop()  
sc.stop()
```

4.3.3 Cleaning the data (1)

The data-cleaning and data preparation (eliminating of null values, visualization techniques) the structure is pretty similar to the “Small data” (Sklearn) approach.

4.3.3.1 Show Number of Rows and Columns and Do Some Visualizations (1.1)

```
data.printSchema()
```

```
root  
|-- id: string (nullable = true)  
|-- name: string (nullable = true)  
|-- host_id: string (nullable = true)  
|-- host_name: string (nullable = true)  
|-- neighbourhood_group: string (nullable = true)  
|-- neighbourhood: string (nullable = true)  
|-- latitude: string (nullable = true)  
|-- longitude: string (nullable = true)  
|-- room_type: string (nullable = true)
```

```
|-- price: string (nullable = true)
|-- minimum_nights: integer (nullable = true)
|-- number_of_reviews: string (nullable = true)
|-- last_review: string (nullable = true)
|-- reviews_per_month: string (nullable = true)
|-- calculated_host_listings_count: double (nullable = true)
```

```
data.describe().show(truncate=False, n=1, vertical=True)
```

```
-RECORD 0-----
summary           | count
id               | 25206
name              | 25146
host_id           | 25163
host_name         | 25141
neighbourhood_group | 25163
neighbourhood     | 25163
latitude          | 25163
longitude          | 25163
room_type          | 25163
price              | 25163
minimum_nights      | 25163
number_of_reviews    | 25158
last_review         | 20690
reviews_per_month    | 20641
calculated_host_listings_count | 25109
only showing top 1 row
```

```
data = data.withColumn('price',
                      data['price'].cast('double'))
data.sample(fraction=0.3,
            seed=42,
            withReplacement=False) \
    .sort('price', ascending=False) \
    .select('price').toPandas().plot.box();
```

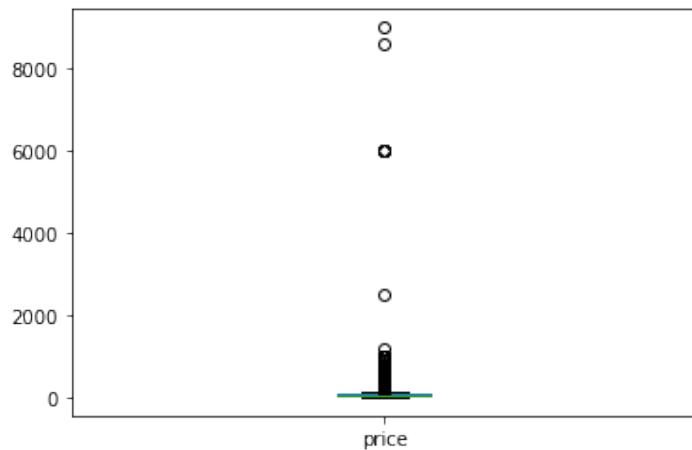


Figure 4.20: Big Data - Price Box Plot

4.3.3.2 Replacing and Casting (1.2)

```
data = data.withColumn('price', F regexp_replace('price', '\$', ''))  
data = data.withColumn('price', F regexp_replace('price', ',', ','))  
data = data.withColumn('price', data['price'].cast('double'))
```

4.3.3.3 Null-Values (1.3)

```
print("{} missing values for price"  
      .format(data  
              .filter(F.isnull(data['price']))  
              .count()))
```

104 missing values for price

```
data = data.fillna(0, subset='price')
```

```
print("{} missing values for  
→ price".format(data.filter(F.isnull(data['price'])).count()))
```

```
0 missing values for price
```

4.3.3.4 String Values (1.4)

```
string_types = [x[0]
                 for x in data.dtypes
                 if x[1] == 'string']
data.select(string_types).describe().show()
```

4.3.4 Model-specific preprocessing (2)

4.3.4.1 Check Missing Entries and Define Userdefined Scatter Plot (2.0)

```
data.select(*(F.sum(F.col(c).isNull() \
                     .cast("int")).alias(c) for c in data.columns)) \
    .toPandas() \
    .transpose()
import numpy as np
import matplotlib.pyplot as plt

def scatter_predicted_vs_actual_price(model, data):
    prediction = np.array([p[0] for p in
                           model.transform(data).select('prediction').collect()])
    truth = np.array([p[0] for p in
                     model.transform(data).select('label').collect()])
    fig, ax = plt.subplots(figsize=(5, 5))
    ax.scatter(prediction, truth, facecolor='steelblue', s=30, alpha=0.6)
    ax.set_xlabel('Predicted price', fontsize=12)
    ax.set_ylabel('Actual price', fontsize=12)
    m = max(max(prediction), max(truth))
    ax.plot([0,m], [0,m], color='grey')
    ax.set_xlim([0,m])
    ax.set_ylim([0,m])
    plt.show()
```

4.3.4.2 StringIndexer (2.1)

Please read the official documentation in order to understand, what the “StringIndexer”²⁴ does.

```
from pyspark.ml.feature import StringIndexer
room_indexer = StringIndexer(inputCol='room_type', outputCol='room_index')
room_indexer_model = room_indexer.fit(data)
data = room_indexer_model.transform(data)

data.groupby('room_type').agg(F.collect_set('room_index').alias('room_index')).sort('room_index', ascending=False).show(4)
```

room_type	room_index
Shared room	[2.0]
Private room	[1.0]
Hotel room	[3.0]
Entire home/apt	[0.0]

only showing top 4 rows

4.3.4.3 OneHotEncoder (2.2)

Please read the official documentation in order to understand, what the “OneHotEncoder”²⁵ does.

```
from pyspark.ml.feature import OneHotEncoder
one_hot_encoder = OneHotEncoder(inputCol = 'neighbourhour_group_index',
                                outputCol = 'one_hot_neighbourhood_group', dropLast = False)
one_hot_encoder_model = one_hot_encoder.fit(data)
data = one_hot_encoder_model.transform(data)
```

4.3.4.4 VectorAssembler (2.3)

Please read the official documentation in order to understand, what the “VectorAssembler”²⁶ does.

²⁴ “StringIndexer”, <https://spark.apache.org/docs/latest/ml-features.html#stringindexer>

²⁵ “OneHotEncoder”, <https://spark.apache.org/docs/latest/ml-features.html#onehotencoder>

²⁶ “VectorAssembler”, <https://spark.apache.org/docs/latest/ml-features.html#vectorassembler>

```
from pyspark.ml.feature import VectorAssembler
data = data.withColumn('number_of_reviews',
                       data['number_of_reviews'].cast('double'))
data.select('number_of_reviews').show()
numeric_attributes = ['number_of_reviews']
vec_num = VectorAssembler(inputCols=numeric_attributes,
                           outputCol='num_features')
data = vec_num.transform(data)
vec_label = VectorAssembler(inputCols = ['price'],
                            outputCol = 'vec_label')
data = vec_label.transform(data)
```

number_of_reviews
145.0
27.0
133.0
292.0
8.0
24.0
48.0
262.0
86.0
60.0
86.0
307.0
130.0
21.0
5.0
188.0
31.0
74.0
296.0
39.0

4.3.4.5 CountVectorizer (2.4)

Please read the official documentation in order to understand what the “CountVectorizer”²⁷ does.

²⁷ “CountVectorizer”, <https://spark.apache.org/docs/latest/ml-features.html#countvectorizer>

4.3.5 Aligning and numerating Features and Labels (3)

4.3.5.1 Aligning (3.1)

Aligning in this context means, that you use the variable `label` for the labeled column (here it is the “price” column) and the variable `feature_cols` for the features.

```
label = 'price'
feature_cols = ['num_features',
                 'vec_label']
cols = feature_cols + ['label']
data_feat = data.withColumnRenamed(label, 'label').select(cols)
```

```
+-----+-----+
| num_features | vec_label | label |
+-----+-----+
| [145.0] | [90.0] | 90.0 |
| [27.0] | [28.0] | 28.0 |
| [133.0] | [125.0] | 125.0 |
| [292.0] | [33.0] | 33.0 |
| [8.0] | [180.0] | 180.0 |
+-----+-----+
only showing top 5 rows
```

4.3.5.2 Numerating (3.2)

```
feat_len = len(numeric_attributes) +
    data.select('room_type', 'room_index').distinct().count()
features = numeric_attributes + [r[0] for r in
    data.select('room_type', 'room_index').distinct().collect()]
feature_dict = dict(zip(range(0, feat_len), features))
```

4.3.6 Pipelines (4)

```
from pyspark.ml import Pipeline
pipeline = Pipeline(stages=[vec_num,
                           vec_label])
pipeline_model = pipeline.fit(data)
```

If you want to work with another pipeline, then all you need is to replace the second line. For example into this:

```
pipeline = Pipeline(stages=[room_indexer,  
                            vec_num,  
                            vec_label])
```

4.3.7 Training data and Testing data (5)

```
data_train, data_test = data_feat.randomSplit([0.9,0.1],  
                                             seed=42)
```

This creates a random split into a training dataset `data_train` and the testing dataset `data_test`. Remember, that in the movies data example (see sec. 4.1) we used the stratified sample.

4.3.8 Apply models and evaluate (6)

4.3.8.1 Ordinary Least Square Regression (6.1)

After having extracted, transformed and selected features we will apply some models, for example the “OLS Regression”²⁸. Please read the official documentation for learning more about OLS.

```
from pyspark.ml.regression import LinearRegression  
lr = LinearRegression(featuresCol='num_features',  
                      labelCol='label',  
                      maxIter=1000,  
                      fitIntercept=True)  
lr_model = lr.fit(data_train)  
  
lr_model.coefficients  
pred = lr_model.transform(data_test)  
pred.select('label','prediction').show(5)
```

Evaluate the results:

²⁸ “OLS Regression”, <https://spark.apache.org/docs/latest/ml-classification-regression.html#linear-regression>

```
from pyspark.ml.evaluation import RegressionEvaluator
re = RegressionEvaluator(metricName='rmse')
rmse = re.evaluate(pred)
print(rmse)
scatter_predicted_vs_actual_price(lr_model,
                                   data_test)
```

4.3.8.2 Ridge Regression (6.2) and Lasso Regression

The coding syntax for the “Ridge Regression” and the “Lasso Regression” are pretty similar (there are also the `fit` and the `transform` methods). Please read the official documentation in order to understanding the “Ridge Regression and the Lasso Regression”²⁹ in detail.

4.3.8.3 Decision Tree (6.4)

The coding syntax for the “Decision Tree” is pretty similar (there is also the `fit` and the `transform` method). Please read the official documentation in order to understanding the “Decision Tree”³⁰ in detail.

4.3.9 Minhash und Local-Sensitive-Hashing (7)

I worked on the Minhashing and on the Local-Sensitive Hasing in a separate example³¹.

4.4 Summary Mind-Map

To summarize the whole coding structure have a look at the following and also the provided mindmaps. My mind map below may help you to structure your code:

²⁹ Ridge Regression and the Lasso Regression, <https://spark.apache.org/docs/latest/mllib-linear-methods.html#linear-least-squares-lasso-and-ridge-regression>

³⁰ Decision Tree, <https://spark.apache.org/docs/latest/ml-classification-regression.html#decision-trees>

³¹ Minhashing and Local-Sensitive Hashing Example, https://github.com/AndreasTraut/Deep_learning_explorations

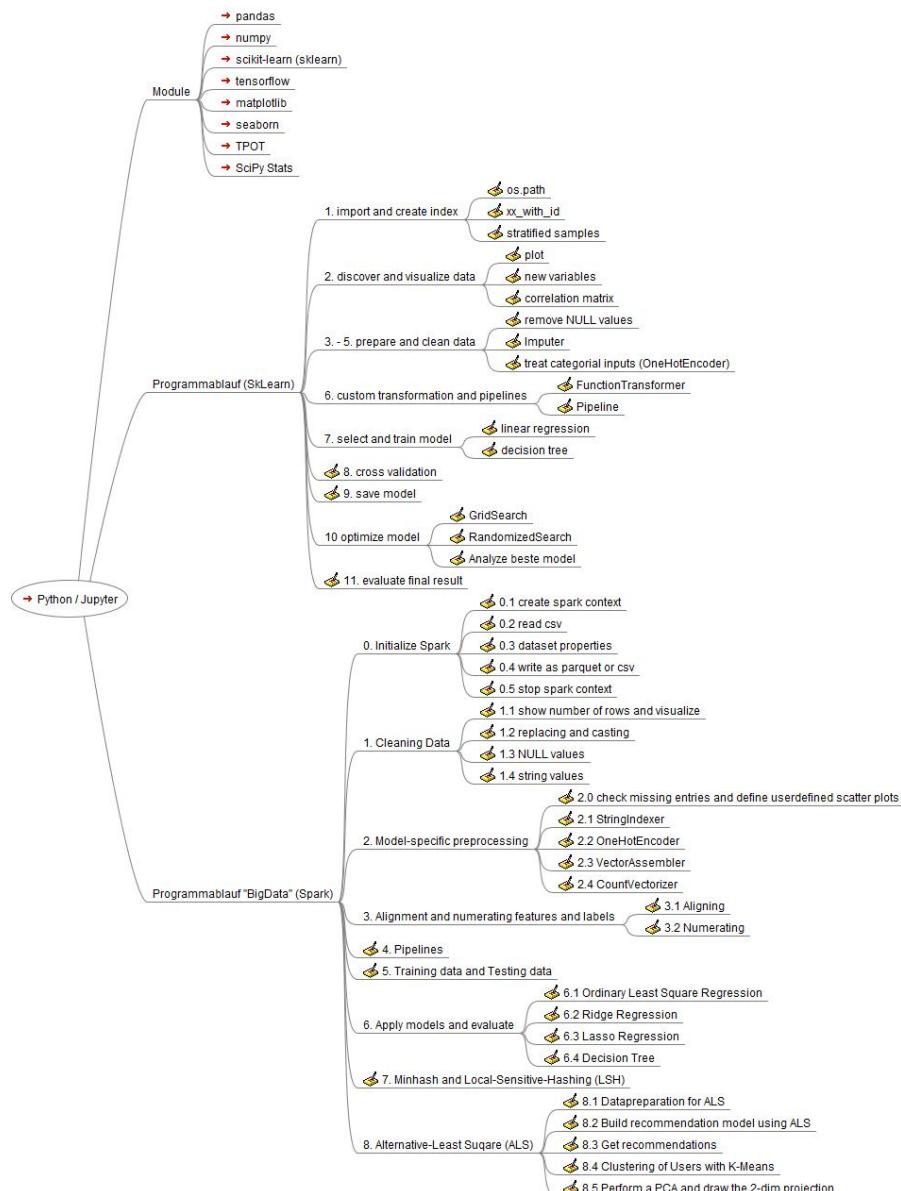


Figure 4.21: Mind Map - Scikit-Learn and Apache SparkML

4.5 Future learnings and coding & data sources

For all of these topics various tutorials, documentation, coding examples and guidelines can be found in the internet **for free!** The Open Source Community is an incredible treasure trove and enrichment that positively drives many digital developments: Scikit-Learn, Apache Spark, Spyder, GitHub, Ten-

sorflow and also Firefox³², Signal-Messenger³³, Threema-Messenger³⁴, Corona-Warnapp³⁵, ... to be mentioned. There are many positive examples of sharing code and data “for free”.

Coding:

If you Google for example “*how to prepare and clean the data with spark*”, you will find tons of documents around “*removing null values*” or “*encoders*” (like the “OneHotEncoder” for treating categorical inputs) or “*Pipelines*” (for putting all the steps in an efficient, customizable order) so on. You will be overwhelmed of all this. Some resources to mention are the official Apache Spark ML documentation and a few more Github repositories like the ones, which I mentioned already (tirthajyoti/Spark-with-Python (MIT Licence), Apress/learn-pyspark (Freeware License), mahmoudparsian/pyspark-tutorial (Apache License v2.0)).

Data:

If you would like to do further analysis or produce alternate visualizations of the Airbnb-data, you can download them under consideration of the *Creative Commons 1.0 Universal Public Domain Dedication Licence*³⁶. The data for the Vermont-Vendor-Payments can be downloaded under consideration of the *Open Data Commons Open Database License*³⁷. The movies database doesn’t even mention a license. There you find a lot of more datasets and also coding examples for your studies.

³² Firefox, <https://github.com/mozilla>

³³ Signal-Messenger, <https://github.com/signalapp>

³⁴ Threema-Messenger, <https://github.com/threema-ch>

³⁵ Corona-Warnapp, <https://github.com/corona-warn-app>

³⁶ Creative Commons 1.0 Universal Public Domain Dedication Licence, <http://creativecommons.org/publicdomain/zero/1.0/>

³⁷ Open Data Commons Open Database License , <http://opendatacommons.org/licenses/odbl/1.0/>

5 Big Data: Map-Reduce and K-Means Clustering

5.1 Map-Reduce Example

Map-Reduce is a programming model for generating big data sets with parallel distributed algorithm on a cluster. This is very important for Big Data and therefore I added some examples in order to explain how Map-Reduce works: the first is the “Word Count Example” (see sec. 5.1.1) for learning the basics of Map-Reduce and the second is the “TF-IDF Example” (sec. 5.1.2) for applying Map-Reduce on an interesting text-mining problem.

Please learn the basis of the Map-Reduce programming model for example from Wikipedia¹ before continuing reading.

5.1.1 Word Count Example

The “Word Count” example is one of the easiest for explaining Map-Reduce: given a text as input the aim is to return a list with words and the number of occurrences of each of these words. I used the “Moby Dick” as input text:

First we start a spark session and open the text file:

```
from pyspark import SparkContext
sc=SparkContext(master="local[4]")
text_file = sc.textFile("data/Moby-Dick.txt")
type(text_file)

pyspark.rdd.RDD
```

The text is now available in `text_file`, which is a resilient distributed dataset (RDD)². RDD is a fault-tolerant collection of elements (here the lines of words), that can be operated on in parallel.

¹ Map Reduce, Wikipedia, <https://en.wikipedia.org/wiki/MapReduce>

² Resilient Distributed Dataset (RDD), <https://spark.apache.org/docs/latest/rdd-programming-guide.html#resilient-distributed-datasets-rdds>

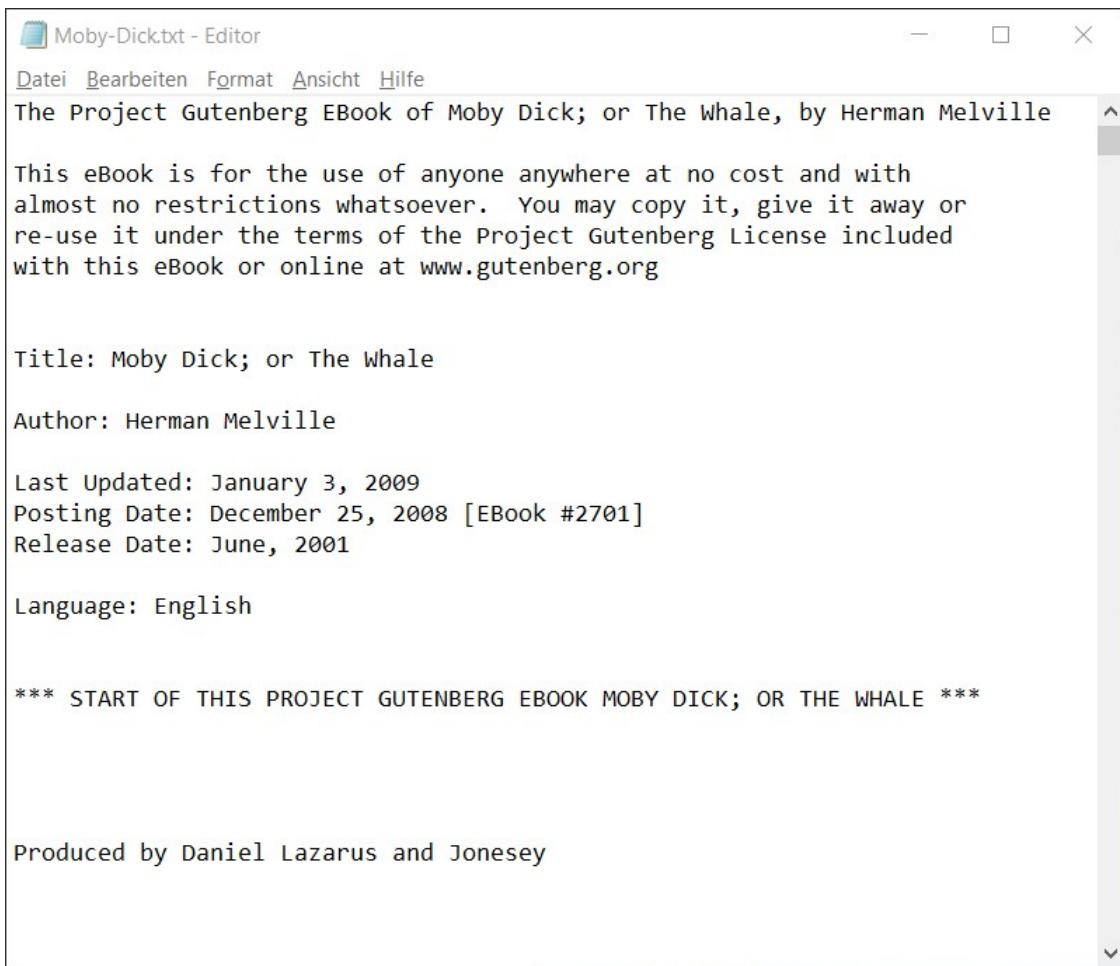


Figure 5.1: Map-Reduce - Word Count Example “Moby Dick”

Let's have a look into the `text_file`:

```
print(text_file.collect())
```

```
[ 'The Project Gutenberg EBook of Moby Dick; or The Whale, by Herman
  ↵ Melville', '',
  'This eBook is for the use of anyone anywhere at no cost
  ↵ and with', 'almost no restrictions whatsoever. You may copy it, give it
  ↵ away or', 're-use it under the terms of the Project Gutenberg License
  ↵ included', 'with this eBook or online at www.gutenberg.org', '',
  'Title: Moby Dick; or The Whale', '',
  'Author: Herman Melville', '',
  'Last Updated: January 3, 2009', 'Posting Date: December 25, 2008 [EBook
  ↵ #2701]', 'Release Date: June, 2001', '',
  'Language: English', '' ]
```

As we can see each line is separated with a comma, as follows: ‘line1’, ‘line2’, ‘line3’, ...

In the next steps we will

1. split the lines by using the spaces as separators
2. eliminate the empty elements
3. **map** the words to a tuple (word, 1)
4. **reduce** by key in order to count the number of occurrences of each word (which is a, b will be transformed to a + b)

The python code is as follows:

```
#1:  
words = text_file.flatMap(lambda line: line.split(" "))  
  
#2:  
not_empty = words.filter(lambda x: x != '')  
  
#3:  
key_values = not_empty.map(lambda word: (word, 1))  
  
#4:  
counts = key_values.reduceByKey(lambda a, b: a + b)
```

I will explain each of these four steps in more detail now:

5.1.1.1 Split the lines by using the spaces as separators (#1)

With the following line of code each line is split up into the words:

```
words = text_file.flatMap(lambda line: line.split(" "))
```

As a word is separated from another word with a blank, we have to use " " as a separator.

```
print(words.take(35))
```

```
['The', 'Project', 'Gutenberg', 'EBook', 'of', 'Moby', 'Dick;', 'or', 'The',  
↳ 'Whale,', 'by', 'Herman', 'Melville', '', 'This', 'eBook', 'is', 'for',  
↳ 'the', 'use', 'of', 'anyone', 'anywhere', 'at', 'no', 'cost', 'and',  
↳ 'with', 'almost', 'no', 'restrictions', 'whatsoever.', '', 'You']
```

We can find some empty entries . . . , ' ', '' , . . . which we will filter out in the next step. But before we do this we want to understand why flatMap has been used here instead of map:

The reason is that the operation `line.split(" ")` generates a **list** of strings, so had we used `map` the result would be an RDD of lists of words. Not an RDD of words.

The difference between `map` and `flatMap` is that the second expects to get a list as the result from the map and it **concatenates** the lists to form the RDD.

5.1.1.2 Eliminate the empty elements (#2)

The next line of code is the following:

```
not_empty = words.filter(lambda x: x != '')
```

This will filter in order to get rid of all empty entries.

```
print(not_empty.take(35))
```

```
['The', 'Project', 'Gutenberg', 'EBook', 'of', 'Moby', 'Dick;', 'or', 'The',
 ← 'Whale,', 'by', 'Herman', 'Melville', 'This', 'eBook', 'is', 'for',
 ← 'the', 'use', 'of', 'anyone', 'anywhere', 'at', 'no', 'cost', 'and',
 ← 'with', 'almost', 'no', 'restrictions', 'whatsoever.', 'You', 'may',
 ← 'copy']
```

As we can see the empty entries are gone (we had two empty entries in the `words.take(35)` above).

5.1.1.3 Map the words to a tuple (#3)

In the next line of code we will apply the `map`:

```
key_values = not_empty.map(lambda word: (word, 1))
```

This means, that each `word` will be mapped to `(word, 1)`. We need this in order to be able to count the words.

```
print(key_values.take(35))
```

```
[('The', 1), ('Project', 1), ('Gutenberg', 1), ('EBook', 1), ('of', 1),
←  ('Moby', 1), ('Dick;', 1), ('or', 1), ('The', 1), ('Whale,', 1), ('by',
←  1), ('Herman', 1), ('Melville', 1), ('This', 1), ('eBook', 1), ('is',
←  1), ('for', 1), ('the', 1), ('use', 1), ('of', 1), ('anyone', 1),
←  ('anywhere', 1), ('at', 1), ('no', 1), ('cost', 1), ('and', 1), ('with',
←  1), ('almost', 1), ('no', 1), ('restrictions', 1), ('whatsoever.', 1),
←  ('You', 1), ('may', 1), ('copy', 1)]
```

5.1.1.4 Reduce by key (#4)

Reduce by key in order to count the number of occurrences of each word (which is a, b will be transformed to a + b).

```
counts = key_values.reduceByKey(lambda a, b: a + b)
```

Have a look at the #3 above and try to find the word 'The':

```
('The', 1), ('Project', 1), ('Gutenberg', 1), ('EBook', 1), ('of', 1), ('Moby', 1), ('Dick;', 1), ('or', 1), ('The', 1),
('Whale,', 1), ('by', 1), ('Herman', 1), ('Melville', 1), ('This', 1), ('eBook', 1), ('is', 1), ('for', 1), ('the', 1), ('u
se', 1), ('of', 1), ('anyone', 1), ('anywhere', 1), ('at', 1), ('no', 1), ('cost', 1), ('and', 1), ('with', 1), ('almost',
1), ('no', 1), ('restrictions', 1), ('whatsoever.', 1), ('You', 1), ('may', 1), ('copy', 1), ('it', 1), ('give', 1), ('it
', 1), ('away', 1), ('or', 1), ('re-use', 1), ('it', 1), ('under', 1), ('the', 1), ('terms', 1), ('of', 1), ('the', 1), ('P
roject', 1), ('Gutenberg', 1), ('License', 1), ('included', 1), ('with', 1), ('this', 1), ('eBook', 1), ('or', 1), ('online
', 1), ('at', 1), ('www.gutenberg.org', 1), ('Title:', 1), ('Moby', 1), ('Dick;', 1), ('or', 1), ('The', 1), ('Whale', 1),
('Author:', 1), ('Herman', 1), ('Melville', 1), ('Last', 1), ('Updated:', 1), ('January', 1), ('3,', 1), ('2009', 1), ('Pos
ting', 1), ('Date:', 1), ('December', 1), ('25,', 1), ('2008', 1), ('[EBook', 1), ('#2701]', 1), ('Release', 1), ('Date:',
1), ('June,', 1), ('2001', 1), ('Language:', 1), ('English', 1), ('***', 1), ('START', 1), ('OF', 1), ('THIS', 1), ('PROJEC
T', 1), ('GUTENBERG', 1), ('EBOOK', 1), ('MOBY', 1), ('DICK;', 1), ('OR', 1), ('THE', 1), ('WHALE', 1), ('***', 1), ('Produ
ced', 1), ('by', 1), ('Daniel', 1)]
```

Figure 5.2: Map-Reduce - Word Count Example Reduce by Key

The aim is to count these. Please note that the word 'the' and 'THE' will not be part of this count as they are different. If we wanted to have them aggregated all together we would have needed another map step, which changes every word into its capital cases.

```
[('The', 549), ('Project', 79), ('EBook', 1), ('of', 6587), ('Moby', 79),
←  ('is', 1586), ('use', 35), ('anyone', 5), ('anywhere', 11), ('at',
←  1227), ('no', 447), ('restrictions', 2), ('whatsoever.', 5), ('may',
←  223), ('it,', 237), ('give', 68), ('away', 117), ('re-use', 2), ('this',
←  1169), ('online', 4), ('www.gutenberg.org', 2), ('Author:', 1), ('Last',
←  1), ('January', 1), ('3,', 2), ('Posting', 1), ('Date:', 2), ('#2701]',
←  1), ('June,', 3), ('Language:', 1), ('English', 42), ('***', 6), ('OF',
←  59), ('THIS', 12), ('GUTENBERG', 3), ('MOBY', 3), ('DICK;', 3)]
```

5.1.1.5 Sort by keys and examine top 10 words

Interesting to have would be the top 10 words:

```
countsSort = counts.sortBy(lambda a: a[1], ascending=False)
countsSort.take(10)
```

```
[('the', 13766),
 ('of', 6587),
 ('and', 5951),
 ('a', 4533),
 ('to', 4510),
 ('in', 3878),
 ('that', 2693),
 ('his', 2415),
 ('I', 1724),
 ('with', 1692)]
```

Having words like ‘the’, ‘of’, ‘and’.. in the top 10 of a text is not surprising. Here is how we can look up the word ‘The’ and ‘THE’: `countsSort.lookup('The')` results in 549 and `countsSort.lookup('THE')` results in 98. It would make sense to aggregate the counts for ‘the’ and ‘The’ and ‘THE’. In this case we would expect have 14413:

```
countsSort.lookup('the')[0] + countsSort.lookup('The')[0] +
    countsSort.lookup('THE')[0]
```

After having learnt this, we might think to restart the whole counting steps and include the UPPERCASE. We will do this in #3.

5.1.1.6 Steps for counting the words using UPPERCASE

```
#1:
words = text_file.flatMap(lambda line: line.split(" "))

#2:
not_empty = words.filter(lambda x: x != '')

#3:
tmp = not_empty.map(lambda word: word.upper())      # UPPERCASE
```

```
key_values_upper = tmp.map(lambda word: (word, 1))

#4:
counts_upper = key_values_upper.reduceByKey(lambda a, b: a + b)
```

Have a look into the results:

```
counts_upperSort = counts_upper.sortBy(lambda a: a[1], ascending=False)
counts_upperSort.take(20)
```

```
[('THE', 14413),
 ('OF', 6668),
 ('AND', 6309),
 ('A', 4658),
 ('TO', 4595),
 ('IN', 4115),
 ('THAT', 2759),
 ('HIS', 2485),
 ('IT', 1776),
 ('WITH', 1750),
 ('I', 1724),
 ('AS', 1713),
 ('HE', 1683),
 ('BUT', 1672),
 ('IS', 1605),
 ('WAS', 1577),
 ('FOR', 1557),
 ('ALL', 1359),
 ('AT', 1312),
 ('THIS', 1283)]
```

On my Docker machine-learning repository³ you will find the Jupyter-Notebook, with this example.

5.1.2 Term Frequency–Inverse Document Frequency (TF-idf) Example

In this example I will expand the Map-Reduce from sec. 5.1. The “Term Frequency–Inverse Document Frequency (TF-idf)”⁴ is a very interesting measure. It shows, how important a term is to a document.

³ My docker machine-learning repository, <https://hub.docker.com/repository/docker/andreastraut/machine-learning-pyspark>

⁴ Term frequency - Inverse Document Frequency, <https://en.wikipedia.org/wiki/Tf%E2%80%93idf>

The TF-idf is a numerical statistic, which is often used in text-based recommender systems and for information retrieval⁵ or text mining or to classify texts.

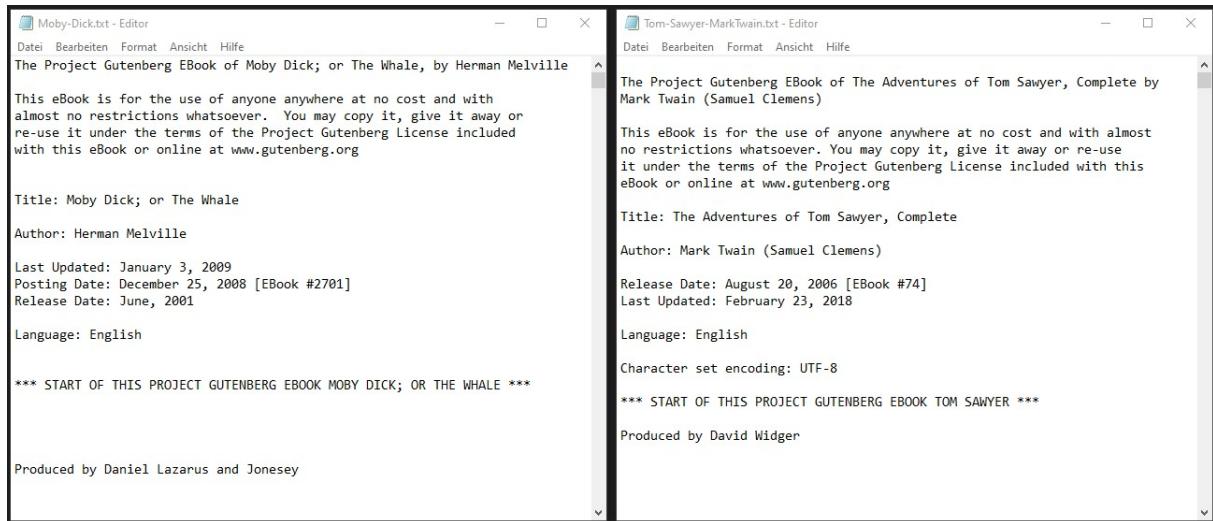


Figure 5.3: TF-idf example - Moby Dick and Tom Sawyer

In my example I used the texts of “Moby Dick” and “Tom Sawyer” as input files. By applying the TF-idf measure, the result are two lists of most important words for each of these documents. This is what the TF-idf found for each of the two documents:

Moby Dick:

WHALE, AHAB, WHALES, SPERM, STUBB, QUEEQUEG, STRARBUCK, AYE

Tom Sawyer:

HUCK, TOMS, BECKY, SID, INJUN, POLLY, POTTER, THATCHER

These list of words are (according to the TF-idf measure) the most important ones for each of the texts and they are also the words, which are the least important ones for the other document. If you have read “Moby Dick”, then you would agree, that WHALE, AHAB, WHALES, SPERM, STUBB, QUEEQUEG, STRARBUCK, AYE are very important words for the “Moby Dick” text. And if you have read “Tom Sawyer”, then you would also agree, that HUCK, TOMS, BECKY, SID, INJUN, POLLY, POTTER, THATCHER are very important words for the “Tom Sawyer” text. And if you have read both of these texts, then you would agree, that neither of these words are relevant for the opposite text: e.g. in the “Moby Dick” text the name HUCK does not appear and in the “Tom Sawyer” text there is no WHALE.

I will explain now, how this works by applying the “Map-Reduce” approach as described in sec. 5.1.

⁵ Information Retrieval, https://en.wikipedia.org/wiki/Information_retrieval

5.1.2.1 Count the words

First I start a Spark Session and read the text:

```
document1 = "data/Moby-Dick.txt"
text_file = sc.textFile(document1)
```

Next is to clean the data by eliminating quotation marks, commas and so on (",.") which is done in #3:

```
#1:
words = text_file.flatMap(lambda line: line.split(" "))

#2:
not_empty = words.filter(lambda x: x != '')

#3:
tmpA = not_empty.map(lambda word: word.upper()) # UPPERCASE
tmpB = tmpA.map(lambda word:
                 word.replace('"', '')) # eliminate the "
tmpC = tmpB.map(lambda word:
                 word.replace("'", '')) # eliminate the '
tmpD = tmpC.map(lambda word:
                 word.replace(',', '')) # eliminate the ,
tmpE = tmpD.map(lambda word:
                 word.replace('.', '')) # eliminate the .
tmpF = tmpE.map(lambda word:
                 word.replace("\\"", "")) # eliminate the "
tmpG = tmpF.map(lambda word:
                 word.replace("?", '')) # eliminate the ?

tmp_char = tmpG.map(lambda word:
                     word.replace("'", '')) # eliminate the '

key_values_upper = tmp_char.map(lambda word: (word, 1))

#4:
counts_upper = key_values_upper.reduceByKey(lambda a, b: a + b)
```

I know that the `replace(...)` in #3 is not very elegant. Better would have been to work with a regular expression. But I didn't want to over-complicate the code here. I think, what I did here is easy to understand. I do the same as above for the other text:

```
document2 = "data/Tom-Sawyer-MarkTwain.txt"
text_file2 = sc.textFile(document2)
```

At the end I have counts_upper2 for the second text.

5.1.2.2 Calculate TF

If we want to calculate the **TF-idf** measure (see [Wikipedia](#)) we would need to do calculate **TF** (Term-Frequency) and **idf** (inverse-document-frequency) separately and then divide them: **TF / idf**. So first I need to calculate **TF**:

```
#Number of words for each documents:
n1 = key_values_upper.count()
n2 = key_values_upper2.count()
print("Document 1: {}".format(document1))
print("Total number of words: {}\\n".format(n1))
print("Document 2: {}".format(document2))
print("Total number of words: {}\\n".format(n2))
```

Document 1: data/Moby-Dick.txt

Total number of words: 215133

Document 2: data/Tom-Sawyer-MarkTwain.txt

Total number of words: 73840

Therefore the **TF** for terms of document 1 are as follows:

```
TF1 = counts_upperSort.mapValues(lambda value: value/n1)
TF1.take(7)
```

```
[('THE', 0.06733973867328583),
 ('OF', 0.031120283731459145),
 ('AND', 0.029888487586748662),
 ('A', 0.021772577893675076),
 ('TO', 0.02159594297481093),
 ('IN', 0.019443785937071485),
 ('THAT', 0.013586943890523536)]
```

And the **TF** for terms of document 2 are as follows:

```
TF2 = counts_upperSort2.mapValues(lambda value: value/n2)
TF2.take(7)
```

```
[('THE', 0.05310130010834236),
 ('AND', 0.04156283856988082),
 ('A', 0.025338569880823402),
 ('TO', 0.02421451787648971),
 ('OF', 0.021424702058504875),
 ('HE', 0.015912784398699892),
 ('WAS', 0.015723185265438786)]
```

5.1.2.3 Calculate idf

As we have $N=2$ documents the **idf** is defined as $\log(N/n) = \log(2/n)$. For calculating the **idf** we need to know if a term (e.g. the word 'THE') occurs in:

- both documents: then $n=2$ and the **idf**= $\log(1)=0$
- only one of the documents: then $n=1$ and the **idf**= $\log(2)$
- in none of the documents: then $n=0$ and the **idf**= $\log(2/0)$ which is not defined (division by 0).

We conclude that when having $N=2$ documents we are only interested in the case, where a term occurs in only **one** document ($n=1$). I define a idf-funktion for this:

```
import math
def idf(term):
    N = 0
    if (counts_upperSort.lookup(term)):
        N += 1
    if (counts_upperSort2.lookup(term)):
        N += 1
    if (N==2):
        print("Term: {}\\nIn both documents.\\nidf('{}') = {:.2f}""
              .format(term, math.log10(2/N)))
    elif (N==1):
        print("Term: {}\\nOnly in one of the documents.\\nidf('{}') = {:.2f}""
              .format(term, math.log10(2/N)))
    elif (N==0):
        print("Term: {}\\nDoes not appear.\\nidf('{}') = not defined"
              .format(term))
```

Now I want to test the idf-function on the word “THE”

```
idf('THE')
```

```
Term: THE
In both documents.
idf('THE') = 0.00
```

The idf for ‘THE’ is 0.00 which means, that it is not very informative (it appears in both documents).

```
idf('MOBY')
```

```
Term: MOBY
Only in one of the documents.
idf('MOBY') = 0.30
```

And a final test on an arbitrary word: `idf('BNWKDMGJ')`

```
Term: BNWKDMGJ
Does not appear.
idf('BNWKDMGJ') = not defined
```

5.1.2.4 Calculate TF-idf

Now I want to put it together and calculate the TF-idf measure. First I will create the tuple ($x[0]$, $x[1]$), where $x[0]$ are the **TF** for each of the terms of document 1 and $x[1]$ are the **TF** for each of the terms of document 2.

If ($x[0]$, $x[1]$) has values for $x[0]$ and $x[1]$ then the term occurs in both documents.

If ($x[0]$, None) then the term occurs only in document 1.

If (None, $x[0]$) then the term occurs only in document 2.

```
both_tf = TF1.fullOuterJoin(TF2)
both_tf.take(7)
```

```
[('TO', (0.02159594297481093, 0.02421451787648971)),
 ('THAT', (0.013586943890523536, 0.012080173347778982)),
 ('AS', (0.008013647371625925, 0.005444203683640303)),
```

```
('ONE', (0.004062603133875324, 0.002397074756229686)),  
('WERE', (0.00313759395350783, 0.0039003250270855903)),  
('THEY', (0.003058573068752818, 0.007367280606717226)),  
('THEIR', (0.0028772898625501436, 0.002627302275189599))]
```

Now calculate the TF-idf for both documents (I will leave out some words, because the result would be too long to show here):

```
both_tfidf = both_tf.mapValues(lambda x:  
    (x[0]*math.log10(2/1),"doc 1")  
    if x[1] is None  
    else (x[1]*math.log10(2/1), "doc 2")  
    if x[0] is None  
    else 0)  
both_tfidf.take(40)
```

```
[('TO', 0),  
 ('THAT', 0),  
 ('AS', 0),  
 ...  
 ('COULD', 0),  
 ('SAME', 0),  
 ('STUBB', (0.00029244824872879597, 'doc 1')),  
 ('NEVER', 0),  
 ('CAN', 0),  
 ('HAND', 0),  
 ('TOO', 0),  
 ('FAR', 0),  
 ('SHIPS', 0),  
 ('PART', 0),  
 ('STARBUCK', (0.00020569326585231104, 'doc 1'))]
```

5.1.2.5 Information retrieval: get most important words

If a term occurs in both documents we already know from above, that the **TF-idf** is 0. It means, that this term is not very interesting and therefore we will get rid of them with the `filter` function:

```
both_tfidf_not0 = both_tfidf.filter(lambda x: x[1]!=0)  
both_tfidf_not0.take(7)
```

```
[('STUBB', (0.00029244824872879597, 'doc 1')),
 ('STARBUCK', (0.00020569326585231104, 'doc 1')),
 ('TOWARDS', (0.00015951722657934328, 'doc 1')),
 ('WHALING', (0.00015112158307516732, 'doc 1')),
 ('NANTUCKET', (0.00012173683081055145, 'doc 1')),
 ('LEVIATHAN', (0.00011753900905846347, 'doc 1')),
 ('OIL', (0.00010354626988483686, 'doc 1'))]
```

Now let's sort the **TF-idf** and see, which terms are most interesting for terms of both documents (the 'doc 1' and 'doc 2' will let you know, in which document this term occurred:

```
both_tfidf_not0Sort = both_tfidf_not0.sortBy(lambda a: a[1],
    ↵ ascending=False)
both_tfidf_not0Sort.take(7)
```

```
[('WHALE', (0.0012075733906839758, 'doc 1')),
 ('HUCK', (0.0008561253939522759, 'doc 2')),
 ('AHAB', (0.0005401197321019868, 'doc 1')),
 ('WHALES', (0.0005177313494241844, 'doc 1')),
 ('TOMS', (0.0003832180334833997, 'doc 2')),
 ('BECKY', (0.0003791412458931507, 'doc 2')),
 ('SPERM', (0.0003204337270760492, 'doc 1'))]
```

But more interesting would be to have the top 30 words of documents 1. Here they are:

```
tmp = both_tfidf_not0Sort.mapValues(lambda x: x[0] if x[1]=='doc 1' else 0)
doc1_tfidf_not0Sort = tmp.filter(lambda x: x[1]!=0)
doc1_tfidf_not0Sort.take(10)
```

```
[('WHALE', 0.0012075733906839758),
 ('AHAB', 0.0005401197321019868),
 ('WHALES', 0.0005177313494241844),
 ('SPERM', 0.0003204337270760492),
 ('STUBB', 0.00029244824872879597),
 ('QUEEQUEG', 0.00028405260522462),
 ('STARBUCK', 0.00020569326585231104),
 ('AYE', 0.00016651359616615659),
 ('PEQUOD', 0.00016651359616615659),
 ('CREW', 0.00016371504833143124)]
```

You remember: the first document was the “Moby Dick” text. It is not surprising, that ‘WHALE’ and ‘AHAB’ (the captain) are in the top 20 words in document 1. As “Moby Dick” is “Sperm Whale” we will find the word ‘SPERM’ very often. Therefor we are able to extract the most important words of the “Moby Dick” text by using the **TF-idf** measure. This is what we call **information retrieval**: I extracted the top 20 most important words from the “Moby Dick” text.

The second document was “Tom Sawyer”. Let’s have a look into the top 30 words of this document 2:

```
tmp = both_tfidf_not0Sort.mapValues(lambda x: x[0] if x[1]=='doc 2' else 0)
doc2_tfidf_not0Sort = tmp.filter(lambda x: x[1]!=0)
doc2_tfidf_not0Sort.take(10)
```

```
[('HUCK', 0.0008561253939522759),
 ('TOMS', 0.0003832180334833997),
 ('BECKY', 0.0003791412458931507),
 ('SID', 0.0002976054940881721),
 ('INJUN', 0.00028129834372717636),
 ('POLLY', 0.00018345544156120197),
 ('POTTER', 0.00017122507879045517),
 ('THATCHER', 0.00015491792842945945),
 ('SAWYER', 0.0001345339904782148),
 ('HARPER', 0.0001345339904782148)]
```

You can see ‘TOM’ and his best friend ‘HUCK’, as well as ‘BECKY’ (the girl to whom Tom fell in love), which is also not surprising. Again: we found the most important top 20 words for the “Tom Sawyer” text.

If you like to have the result as a dictionary of keys, you can use the following code (I won’t show the whole dictionary, it’s very long):

```
# Moby Dick
doc1_tfidf_not0Sort.collectAsMap().keys()

dict_keys(['WHALE', 'AHAB', 'WHALES', 'SPERM', 'STUBB', 'QUEEQUEG',
           'STARBUCK', 'AYE', 'PEQUOD', 'CREW', 'TOWARDS', 'WHALING', 'FLASK',
           'NANTUCKET', 'LEVIATHAN', 'MOBY', 'ERE', 'OIL', 'SAILOR', 'AHABS',
           'HARPOONEER', 'CABIN', ....])
```

```
# Tom Sawyer
doc2_tfidf_not0Sort.collectAsMap().keys()

dict_keys(['HUCK', 'TOMS', 'BECKY', 'SID', 'INJUN', 'POLLY', 'POTTER',
↪ 'THATCHER', 'SAWYER', 'HARPER', 'TOWARD', 'HUCKLEBERRY', 'HUCKS', 'BEN',
↪ 'MUFF', 'NOBODY', 'HADNT', 'JIM', 'WARNT', 'TAVERN', 'DOUGLAS',
↪ 'WELSHMAN', 'HANTED', ...]
```

On my Docker machine-learning repository⁶ you will find the Jupyter-Notebook, with this example.

5.1.2.6 Conclusion

By using the **TF-idf** measure I was able to extract the most important terms for the two documents “Moby Dick” and “Tom Sawyer” as shown below.

We can say two things now:

One: We can say, that the words on each of these lists “**describe the documents**”. The list has a very “close relationship” to the document and these words are the “information”, which we will find in the document. Therefore we have the **information retrieval** part here. For example if you haven’t read “Moby Dick” then you would know from this list, that it’s about a WHALE, because WHALE is part of the list and the two main characters in “Moby Dick” are AHAB, STUBB and QUEEQUEG because these words are also on the list.

Two: We can also say, that for one document the most important words on the lists are the “**most different**” words to all other documents. These words “separate” the documents. Therefore these words could be helpful to “**classify documents**”. For example: image you do the same TF-idf calculation for thousands of documents. Then you will get thousands of lists like shown here. You would expect to have some lists which have a big overlap and others which don’t. Now you can say: the smaller the overlap of one list to all the other, the more “different” this document is compared to all the others. And vice-versa: the bigger the overlap of one list to another, the more “similar” the document is to the other. One application for doing this is to “classify documents” or to “**create a better search algorithms**”. For example: if you are working in a law firm, then you probably have thousands of court judgments. It might be interesting for you to find all court judgments, which are the most similar to your current running legal process. Building a “search algorithm” based on these TF-idf lists can very efficient and helpful for finding the “**most similar**” and “**most relevant**” documents.

⁶ My docker machine-learning repository, <https://hub.docker.com/repository/docker/andreasraut/machine-learning-pyspark>

5.2 K-Means Clustering Algorithm

Additionally I worked on this dataset to show how the K-Means Clustering Algorithm⁷ can be applied by using the Spark Machine-Learning Library. I will show how the “Vermont Vendor Payments” dataset can be clustered. In the images below every color represents a different cluster:

```
from pyspark.ml.clustering import KMeans
```

5.2.0.1 Three Clusters (K=3)

Now I set the parameter K initially to 3:

```
kmeans = KMeans().setK(3).setSeed(1)
```

Do the fit:

```
model = kmeans.fit(dataset)
model.setPredictionCol('prediction')
```

```
KMeansModel:
  uid=KMeans_00206e60dd5b,
  k=3,
  distanceMeasure=euclidean,
  numFeatures=2
```

Do the transform:

```
transformed = model.transform(dataset) \
    .select('features', 'prediction')
transformed
print((transformed.count(), len(transformed.columns)))
```

```
DataFrame[features: vector, prediction: int]
(3387, 2)
```

How many clusters do we have and how many points does each of these cluster have?

⁷ K-Means Clustering Algorithm, <https://spark.apache.org/docs/latest/ml-clustering.html#k-means>

```
model.summary.clusterSizes
```

```
[3385, 1, 1]
```

This means, that the first cluster contains 3385 data point, the second only one and the third also only one. These are extreme values.

```
transformed.sort('prediction', ascending=False) \
    .show(n=transformed.count(), truncate=False)
```

features	prediction
[1.0,2.554683953107E10]	2
[0.0,2.1964140145164E11]	1
[49.0,5948.0]	0
[49.0,3736.0]	0
[49.0,471766.88]	0
[49.0,3916.0]	0
[49.0,501.25]	0
[49.0,2421.4]	0
...	

As expected the two extreme values form two separate clusters (each containing only one single data point). The third cluster is the remaining rest.

5.2.0.2 Five Cluster (K=5)

Let's have a look at the result for K=5 cluster:

```
kmeans = KMeans().setK(5).setSeed(1)
model = kmeans.fit(dataset)
model.setPredictionCol('prediction')
model.summary.clusterSizes
transformed = model.transform(dataset) \
    .select('features', 'prediction')
transformed.sort('prediction', ascending=False) \
    .show(n=transformed.count(),
          truncate=False)
```

```
[3373, 1, 1, 12, 0]
+-----+
| features           | prediction |
+-----+
|[7.0,2.32484304242E9] | 3
|[2.0,2.33823407921E9] | 3
|[9.0,2.32676185131E9] | 3
...
|[0.0,1.49675535347E9] | 3
|[0.0,1.66099306197E9] | 3
|[1.0,2.554683953107E10] | 2
|[0.0,2.1964140145164E11] | 1
|[64.0,1915533.81]      | 0
|[64.0,40981.07]        | 0
...
...
```

Hide the cluster “prediction=0” and only show the other clusters:

```
transformed.filter(F.col('prediction') != '0') \
    .sort('features',
          ascending=True) \
    .show(n=transformed.count(),
          truncate=False)
```

```
+-----+
| features           | prediction |
+-----+
|[0.0,1.49603759082E9] | 3
|[0.0,1.49675535347E9] | 3
|[0.0,1.66099306197E9] | 3
|[0.0,2.17201705172E9] | 3
|[0.0,3.12255843328E9] | 3
|[0.0,4.58524599324E9] | 3
|[0.0,5.05357651587E9] | 3
|[0.0,6.11474963014E9] | 3
|[0.0,2.1964140145164E11] | 1
|[1.0,2.554683953107E10] | 2
|[2.0,2.33823407921E9] | 3
|[7.0,2.32484304242E9] | 3
|[9.0,2.32676185131E9] | 3
|[22.0,1.83834865364E9] | 3
+-----+
```

As expected we have a **cluster number 1** (see “prediction=1”), with the extreme value $2 * 10^{11}$.

And another **cluster number 2** (see “prediction=2”) with the other extreme value $2.5 * 10^{10}$.

And the **cluster number 3** (see “prediction=3”) has values around $1 * 10^9$ and $6 * 10^8$.

Therefore the K-Means algorithm found exactly the clusters, which we expected. Great!

5.2.0.3 Ten Clusters (K=10)

Let's have a look at K=10 clusters:

```
kmeans = KMeans().setK(10).setSeed(1)
model = kmeans.fit(dataset)
model.setPredictionCol('prediction')
model.summary.clusterSizes
transformed = model.transform(dataset) \
    .select('features', 'prediction')
transformed.filter(F.col('prediction') != '0') \
    .sort('features', ascending=True) \
    .show(n=transformed.count(),
          truncate=False)
```

[3363, 1, 1, 1, 18, 3, 0, 0, 0, 0]	
+	- - - - -
features	prediction
+	- - - - -
[0.0,7.826641758E8]	4
[0.0,8.3250279165E8]	4
[0.0,8.7528370872E8]	4
[0.0,9.3917221933E8]	4
[0.0,9.9943549988E8]	4
[0.0,1.04293246515E9]	4
[0.0,1.2885636375E9]	4
[0.0,1.35549635179E9]	4
[0.0,1.49603759082E9]	4
[0.0,1.49675535347E9]	4
[0.0,1.66099306197E9]	4
[0.0,2.17201705172E9]	4
[0.0,3.12255843328E9]	5
[0.0,4.58524599324E9]	5
[0.0,5.05357651587E9]	5
[0.0,6.11474963014E9]	3
[0.0,2.1964140145164E11]	1

```
[1.0,2.554683953107E10] |2
[2.0,2.33823407921E9] |4
[4.0,1.16802981964E9] |4
[5.0,1.17368104165E9] |4
[7.0,2.32484304242E9] |4
[9.0,2.32676185131E9] |4
[22.0,1.83834865364E9] |4
+-----+
```

This shows that even if I take K=10 the K-Means algorithm will only find 6 clusters (named with number 0,...,5). This is also what we expected, so the K-Means algorithm works perfectly fine. Great.

The first cluster has 3363 data points, the second cluster has one data point and so on: [3363, 1, 1, 1, 18, 3, 0, 0, 0]

I want to transform the result into vectors and plot it with colors: each cluster should have a separate color. Therefore I create the function “ith_”:

```
from pyspark.ml.linalg import Vectors
from pyspark.sql.types import DoubleType
from pyspark.sql.functions import lit, udf

def ith_(v, i):
    try:
        return float(v[i])
    except ValueError:
        return None

ith = udf(ith_, DoubleType())

x = transformed.select(ith("features", lit(0))) \
    .toPandas()
x_name = features[0]
y = transformed.select(ith("features", lit(1))) \
    .toPandas()
y_name = features[1]
z = transformed.select("prediction") \
    .toPandas()
```

I will plot the same data as above but as we know from the K-Means Algorithm for each datapoint (State_index, sumofAmount) to which cluster it belongs, I can plot them using different colors:

- violett

- yellow
- blue
- ...

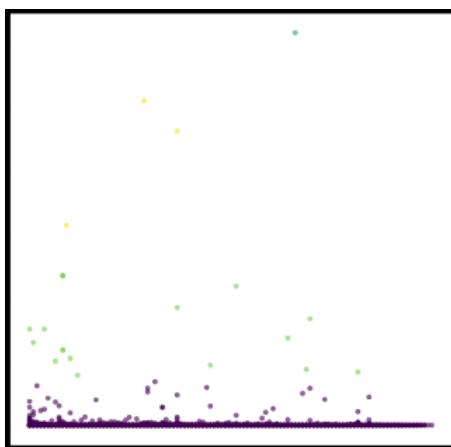
```
my_plotter(x, x_name, y, y_name, z['prediction'])
```

And again: as the extreme values ($2 * 10^{11}, \dots$) create separate cluster (number 1 and number 2), which are disturbing a bit, I filter them out in order to plot only the remaining clusters:

```
transformed_betterVis = transformed.filter((F.col('prediction') >= '3') |  
                                         (F.col('prediction') == '0')) \  
                                         .sort('features', ascending=True)  
transformed_betterVis.show(n=transformed.count(), truncate=False)
```

Now I can plot the remaining data points:

```
x = transformed_betterVis.select(ith("features", lit(0))) \  
                           .toPandas()  
x_name = features[0]  
y = transformed_betterVis.select(ith("features", lit(1))) \  
                           .toPandas()  
y_name = features[1]  
z = transformed_betterVis.select("prediction").toPandas()  
my_plotter(x, x_name,  
           y, y_name,  
           z['prediction'])
```



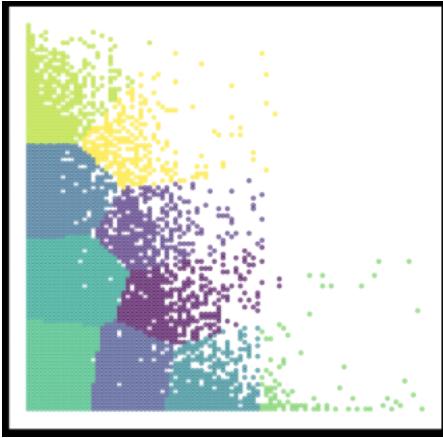
Each color represents a separate cluster and the plot proves, that these clusters make sense:

The violet cluster is **cluster number 3** and contains a data point at about $6 * 10^9$.

The blue cluster is the the **cluster number 4** and contains values between $1 * 10^{e9}$ and $2.5 * 10^{e9}$.

The yellow cluster is **cluster number 5** and contains values between $3 * 10^{e9}$ and $5 * 10^{e9}$.

Of course you can also use other features for the clustering. Another clustering (based on State and Department) results in the following image:



On my Docker machine-learning repository⁸ you will find the Jupyter-Notebook, with this example.

⁸ My docker machine-learning repository, <https://hub.docker.com/repository/docker/andreasraut/machine-learning-pyspark>

6 Use Cases of Artificial Intelligence in the Industry

In the following text I explain in an easily understandable way what is meant by "artificial intelligence (AI) in the industry", describe areas of application and illustrate with practical examples and programming applications how AI can be implemented in concrete terms.

In the **first part** (see sec. 6.1), I will explain the basic concepts of AI and describe some areas where AI is already being successfully applied. I will also describe the peculiarities of big data, deep learning and process mining.

In the **second part** (see sec. 6.2), I show an example of how the car manufacturer BMW has benefited from AI techniques.

In the **third part** (see sec. 6.3), I show how AI techniques can be implemented in the Python programming language when dealing with the topic of image recognition and provide my programming code in the process.

In the **fourth part** (see sec. 6.4), I give some recommendations on what to look for when introducing AI techniques in a company.

I think the choice of further articles to delve into the topic "Use Cases of Artificial Intelligence in the Industry" is huge and I hope this short introduction is helpful to get started.

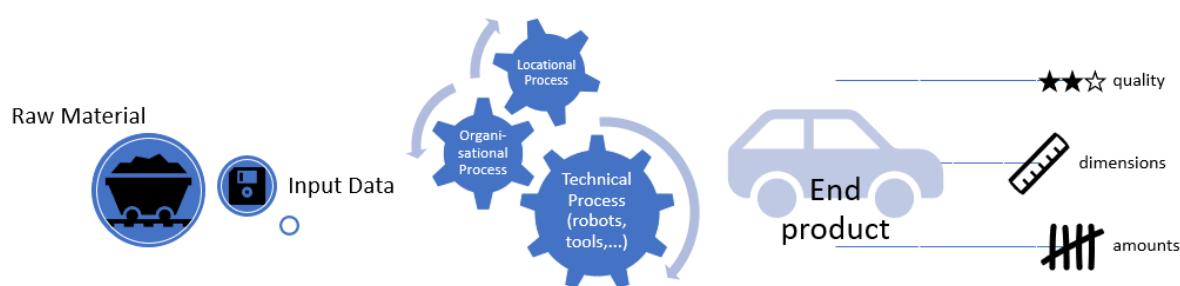


Figure 6.1: Use Cases of AI - Chart

6.1 AI Explained

My diagram fig. 6.1 shows **raw material / input data** on the left and the **end product** on the right. In between, **different processes**. During the runtime of these processes, the intermediate results are usually logged and stored by means of log files (which is also data).

6.1.1 What kind of processes are meant here?

For example, machines could process raw materials that are exposed to a certain temperature and pressure during their processing. Temperature and pressure are determined by sensors, histories with a time stamp and stored in log files. These processes can take place regionally at different locations or organizationally in different company units. There might also be technically differences in these processes (e.g. robots, tools,...).

I would like to mention a few examples (there are a lot more) where AI is successfully applied in practice:

- **Sales forecasts:** Artificial intelligence calculates the expected sales of products based on a large number of input data (e.g. stock market data, weather, commodity prices, customs restrictions, price development on the sales markets, inflation, interest rates or social media trends). This makes it possible to better determine expected sales and optimally control production.¹
- **Automatic orders:** The order quantities and order times for raw materials are automatically determined and optimized by artificial intelligence. This is to prevent storage capacities from being exceeded or delivery bottlenecks from occurring. In addition, as many supplier discount offers as possible are to be optimally utilized.²
- **Product development for series production:** Automated tests are carried out on the products and validated by the artificial intelligence so that it can point out where adjustments still need to be made to the products so that they can be produced in series cost-effectively and without errors.³
- **Quality control:** Images of the products are generated using sensors, X-rays or high-resolution cameras. Artificial intelligence can then use image recognition algorithms to detect defects in the products and sort them out.⁴⁵

¹ Big Data Insider, This is the role of machine learning in sales forecasting, <https://www.bigdata-insider.de/diese-rolle-spielt-machine-learning-bei-absatzprognosen-a-625751/>

² Procurement up-to-date, Optimal raw material inventory thanks to precise AI-based demand forecasts, <https://beschaffung-aktuell.industrie.de/logistik/optimaler-rohstoffbestand-dank-praeziser-bedarfsprognosen-auf-ki-basis/>

³ Intel, Artificial intelligence reduces costs and accelerates time to market, <https://www.intel.com/content/dam/www/public/us/en/documents/white-papers/artificial-intelligence-reduces-costs-and-accelerates-time-to-market-paper.pdf>

⁴ Fraunhofer Institute, AI-based visual quality control, <https://www.iais.fraunhofer.de/de/geschaeftsfelder/Computer-Vision/visuelle-qualitaetskontrolle.html>

⁵ Elektronik Praxis, AI in quality control: when no detail should be overlooked, <https://www.elektronikpraxis.vogel.de/ki>

I have linked further articles in the footnotes. The list could go on, but this selection should cover the most important areas.

6.1.2 What data does artificial intelligence have access to?

Artificial intelligence now has access to all data:

- all input data: also includes all data describing the material properties (length, width, weight...)
- All log data: also includes all data resulting from the processing steps, e.g. temperature or pressure with which materials are processed.
- all output data: includes data that measures the quality of the product but as well the dimension (length,...) and amounts. For example, a quality assurance staff might rate the product as not ok because it is defective or because an important KPI metric is not satisfactory.

The artificial intelligence knows everything and can thus establish a **connection** between "input data" and "output data" (or "raw material" and "end product") at any time and always has the process (the log files) in view. For example: as soon as a human evaluates the "output data" or the "end product" as "not ok", the artificial intelligence can draw a conclusion as to which input parameter or which process step was most relevant for the anomaly and can make a suggestion as to what should be changed.

6.1.3 Is Artificial Intelligence Really Intelligent?

Artificial intelligence is not "intelligent" as we humans commonly understand it: AI is only an algorithm that can represent these relationships with models. There are different approaches, depending on what is relevant at the time:

- When the amount of input data is huge, we speak of Big Data⁶. This is the case, for example, with sensor data, i.e. when temperature, pressure or travel distances of machines are continuously collected. Each data point in itself is often only a simple number, but over time it adds up to a huge amount of data that can no longer be processed using conventional data processing methods. Big Data approaches are sometimes quite different from conventional approaches to data processing: other systems are used, such as Apache Spark⁷ to calculate with networked computers) or Hadoop⁸ (to process very large data sets distributed across several computers). I have been working in my Machine-Learning Repository⁹ to describe the difference when working with a Big Data system compared to conventional data processing.

in-der-qualitaetskontrolle-wenn-kein-detail-uebersehen-werden-darf-a-860403/

⁶ Big Data https://de.wikipedia.org/wiki/Big_Data

⁷ Apache Spark, <https://spark.apache.org/>

⁸ Hadoop, <https://hadoop.apache.org/>

⁹ Machine-Learning Repository, <https://github.com/AndreasTraut/Machine-Learning-with-Python>

- If, on the other hand, the input data takes a back seat (i.e. no Big Data), but the processes come to the fore, this is called *Process mining*¹⁰. Here, the log files that record the processes are often transformed into models and then evaluated.
- One speaks of a *Deep Learning*¹¹ approach when neural networks are used. This is often the case when the input data are not just simple data points, as is the case with sensor data, but have a complex structure, such as images or documents. An image consists of several thousand pixels in each of the two image axes and for each of these pixels there are many possible shades of color. A document has sentences, words and letters that follow grammatical and orthographic rules. In the Deep Learning approach, groups of pixels or letters are linked to so-called "neurons", which together form a layer. A neuron layer then passes on data to the next neuron layer above it according to a given arithmetic operation, and when hundreds of such layers are stacked on top of each other, you get a neural network (hence the name "deep"). I have also dealt with this and documented my experiences[^repodeeplearning].

AI approaches can therefore be used to penetrate the **interrelationships of** the input data, the log files and the output data. I will explain in the next section how profit can be generated with this.

6.2 Example: Benefits of AI for BMW

The "Capgemini Research Institute" published in December 2019 an interesting study¹²: 300 companies from the industrial manufacturing, automotive, consumer goods, aerospace and defense sectors were examined and it was found that companies in Germany already use a great deal of artificial intelligence in their value chains and production processes compared to other countries, but should deepen this.

I would like to briefly pick out a concrete example from the CBR Exclusive Article "How BMW optimised supply chain big data with Teradata"¹³, on which this study is based:

¹⁰ Process mining, Wikipedia, <https://en.wikipedia.org/wiki/Process-Mining>

¹¹ Deep-Learning, Wikipedia, https://en.wikipedia.org/wiki/Deep_learning

¹² "Capgemini Research Institute" published an interesting study, <https://www.capgemini.com/de-de/news/ki-in-der-industrie/>

¹³ CBR-Online: "How BMW optimised supply chain big data with Teradata", <https://www.cbronline.com/big-data/analytics/bmw-optimise>

CBR Exclusive: How BMW optimised supply chain big data with Teradata

JAMES NUNNS EDITOR
10TH NOVEMBER 2016

+ INCREASE / DECREASE TEXT SIZE -



Figure 6.2: CBR-Online: “How BMW optimised supply chain big data with Teradata”, Source: BMW Group

The car manufacturer BMW works at 31 production sites where very complex processes take place. In the picture above I have tried to show the countless different raw materials in red. The process from the raw materials to the final product (the car) is very long, complicated and confusing and is sometimes scattered over different continents. Inventory is stored temporarily at many points in the process. A major challenge for BMW was to store the large amount of data that is generated in a meaningful form (keyword: Data Warehouse¹⁴ and Data Lake¹⁵).

BMW celebrated a success in 2016 when it was able to gain valuable insights from analyzing its inventory: the teams were able to reduce inventory costs by 70% by gaining more transparency across their many production sites and optimizing processes.

¹⁴ Data Warehouse, https://de.wikipedia.org/wiki/Data_Warehouse

¹⁵ Data Lake, <https://de.wikipedia.org/wiki/Data-Lake>



In 2016, Klaus Straub, CIO at BMW, described in an article¹⁶ his ideas for the digital transformation of the company. Even then in 2016, he saw the great potential that artificial intelligence would create, for example, to **improve quality** or **make processes more efficient**, although linking IT with real production processes would be a major challenge.

But how can this be implemented in concrete terms? I would like to give an insight into this in the following section.

6.3 Implementation of AI Techniques

There are many free and freely available (Open Source¹⁷) tools that you only have to adapt to your own needs depending on the question. How exactly this is done is shown below in my program code.

6.3.1 What can be seen in the picture?

For example, if there is a picture of a component and the question is what is in that picture, a human could quickly find out by just looking at it. The human could also see whether the component is defective or not. The AI can do that too, and for this question I choose a deep-learning approach and use

¹⁶ Global Intelligence for Digital Leaders, Klaus Straub, CIO BWM, “Powering Digital Disruption at BMW”, <https://www.i-cio.com/profession/cio-profiles/item/powering-digital-disruption-at-bmw>

¹⁷ Open Source, Wikipedia https://en.wikipedia.org/wiki/Open_Source

the ResNet50¹⁸ network, which is already trained on the *ImageNet*¹⁹ image data set. This saves me a lot of programming effort and so I only need 10 lines of code for the model to tell me that in fig. 6.3 the program sees a "cup" with 86% probability and a "coffee cup (coffeepot)" with 6.8%. In fig. 6.5²⁰ the model sees a "switch component" with 76% probability. You can download my program code from my GitHub repository²¹.



Figure 6.3: Use Case - ResNet50: Espressocup

```
resnet = ResNet50(weights='imagenet')
x = preprocess_input(np.expand_dims(img.copy(), axis=0))
preds = resnet.predict(x)
decode_predictions(preds, top=5)

[[('n07930864', 'cup', 0.86062807),
 ('n03063689', 'coffeepot', 0.06782799),
 ('n03063599', 'coffee_mug', 0.05167182),
 ('n03950228', 'pitcher', 0.0064888997),
 ('n04398044', 'teapot', 0.005076931)]]
```

Figure 6.4: Use Case - ResNet50: Code for Prediction of Espressocup



Figure 6.5: Use Case - ResNet50: Network Card, Pixabay²²

¹⁸ ResNet50, Tensorflow https://www.tensorflow.org/api_docs/python/tf/keras/applications/ResNet50

¹⁹ "ImageNet", <http://www.image-net.org/>

²⁰ Pixabay, network map, <https://pixabay.com/de/photos/netzwerkkarte-bauteil-schaltkreis-550544/>

²¹ Jupyter-Notebook "Espresso / Network Card Example", https://github.com/AndreasTraut/Deep-Learning/blob/master/Image_classification/Image_classifier_example_2_transfer_learning_ResNet52-German.ipynb

```
[[('n04372370', 'switch', 0.76144683),
 ('n03777754', 'modem', 0.08100146),
 ('n03272010', 'electric_guitar', 0.032640774),
 ('n03492542', 'hard_disc', 0.015652671),
 ('n04070727', 'refrigerator', 0.008717526)]]
```

Figure 6.6: Use Case - ResNet50: Code for Prediction Network Card

6.3.2 Which groups can be formed?

Let's say we have a picture that we don't know much about and which we want to classify in a grouping that we know. For example, an X-ray picture where we ask ourselves whether and which disease is to be seen on it²³. Or a picture of a plant for which we want to know the name and care instructions. Texts can also be grouped. In the case of a document or contract, we may be looking for similar texts. Grouping similar things is a frequently discussed problem. We all know the useful Google function to show similar images. You enter a term (e.g. espresso cup) in the search bar and similar images are displayed:

²² Pixabay, network map, <https://pixabay.com/de/photos/netzwerkkarte-bauteil-schaltkreis-550544/>

²³ Die Zeit, Artificial Intelligence: When Computers Evaluate X-Ray Images, <https://www.zeit.de/wissen/2019-09/kuenstliche-intelligenz-medizin-diagnose-krankheiten-bilddiagnostik>

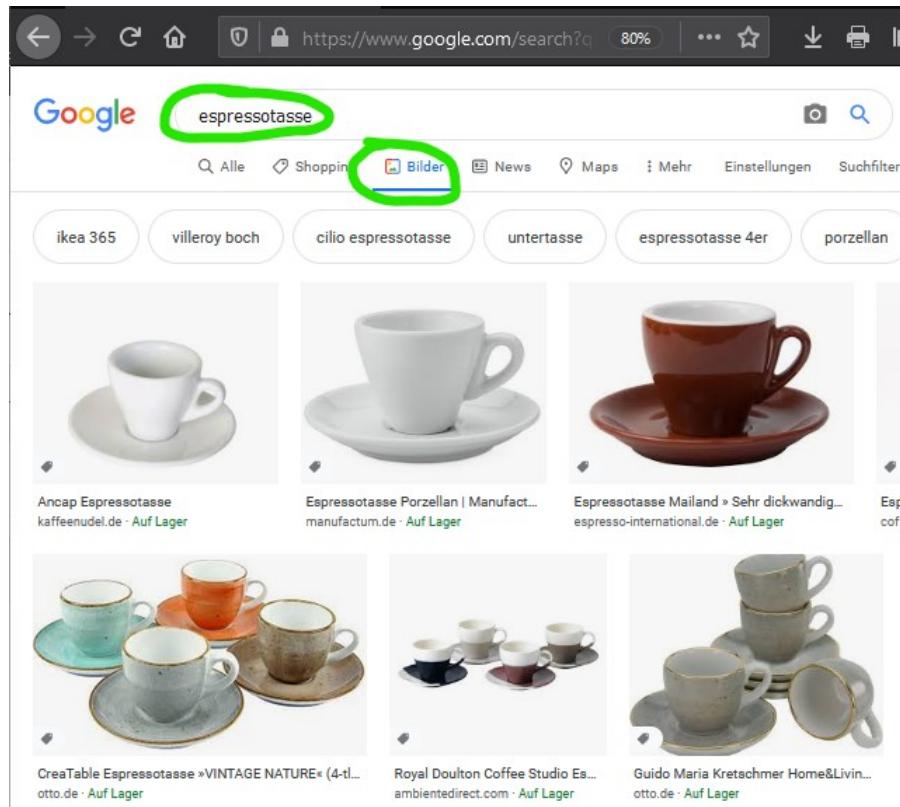


Figure 6.7: Use Case - Google: Similar Pictures

Two questions arise when implementing the program code.

The first question is: How do you compare two pictures? Or two texts? It's not easy, but this problem has been analyzed many times. So there are procedures that you just have to copy, and I will show you how in a moment.

The second question is: How do you go about comparing all the things (pictures or texts) in pairs? Let's take 1 million things that we compare with each other in pairs in order to be able to form the groups. Then we already have 1 million times $999\,999 / 2$, that is about 500 billion arithmetic operations. That can take a very long time. Since this question involves a lot of input data, I choose a Big Data approach, namely "Local-Sensitive-Hashing" (LSH).²⁴ LSH is a technique to classify similar things into groups with a high probability. In other words, one does without absolutely exact results and accepts a small probability of error. This probability can be set with control parameters (as needed). Once these parameters are set, the AI algorithm can very quickly classify new images into groups. The advantage of doing without an absolutely exact 100% grouping is obvious: LSH runs much faster than 100% exact algorithms.

The result of my work was: I applied the LSH algorithm to over 9000 images (each about 300*300 pix-

²⁴ Towards Datascience, Locality Sensitive Hashing An effective way of reducing the dimensionality of your data, <https://towardsdatascience.com/understanding-locality-sensitive-hashing-49f6d1f6134>

els), including my collection of church windows and espresso cups (I like to drink espresso and started photographing the cups at some point). On my own computer, this grouping took a few minutes and was only necessary once. After that, I downloaded a completely new picture of an espresso cup from the internet:

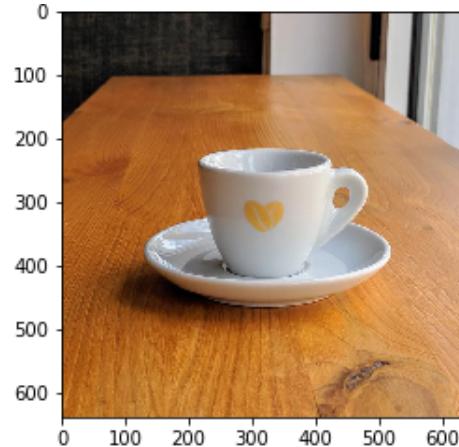


Figure 6.8: Use Case - Espresso Cup

I gave this picture to the program and then used it to search for similar pictures from my picture collection. After a few seconds, the program showed me these 15 pictures:



Figure 6.9: Use Case - Espresso Cup: Similar Pictures

I then tested the same for my church window images with this test image:

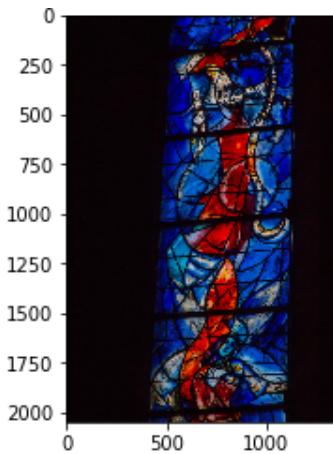


Figure 6.10: Use Case - Church Window

Result:

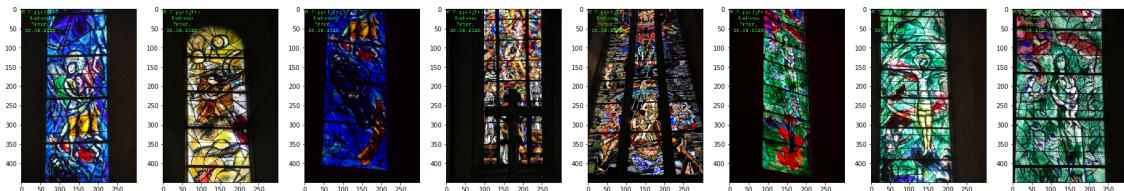


Figure 6.11: Use Case - Church Window: Similar Pictures

If you want to have a quick look at the lines of code to make sure that only a few lines of code are needed for this problem, you can download my program code from my GitHub repository²⁵ and read my further explanations²⁶. I have also explained the deep-learning topic in more depth in my Deep-Learning repository²⁷.

In the next section, I will give you a first concept to introduce AI techniques in your company.

6.4 Recommendations For Implementing AI

If you are now interested in using artificial intelligence in your company as well, the recommendation is that you consider the following:

What concrete benefits do you expect the analysis of your data to bring you? First collect and structure

²⁵ Jupyter-Notebook “Similar Pictures” Example, https://github.com/AndreasTraut/Deep_learning_explorations/blob/master/8_Image_similarity_search/Beispiel_aehnliche_Bilder_finden.ipynb

²⁶ “Similar Pictures” Example: further explanations, https://github.com/AndreasTraut/Deep_learning_explorations/blob/master/README.md

²⁷ Deep-Learning repository, <https://github.com/AndreasTraut/Deep-Learning>

background knowledge about your company: Which company units are affected, who are the key people, who is the "sponsor" of the project?

- Describe the problem and the motivation for the data analysis project. Also think about the current situation, its advantages and disadvantages. You will need this to compare with the new data analysis project.
- Describe what a successfully implemented data analysis project would look like: Are there success metrics (objective goals) or subjective goals that you can define or describe to measure the "success" of the data analysis project for your company? It is important to define measurable business objectives so that further measurable objectives can be derived from them for the further implementation of the data analysis project: What kind of data analysis should be targeted for the problem? What data is specifically needed for these models and what technical and organizational steps are needed to extract, transform, model and evaluate this data from different sources?
- Early on, also ask yourself how the "deployment" is to proceed, i.e. how the programs that were developed in a test environment are to be made to run in daily productive operation. Should you use your own computers or the cloud? If your company has high data protection requirements, a cloud solution may not be the first choice and should be questioned. An expensive investment in your own hardware could then be the next step for you. If, on the other hand, you want to try out different things first and are not yet ready to invest massive capital in new hardware, then a cloud solution could be ideal for you. You can read about my experiences with the Microsoft Azure Cloud Platform²⁸.

There are methodological approaches that can be applied when transforming into such a data analysis project. Since the costs for Big Data systems are enormous (financial costs but also the time your employees are tied up) and usually many company areas are affected, it is advisable to take a structured approach.

I hope that my brief introduction to the topic of "*artificial intelligence in industry*" was helpful for getting started and I think that the selection of further articles for delving into the topic is huge. I wish you much fun and success in your further research and implementation.

²⁸ My Documented Experiences with the Microsoft Azure Cloud Platform, <https://github.com/AndreasTraut/Experiences-with-MicrosoftAzure>

List of Figures

1.1	Data Science - Adapted version. Creative Commons Legal Code	2
1.2	Data Science	3
1.3	Certificate Data Analysis with Python: Zero to Pandas	7
1.4	Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License	7
2.1	Spyder-IDE: An integrated development environment with debugger.	9
2.2	Jupyter-Notebook: Code and Documentation in one place.	10
3.1	Cross-Industry Standard Process for Data Mining - CRISP-Cycle, Wikipedia, CC BY-SA 3.0	13
3.2	Consumer Price Index	15
3.3	Consumer Price Index Figure	20
3.4	Consumer Price Figure - absolute and relative increments	21
3.5	Consumer price index versus Stockprices, Real estate prices	22
3.6	Last-FM Music Statistics - Overview year 2019	25
3.7	Last-FM Music Statistic - Listening clock 2019	25
3.8	Last-FM Music Statistics - Format of the Database	26
3.9	Last-FM Music Statistics - Overall statistics	27
3.10	Last-FM Music Statistics - Reproduced Statistics of 2019	29
3.11	Last-FM Music Statistics - Overview year 2018	30
3.12	Last-FM Music Statistics - Reproduced Statistics of 2018	30
3.13	Last-FM Music Statistics - Overview year 2017	31
3.14	Last-FM Music Statistics - Reproduced Statistics of 2017	31
3.15	Last-FM Music Statistics - Listening clock 2018 (week 21)	32
3.16	Marathon Example - Seaborn Jointplot	34
3.17	Marathon Example - Seaborn Histogram “size”	35
3.18	Marathon Example - Seaborn Histogram “weight”	36
3.19	Marathon Example - Seaborn PairGrid	37
3.20	Marathon Example - Seaborn Kernel Density “size-to-weight”	38
3.21	Marathon Example - Seaborn Kernel Density “size”	38
3.22	Marathon Example - Seaborn Regression Plots “men”	39
3.23	Marathon Example - Seaborn Regression Plots “women”	39
3.24	Marathon Example - Seaborn Violinplot “size”	41
3.25	Marathon Example - Seaborn Violinplot “weight”	41

3.26 Build Status passing	42
3.27 Pedestrians in Corona-Lockdown - Frequency	43
3.28 Pedestrians in Corona-Lockdown - Barplot “Ulm, Münsterplatz”	44
3.29 Pedestrians in Corona-Lockdown - Barplot “Augsburg”	44
3.30 Pedestrians in Corona-Lockdown - Barplot “München, Maximilianstrasse”	45
3.31 Pedestrians in Corona-Lockdown - Barplot “München, Neuhausstrasse”	45
3.32 Pedestrians in Corona-Lockdown - Barplot “Ulm” 08/2019 to 03/2020	46
3.33 Pedestrians in Corona-Lockdown - Barplot “München, Maximilianstrasse” 08/2019 to 03/2020	46
3.34 Pedestrians in Corona-Lockdown - Barplot “Ulm” April 2020	47
3.35 Deutsche Bahn - Elevators 2	50
3.36 Deutsche Bahn - Elevators 1	51
3.37 Deutsche Bahn - Latitude and Longitude	52
3.38 Interactive Pivots - Grouping and aggregating	53
3.39 Interactive Pivots - Graph	54
3.40 Big Data Visualization - Data Format	60
3.41 Big Data Visualization - Line Chart	61
3.42 Big Data Visualization - Stacked Barplot	62
3.43 Data App - Comparison of Streamlit, RStudio-Shiny and others, Quelle: DataRevenue-Blog	63
3.44 Data App - Marathon Example	64
3.45 Data App - SEIR Model Example	65
3.46 Professional BI - Power BI	66
3.47 Professional BI - Power BI Prices	67
3.48 Professional BI - Tableau 1	68
3.49 Professional BI - Tableau 2	69
3.50 Professional BI - Tableau Prices	69
3.51 Professional BI - QLink	70
3.52 Professional BI - QLink Prices	71
4.1 Movies Database - NaN Values	75
4.2 Movies Database - Plot Null and NotNull	76
4.3 Movies Database - Histogram	77
4.4 Movies Database - Scatterplot Revenue Year	77
4.5 Movies Database - Scatterplot Year Score	79
4.6 Movies Database - Stratified Sample (Train, Test, NaNs)	82
4.7 Movies Database - Plot True vs Predicted Value Side-by-Side Testdataset	88
4.8 Movies Database - Plot True vs Predicted Value Side-by-Side Trainingdataset	89
4.9 Mind Map - Scikit-Learn “Small Data”	91

4.10 Small Data - Optimize Model GridSearchCV on LinearRegressor (10.1.2)	102
4.11 Small Data - Optimize Model Randomized Search (10.2)	103
4.12 Small Data - Optimize Model Analyze best models (10.3)	104
4.13 Small Data - Evaluate final model on test dataset (11)	105
4.14 Mind Map - Apache Spark ML “Big Data”	106
4.15 Big Data - Run Docker	107
4.16 Big Data - Docker Dashbord	108
4.17 Big Data - Docker Localhost	108
4.18 Big Data - Docker Localhost Datafiles	109
4.19 Big Data - Docker Jupyter Notebook	109
4.20 Big Data - Price Box Plot	114
4.21 Mind Map - Scikit-Learn and Apache SparkML	121
5.1 Map-Reduce - Word Count Example “Moby Dick”	124
5.2 Map-Reduce - Word Count Example Reduce by Key	127
5.3 TF-idf example - Moby Dick and Tom Sawyer	130
6.1 Use Cases of AI - Chart	147
6.2 CBR-Online: “How BMW optimised supply chain big data with Teradata”, Source: BMW Group	151
6.3 Use Case - ResNet50: Espressocup	153
6.4 Use Case - ResNet50: Code for Prediction of Espressocup	153
6.5 Use Case - ResNet50: Network Card, Pixabay	153
6.6 Use Case - ResNet50: Code for Prediction Network Card	154
6.7 Use Case - Google: Similar Pictures	155
6.8 Use Case - Espresso Cup	156
6.9 Use Case - Espresso Cup: Similar Pictures	156
6.10 Use Case - Church Window	157
6.11 Use Case - Church Window: Similar Pictures	157